



# Segment Routing Bootcamp

# Prerequisite

Please read [www.segment-routing.net/tutorial](http://www.segment-routing.net/tutorial)

SR Concepts >

- SR Introduction
- SR MPLS data plane
- SR Global Block (SRGB)

SR IGP >

- SR IGP Control Plane
- Topology Independent LFA (TI-LFA)

SR and LDP >

- SR and LDP co-existence
- SR Mapping Server
- SR and LDP interworking

SR TE >

- SR TE

Please view [www.segment-routing.net/demos](http://www.segment-routing.net/demos)

SR On-Demand Next Hops (ODN)

SR TI-LFA

SR Microloop Avoidance

Please view

<https://www.youtube.com/watch?v=EUxm9EH1wyQ&t=636s>

Segment Routing - Deployment Experience and Technology Update

# Agenda of Day1 and Day2

## Day-1

- Cisco SD IP Network Framework
- Segment Routing (SR) - TI-LFA
- SR Mapping Server
- Use Case P-PE (SR + LDP/Rsvp-TE)
- Lab Hands-On

## Lunch & Pray

- SR Migration
- Use-case P(SR) + PE(SR)
- SR Policy Concept
- Lab Hands-On

## Day-2

- VPN SR TE based on Low Latency/ Constraint/WECMP
- SR PCE - BGP LS + PCEP
- Use-Case Seamless MPLS - BGP LU
- Lab Hands-on

## Lunch & Pray

- SR On Demand NextHop (ODN)
- SR PCE Rest API
- Lab Hands-on
- Demo SR ODN + SR PCE + NSO Automation

# Housekeeping

Please introduce yourself : Name – Company – Role

Please access and reserve <http://dcloud.cisco.com>  
Cisco WAN Automation Engine 7.1.1.1 - Segment Routing and XTC  
Sandbox v1

Please access dcloud vpn and test Telnet Session to every nodes

Please have python installed in the laptop

And respect others, please be punctual and set HP to silent mode

# SR Champion and SR Redelivery Session

Congratulations ! YOU are Partner SR Champion

We will need your help to run SR Redelivery Session to your colleague engineer

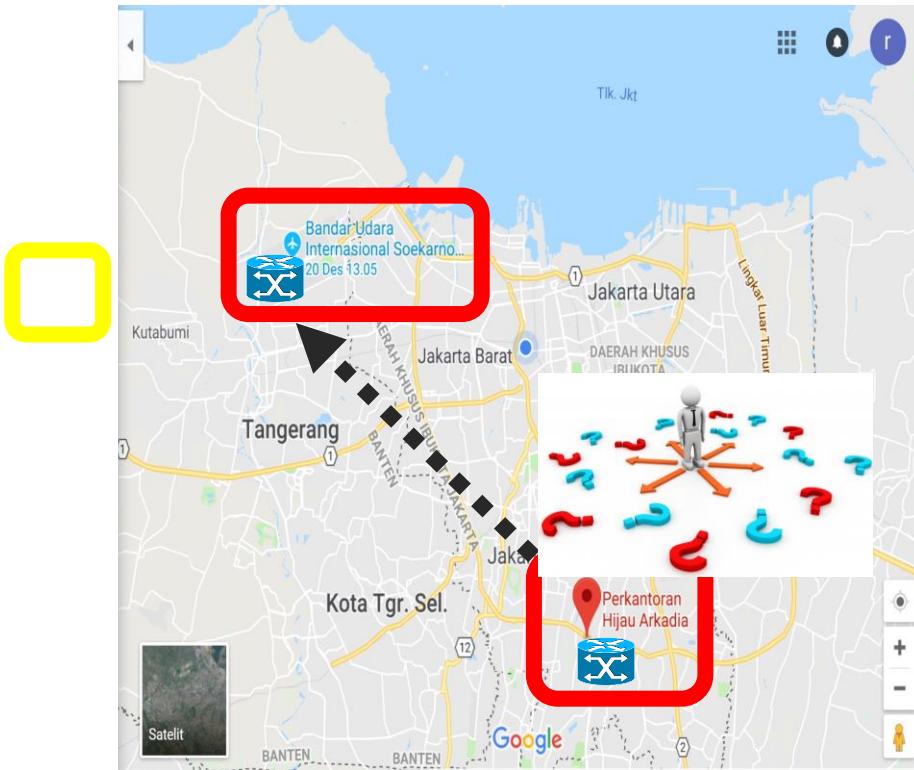
Every 2 SR Champions will deliver 1 SR Redelivery Session

Please advice time and place



# Cisco SD IP Network Framework

# Objective : Go to the Airport ASAP (Business Requirement)



- What is the Best Path ?
- Alternative Path ?
- How is Situation along the Path ?
- What if the Path is Blocked ?
- How about Tomorrow Path ?



# Objective : Go to the Airport ASAP (Business Requirement)

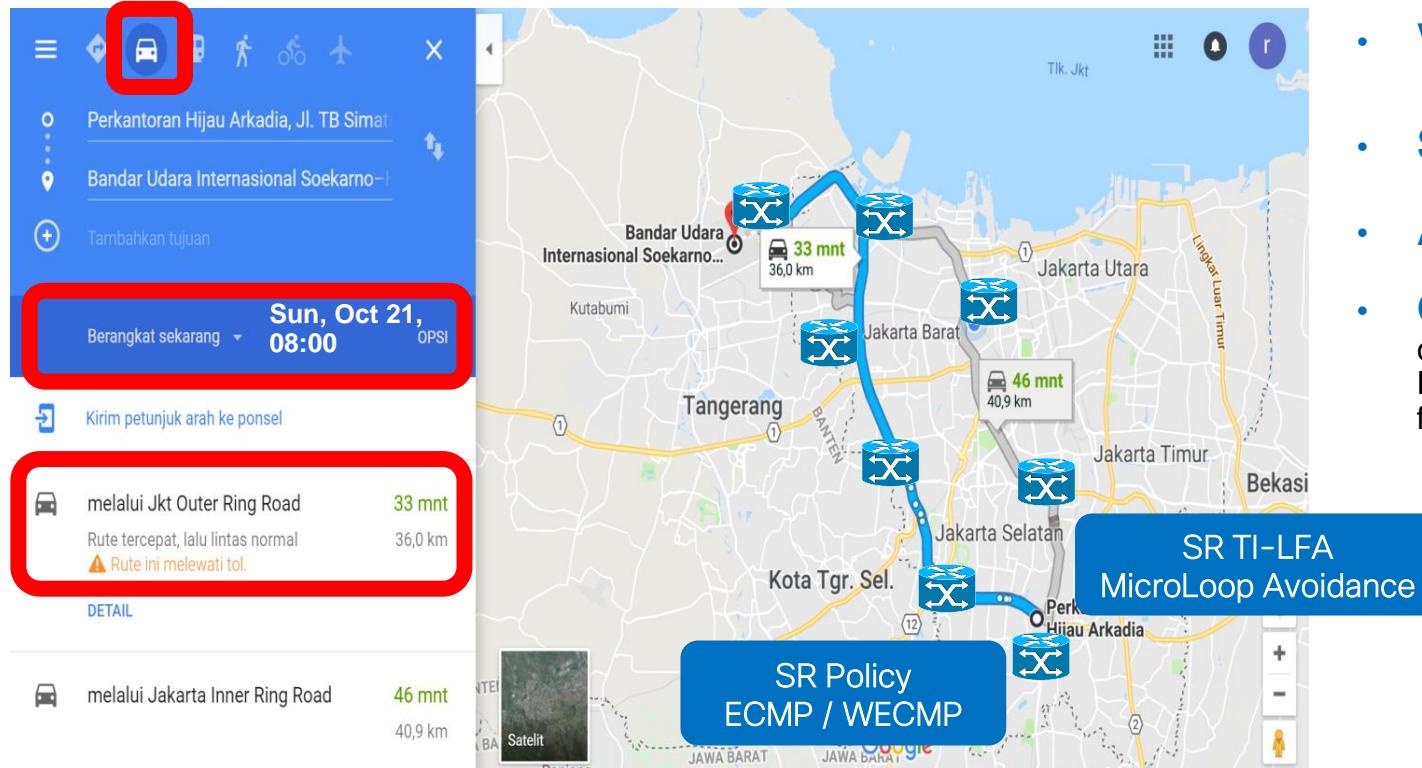
How is the Path From A (Telkom) to Z (Airport) NOW ? (Plan/Design)

The image displays a mobile application interface for navigation. At the top, there's a toolbar with icons for bus, car, walk, bike, and plane. Below it, the starting point is listed as 'Perkantoran Hijau Arkadia, Jl. TB Simatupang' and the destination as 'Bandar Udara Internasional Soekarno-Hatta'. A red box highlights the departure time 'Berangkat sekarang ▾ Thu, Oct 11, 08:00'. Another red box highlights the first route entry: '09.13–11.07 1 jam 54 mnt Loop Line > KA Bandara'. The map shows the path from Telkom to the airport, with several segments highlighted in red, labeled 'IGP Best Path'. The map includes labels for Bandara Internasional Soekarno Hatta, Bandar Udara Internasional Soekarno..., Kutabumi, Tangerang, Stasiun BNI City, Stasiun Sudirman, Jakarta Timur, Bekasi, Kota Tgr. Sel., Perkantoran Hijau Arkadia, Tanjung Barat, Satelite, and various roads like Jl. TB Simatupang, Jl. P. R. Soekarno, and Jl. Prof. Dr. Satrio. A red box also highlights the 'Kirim petunjuk arah ke ponsel' button.

- **Visualization** for overall network visibility
- **Simplification** for only highlight On-Demand Network Path from A to Z

But... it takes long time. (Latency Analysis)

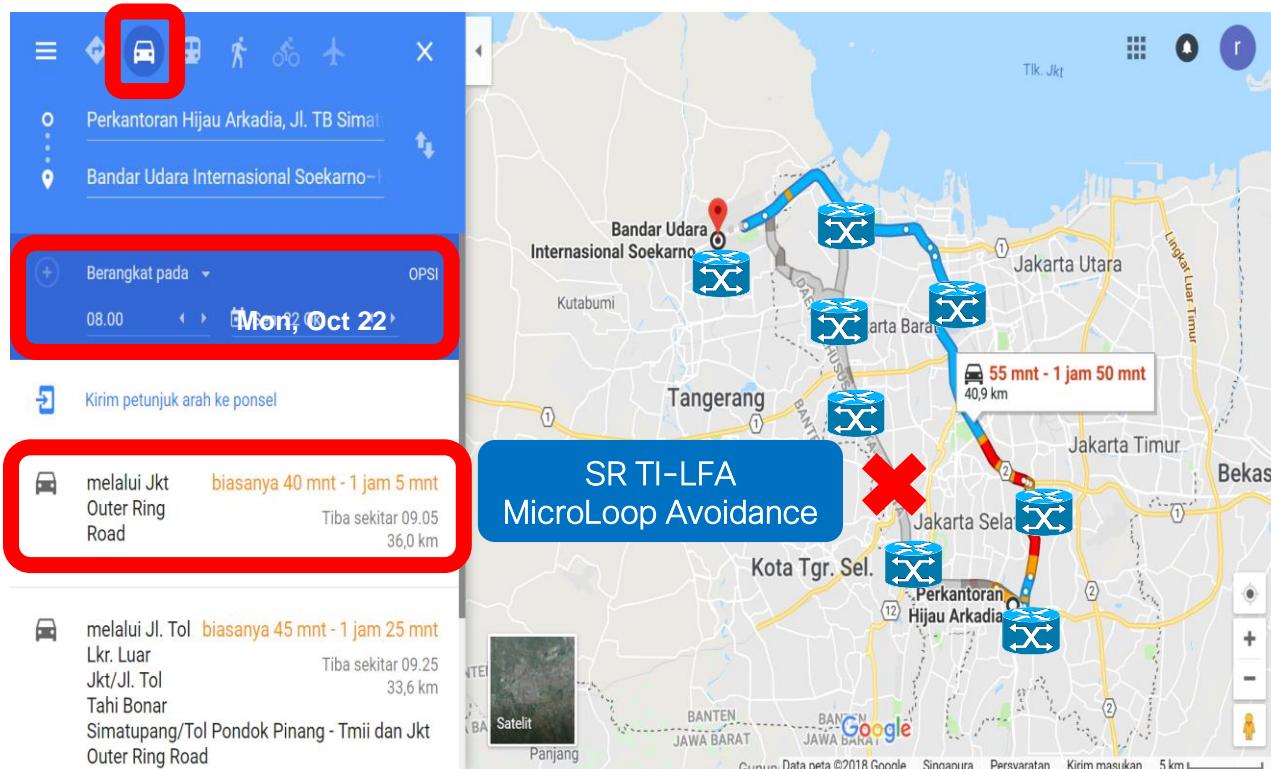
I need shorter time. How to get faster and better Path ? (Optimization)



- **Visualization**
- **Simplification**
- **Analysis**
- **Optimization** for calculating better On-Demand SLA Route Path from A to Z

# Good Optimization, .... What if the Path is blocked ? 😞 (Simulation)

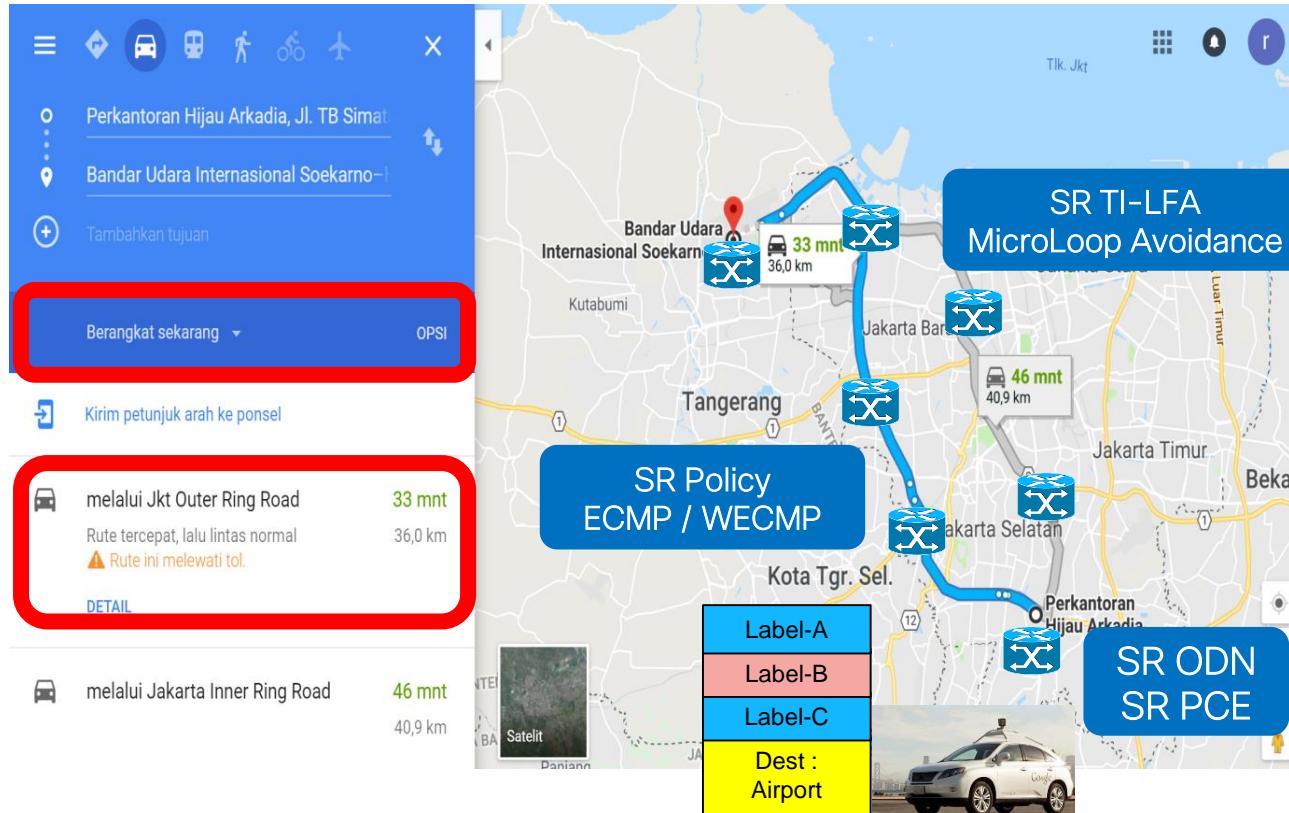
## What will the Path Condition be NEXT DAY ? (Prediction)



- **Visualization**
- **Simplification**
- **Analysis**
- **Optimization** for best SLA Path Requirement
- **Simulation** for critical failure analysis and solutions
- **Prediction** for network traffic growth planning

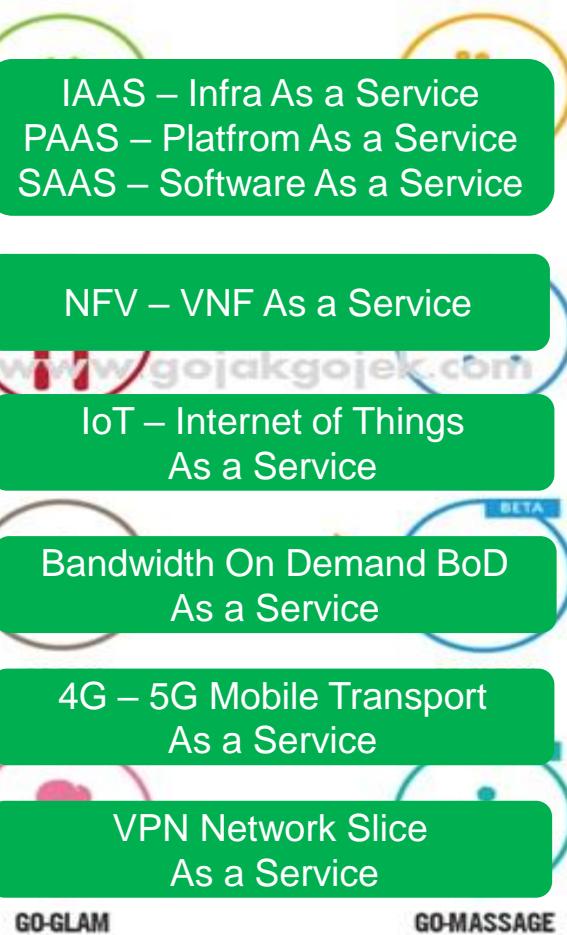
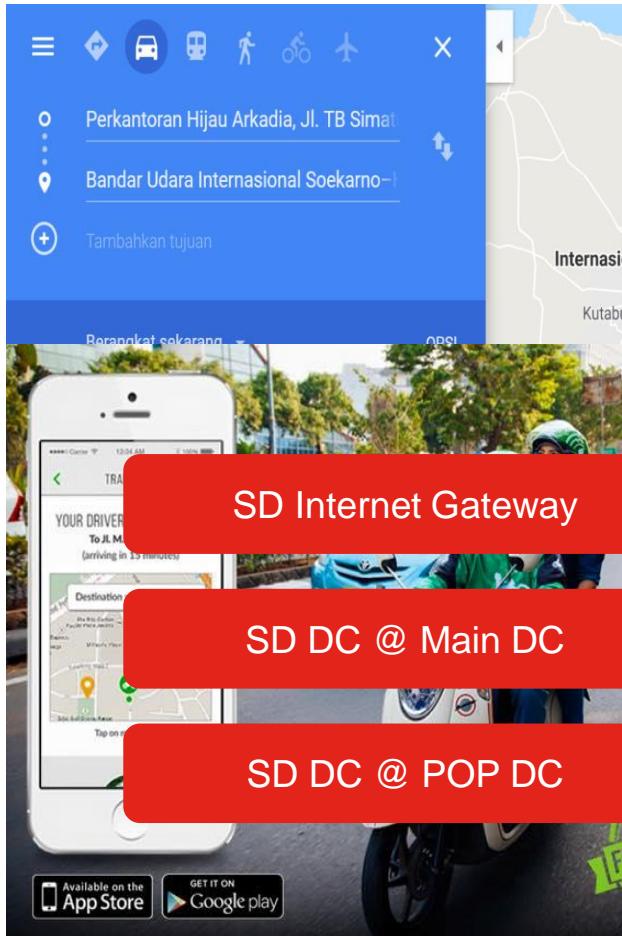
# Next Level Transport : Automated Intelligent Transportation

## (Segment Routing + SDN Automation with Intent-based Networking)



- **Visualization**
- **Simplification**
- **Analysis**
- **Optimization**
- **Simulation**
- **Prediction**
- **On-Demand** for business SLA requirement fulfillment

# More Digital Value Add Service ... More Revenue with Automation



- **Visualization** for overall network visibility
- **Simplification** for network complexity
- **Analysis & Simulation** for critical failure analysis and solutions
- **Prediction** for network traffic growth planning
- **Optimization** for better SLA network performance
- **On-Demand** for instant business SLA requirement fulfillment
- **Value Add Service** for more revenue generation

# Cisco SD IP Network

## Scalable Platform

- Scalable fabric designs to accommodate elastic network service expansion
- Granular building blocks using high-density and efficient Cisco routers

## Network Simplification

- Segment Routing - EVPN into building blocks with simplified feature sets
- Increased visibility and path control with XTC SR PCE
- Consistent operational model

## SDN Automation

- NSO enables modern Service Automation to transform network operations
- Common automatable tasks prescribed
- Standards-based YANG models used for configuration and state

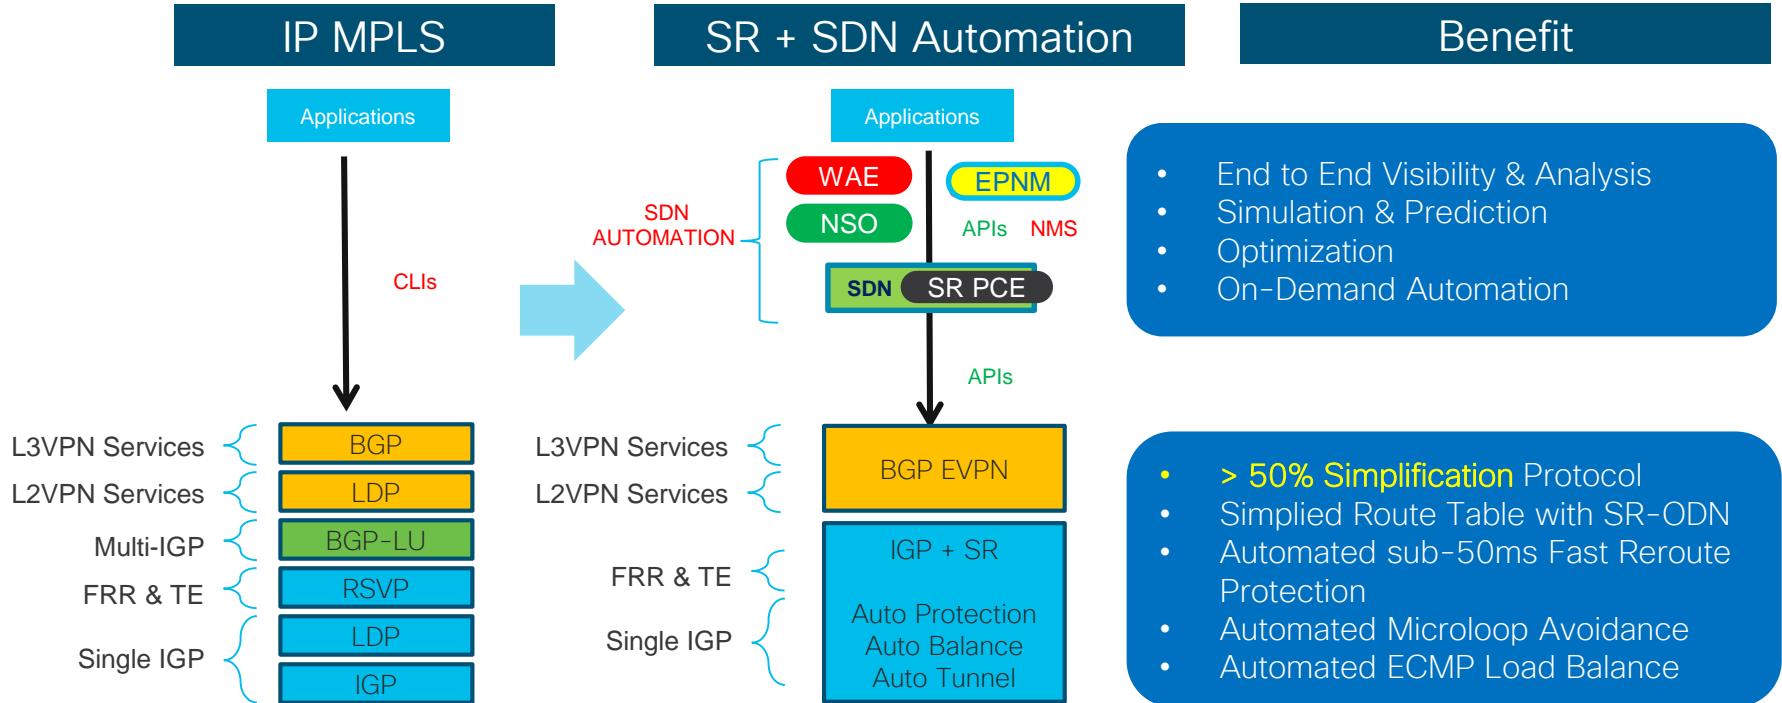
## Analytic Telemetry

- WAE enables Proactive Analytic, Simulation and Predictive Network Modeling
- Rich streaming telemetry available for all protocols using standards-based YANG models
- Unlock proactive fault detection and remediation

## Validated

- Network design and its component features tested individually and as a system
- design is continually evaluated as a component of end to end network

# IP MPLS Evolution to Segment Routing + SDN Automation



# Segment Routing: Key capabilities

Foundation for Network Operations Simplification and Automation,  
and Service Agility via Programmability

Simplification of network protocols  
Improved scalability ( 3 to ONLY 1)

Automated 50ms convergence

Simplification of Traffic Engineering

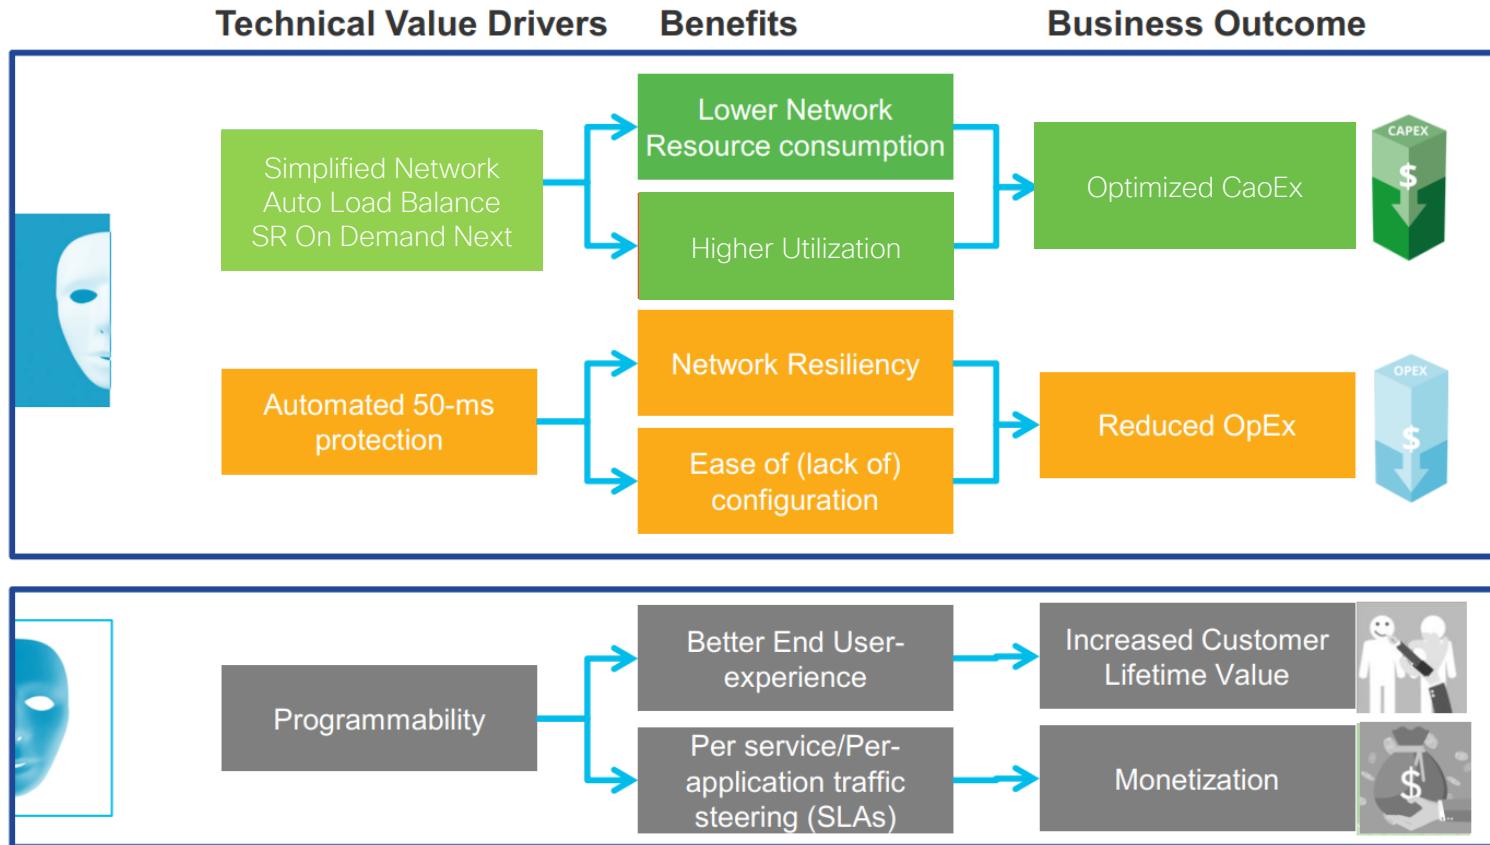
Built-in Redundancy & HA

Application enabled policy using SDN  
techniques

Support MPLS & v6 Forwarding

Universal Forwarding Plane  
From Access to DC

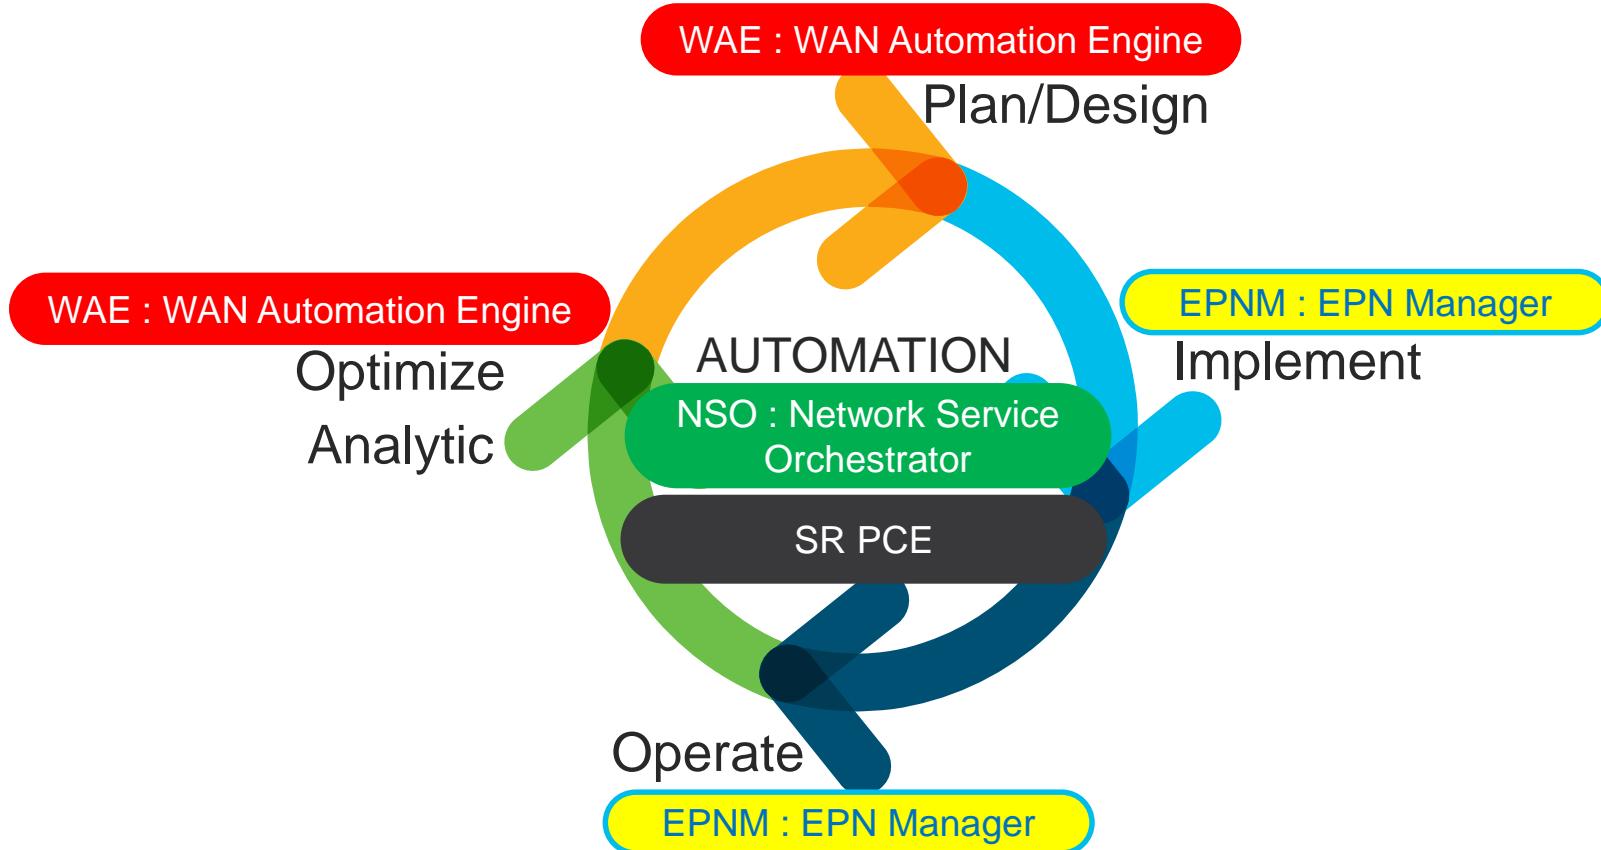
# Segment Routing Value Proposition



# Segment Routing vs RSVP-TE vs LDP

Use Case	SR	TE	LDP
Operational Simplicity	✓	✓	✓
Fast Reroute	✓	✓	✓
ECMP	✓	✗	✓
Traffic Engineering	✓	✓	✗
Multi Domain TE	✓	✓	✗
Network Slicing	✓	✓	✗
On-Demand-Next Hop	✓	✗	✗
Intent Based Automated Steering	✓	✗	✗

# Automation in Network Lifecycle



# Automation in Network Lifecycle

EPNM

Evolved Programmable  
Network Manager



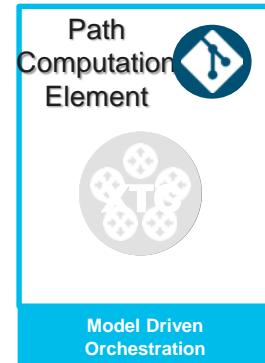
WAE

WAN Automation  
Engine



SR PCE

XR Transport  
Controller



NSO

Network Service  
Orchestrator



"Operate" service,  
network & device  
resources as the single  
point of truth for all  
operations

Traffic planning,  
design &  
optimization within  
the network

Real Time Path  
Compute and  
Topology

Implement Agile  
Service Orchestration  
for Across WAN, DC,  
NFVI

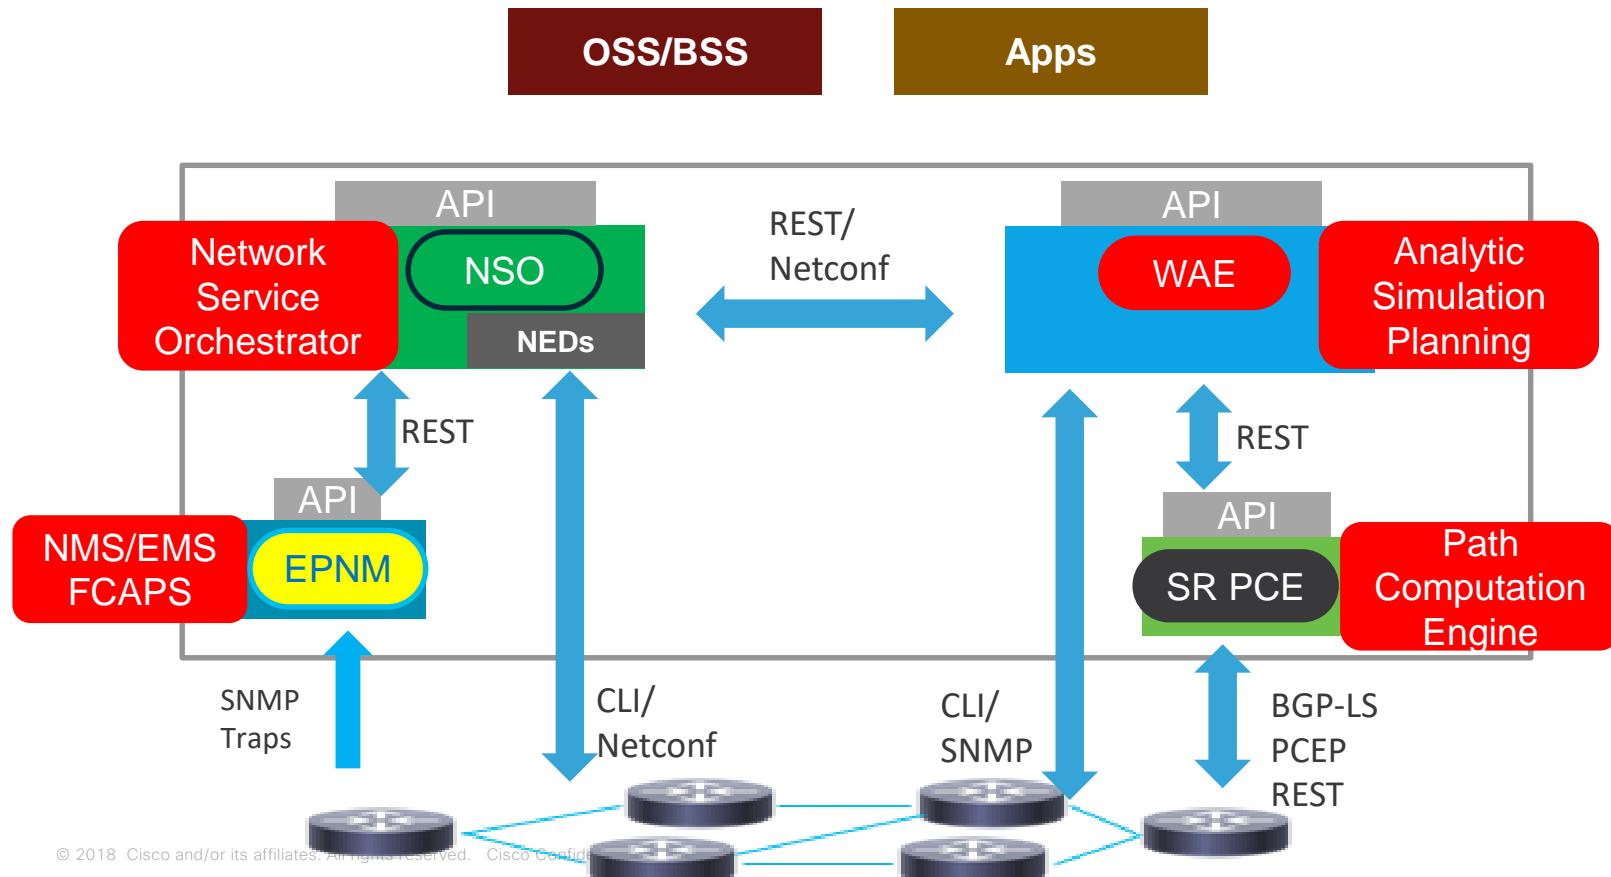
NMS

ANALYTIC  
SIMULATION

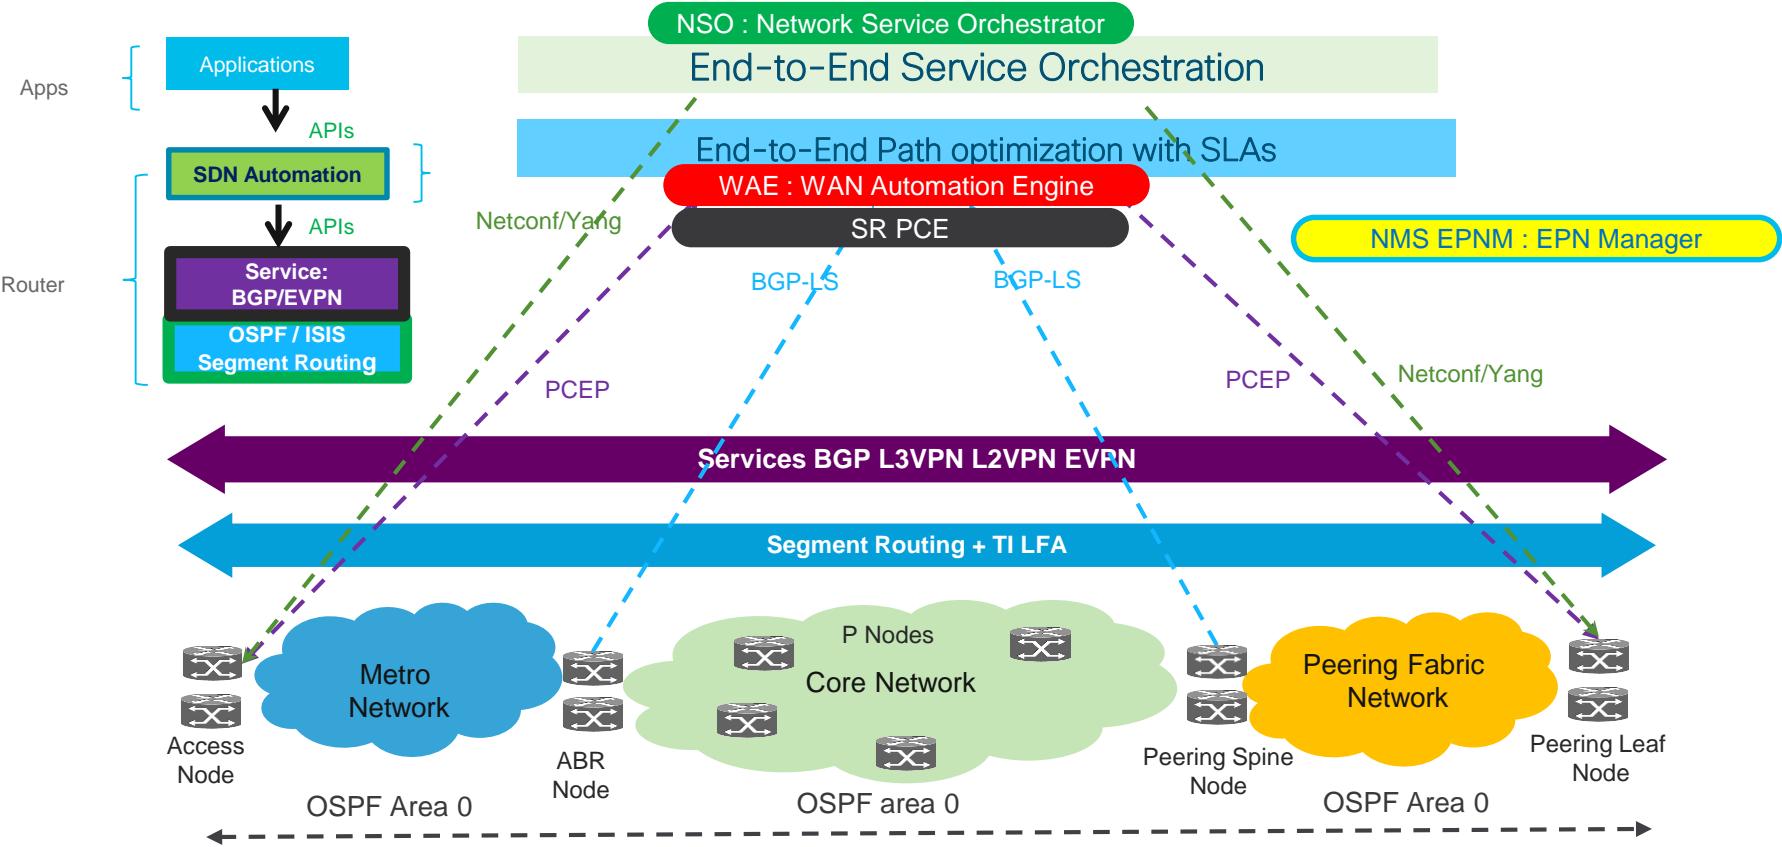
SDN  
CONTROLLER

ORCHESTRATOR

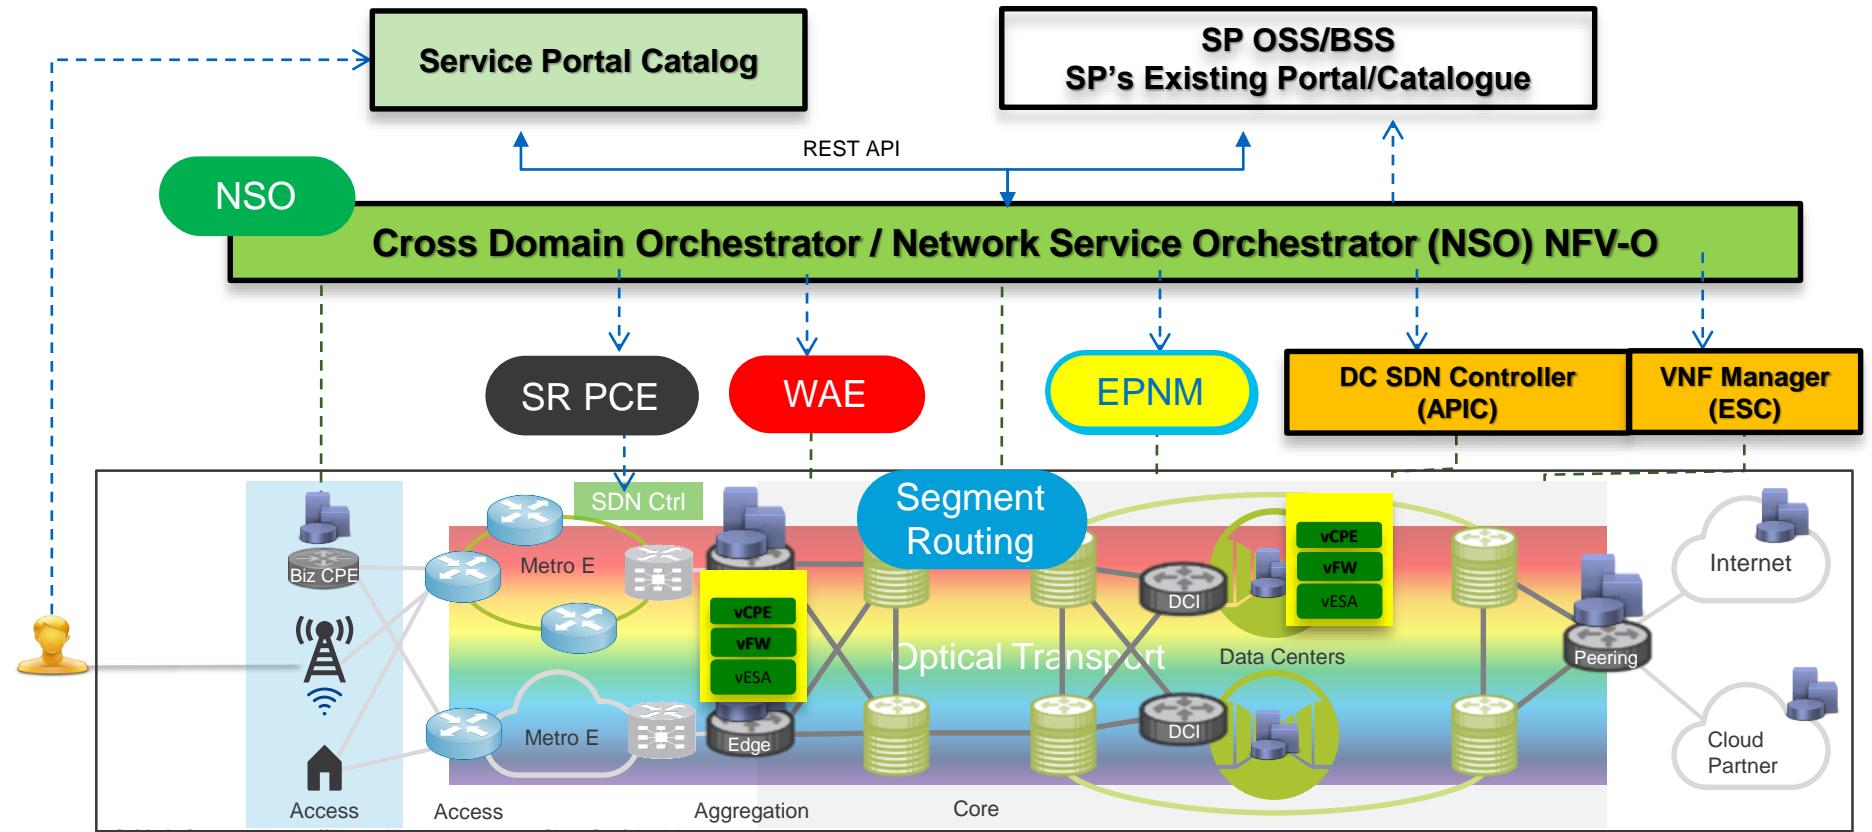
# Integration Architecture



# Segment Routing + SDN Automation



# SDN Automation Framework



# Segment Routing – SPRING RFC 7855

## Clarence Filsfils@Cisco

Internet Engineering Task Force (IETF)  
Request for Comments: 7855  
Category: Informational  
ISSN: 2070-1721

S. Previdi, Ed.  
C. Filsfils, Ed.  
Cisco Systems, Inc.

S. Litkowski  
Orange  
M. Horneffer  
Deutsche Telekom  
R. Shakir  
Jive Communications, Inc.  
May 2016

## Source Packet Routing in Networking (SPRING) Problem Statement and Requirements

### Abstract

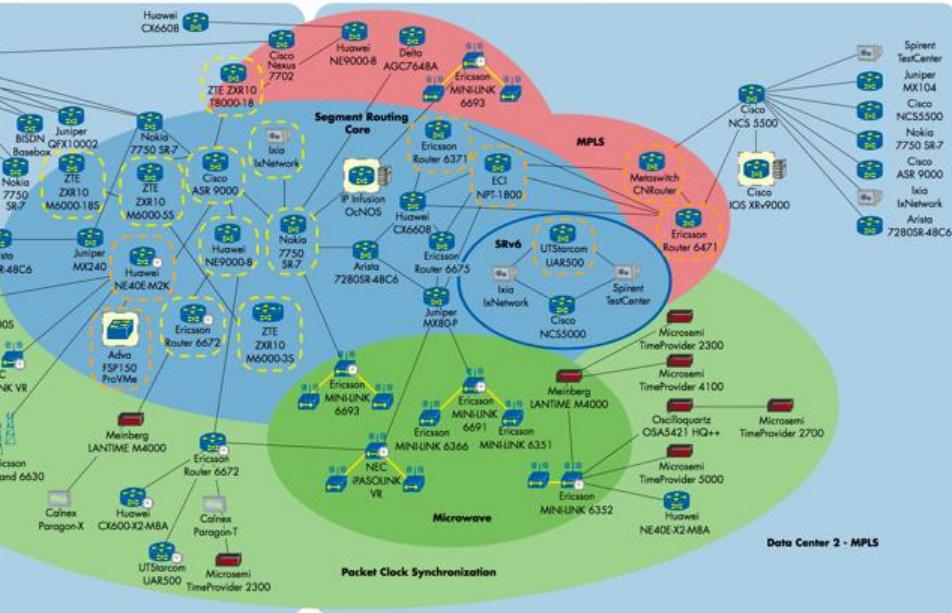
The ability for a node to specify a forwarding path, other than the normal shortest path, that a particular packet will traverse, benefits a number of network functions. Source-based routing mechanisms have previously been specified for network protocols but have not seen widespread adoption. In this context, the term "source" means "the point at which the explicit route is imposed"; therefore, it is not limited to the originator of the packet (i.e., the node imposing the explicit route may be the ingress node of an operator's network).

# Segment Routing + SDN Interop Showcase

- Segment Routing
- Segment Routing Traffic Engineering
- SDN Controller XTC
- Topology Independent LFA
- Segment Routing and LDP Interworking
- BGP Prefix Independent Convergence (PIC)
- Stateful PCE Path Computation Element
- Network Provisioning with NETCONF



# EANTC MPLS + SDN + NFV World Congress Public Multi-Vendor Interoperability Test 2018 Scope



Segment Routing Test Cases ✓

Ethernet VPN Test Cases

SDN Test Cases ✓

Microwave Test Cases

Clock Synchronization Test Cases

# www.segment-routing.net for tutorial & demos

→ ⓘ Not secure | [www.segment-routing.net/tutorials/](http://www.segment-routing.net/tutorials/)

ⓘ Not secure | [www.segment-routing.net/demos/](http://www.segment-routing.net/demos/)

The screenshot shows the website's navigation bar with links for SEGMENT ROUTING, NEWS, DEMOS, and TUTORIALS. Below the navigation, there are four main sections: 'Tutorials', 'SR concepts', 'SR IGP', 'SR and LDP', and 'SR Traffic Engineering'. Each section has a brief description and a video thumbnail. The 'TUTORIALS' section is currently selected.

- Tutorials**
  - > **SR concepts**  
Learn about the foundational SR concepts.
  - > **SR IGP**  
The SR IGP control plane and automated fast-reroute resiliency.
  - > **SR and LDP**  
Introducing SR in an LDP network, how they can co-exist and interwork.
  - > **SR Traffic Engineering**  
SR-TE provides you low-delay paths, disjoint paths,... and many more.
- SR On-Demand Next Hops (ODN)**  
Watch Jose Liste's demonstration of SR ODN. Learn how SR maps traffic from edge to backbone at scale.  
*José Liste, Cisco – November 08, 2016*
- Topology Independent - Loop Free Alternate (TI-LFA)**  
Watch Kris Michielsen's demonstration of TI-LFA. TI-LFA protects against topology changes by finding an optimum backup path, against link, node or topology changes.  
*Kris Michielsen, Cisco – November 07, 2016*
- Microloop Avoidance**  
Watch Kris Michielsen's demonstration of Microloop Avoidance. Microloops are a day-1 problem of IP networks.  
*Kris Michielsen, Cisco – November 06, 2016*

# <https://xrdocs.github.io/design/>



The screenshot shows the homepage of the xrdocs.github.io/design/ website. At the top, there is a navigation bar with links for @XRDOCS, DESIGN, BLOGS, TUTORIALS, SEARCH Q, TAGS, and ARCHIVE. On the left side, there is a section titled "Network Designs" with a subtext: "Service Provider deployments have consistently evolved for the last few decades and keeping abreast of the changes can get hard. Let the best in the industry walk you through Network Designs and Architectures that will lay the foundation for the next-gen SP networks". Below this text are two buttons: "Get Started!" and "Github". The main visual element is a large, stylized compass rose graphic on a blue background, with a world map visible behind it.

## Metro Fabric Design

Read about the goals of the Metro Fabric Design and familiarize yourself with all its components: the platforms, the software, the technologies and use cases

[High Level Design](#)

[Implementation Guide](#)



# EPN 5.0 Design Guide

<https://www.cisco.com/c/en/us/solutions/enterprise/design-zone-service-provider/programmable-network.html#~tab-epn-5guides>

← → C <https://www.cisco.com/c/en/us/solutions/enterprise/design-zone-service-provider/programmable-network.html#~tab-epn-5guides> ☆

## EPN 5 Validated Design Guides

[EPN 5.0 Transport Design Guide \(PDF - 5.6 MB\)](#)

[EPN 5.0 Services Design Guide \(PDF - 2.7 MB\)](#)

[EPN 5.0 System Test Topology Reference Guide \(PDF - 2 MB\)](#)

[EPN 5.0 Network Service Orchestration User Guide \(PDF - 2.3 MB\)](#)

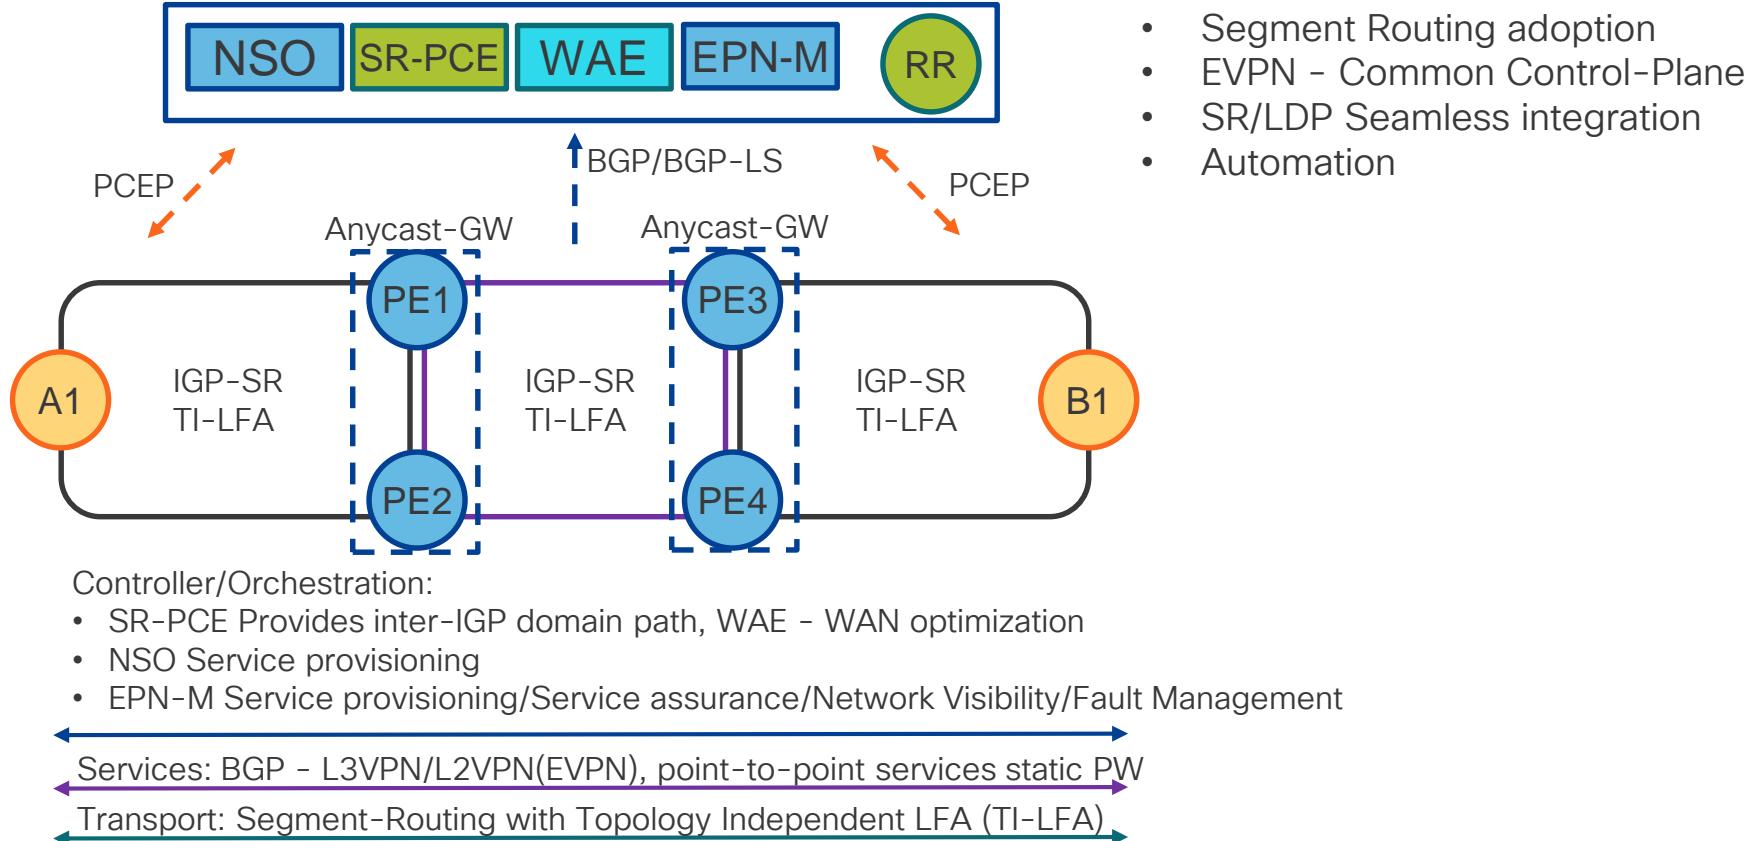
[EPN 5.0 Implementation Guide for Large Network with End to End Segment Routing \(PDF - 2.1 MB\)](#)

[EPN 5.0 Implementation Guide for Small Network with End to End Segment Routing \(PDF - 2 MB\)](#)

[EPN 5.0 Implementation Guide for Large Network with End to End Programmable Segment Routing \(PDF - 2 MB\)](#)

[EPN 5.0 Implementation Guide for Large Network with Segment Routing and LDP Interworking \(PDF - 2.1 MB\)](#)

# EPN 5.0 Framework

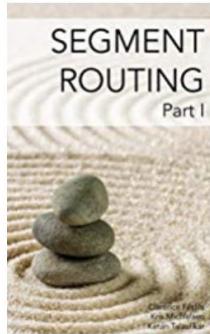


# Segment Routing Recommendation Books



[https://www.amazon.com/gp/product/B07N6H54VP/ref=series\\_rw\\_dp\\_sw](https://www.amazon.com/gp/product/B07N6H54VP/ref=series_rw_dp_sw)

1



[Segment Routing, Part I \(Oct 6, 2016\)](#)

by [Clarence Filsfils, Kris Michelsen, Ketan Talaulikar](#)

(6)

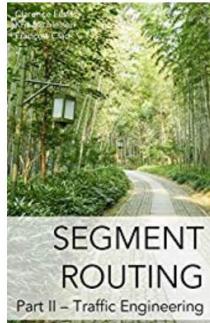
\$46.80

Sold by: Amazon Digital Services LLC

Over the past ten years, operators have heavily invested in network infrastructure because of exploding growth in IP traffic. Moreover, the lack of close interaction

[▼ Read More](#)

2



[Segment Routing Part II: Traffic Engineering \(Apr 22, 2019\)](#)

by [Clarence Filsfils, Kris Michelsen, Francois Clad](#)

\$46.80

Sold by: Amazon Digital Services LLC

As more and more Service Providers and Enterprises operate a single network infrastructure to support an ever-increasing number of services, the ability to custom

[▼ Read More](#)



Segment Routing (SR) - TI-LFA

SR Mapping Server

Use Case P-PE (SR + LDP/RSPV-TE)

SR Migration

Use-case P (SR) + PE (SR)

# SR References

[www.segment-routing.net/tutorials/](http://www.segment-routing.net/tutorials/)

SR Concepts >

- SR Introduction

- SR MPLS data plane

- SR Global Block (SRGB)

SR IGP >

- SR IGP Control Plane

- Topology Independent LFA (TI-LFA)

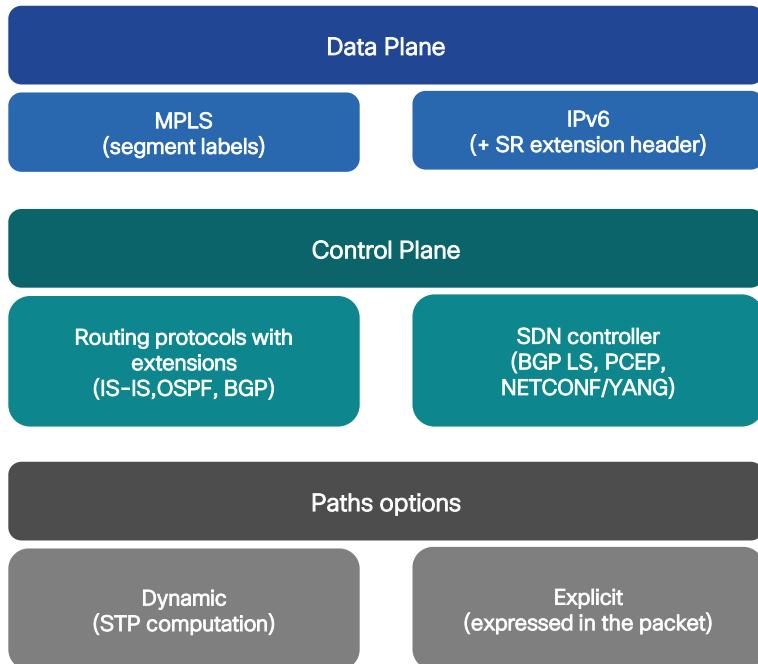
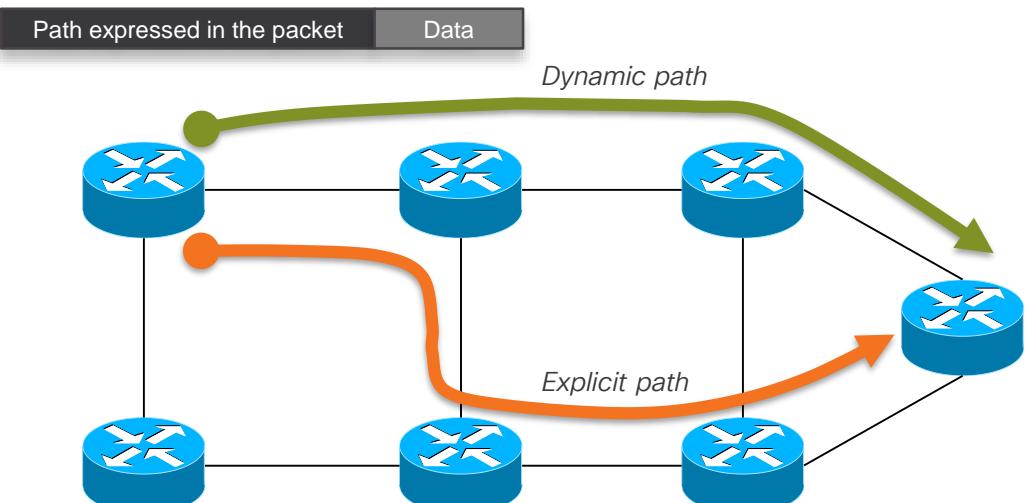
SR and LDP >

- SR and LDP co-existence

- SR Mapping Server

- SR and LDP interworking

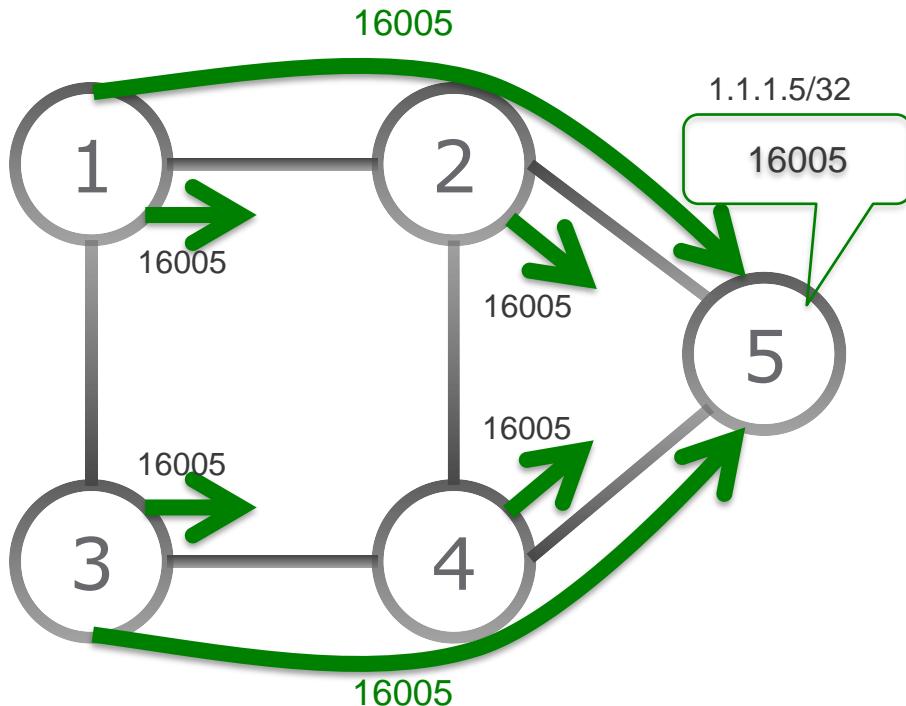
# Segment Routing - Overview



# IGP Prefix Segment

- Signaled by ISIS/OSPF
  - Minor extensions to existing link-state routing protocols
- Shortest-path to IGP prefix
  - Equal Cost MultiPath (ECMP)-aware
- Global significance in SR domain
- Label = SRGB + Index
  - SRGB = Segment Routing Global Block
  - Default SRGB: 16,000 – 23,999
  - Advertised as index

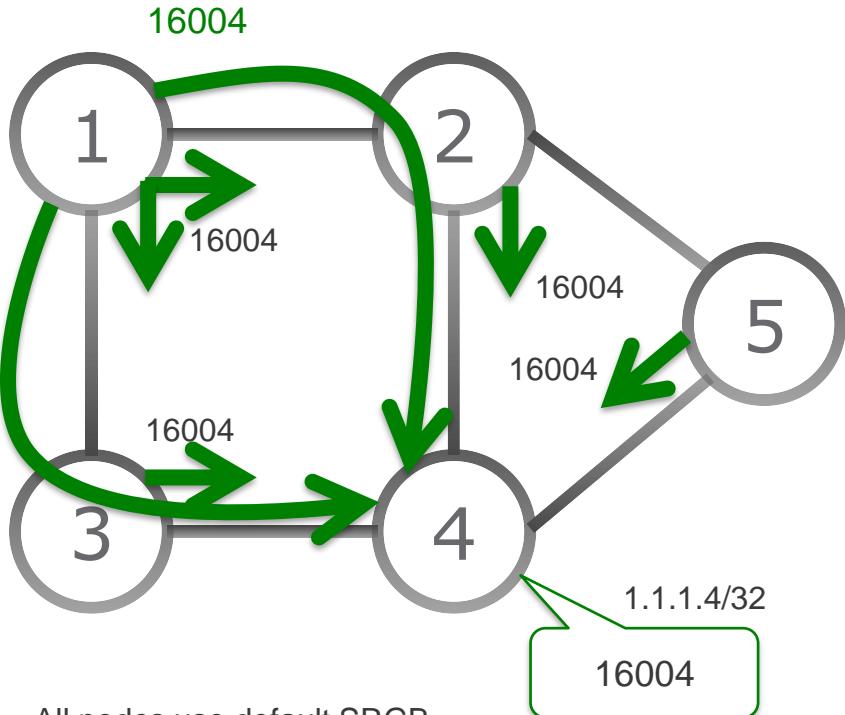
Example 1: Best Path



# IGP Prefix Segment

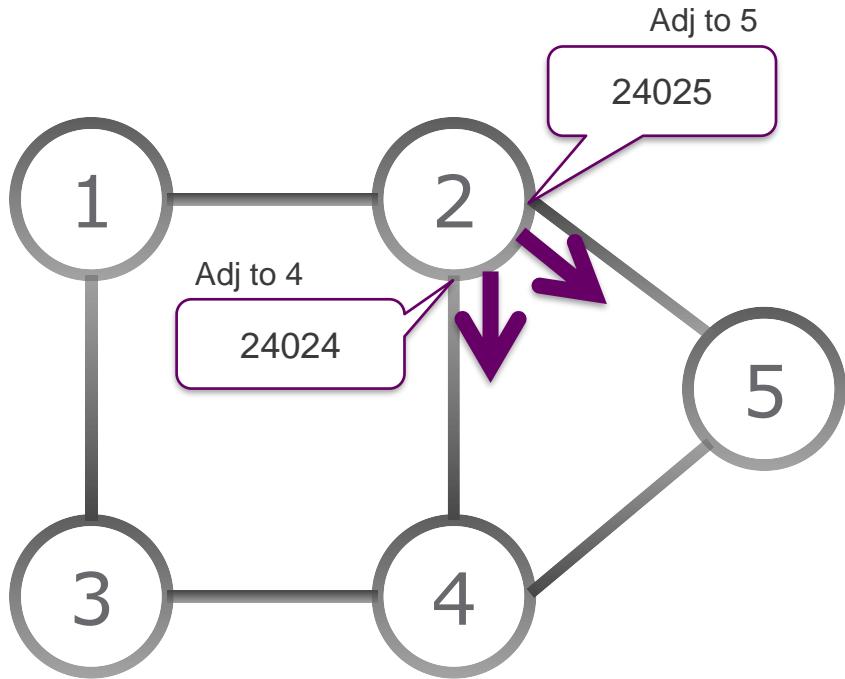
- Signaled by ISIS/OSPF
  - Minor extensions to existing link-state routing protocols
- Shortest-path to IGP prefix
  - Equal Cost MultiPath (ECMP)-aware
- Global significance in SR domain
- Label = SRGB + Index
  - SRGB = Segment Routing Global Block
  - Default SRGB: 16,000 – 23,999
  - Advertised as index

## Example 2: ECMP



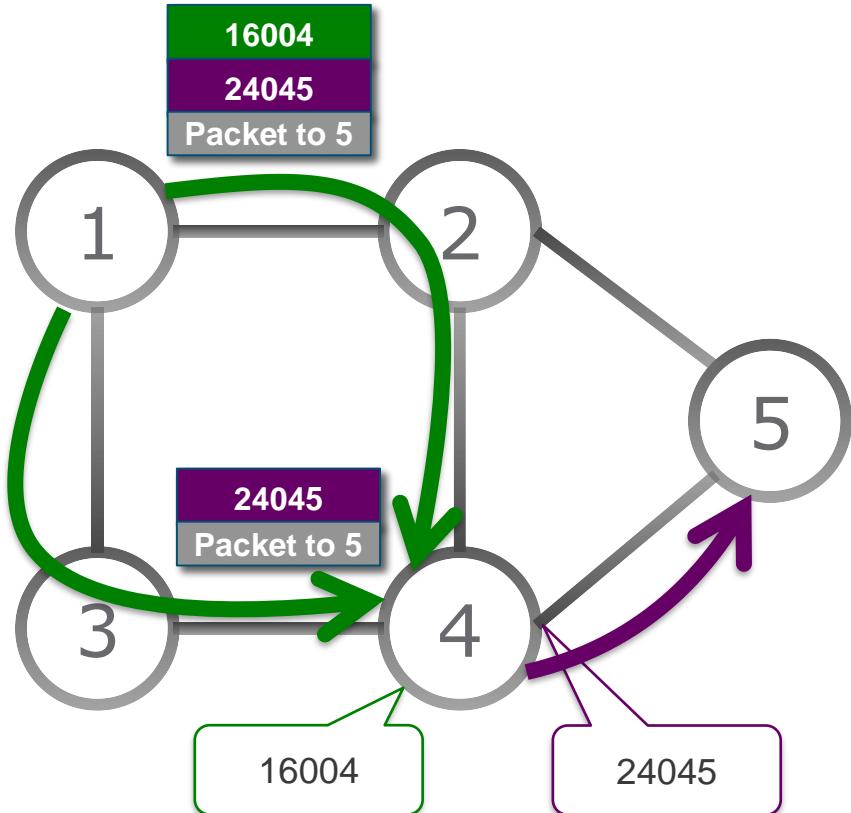
# IGP Adjacency Segment

- Signaled by ISIS/OSPF
  - Minor extensions to existing link-state routing protocols
- Forward on IGP adjacency
- Local significance
- Automatically allocated by router

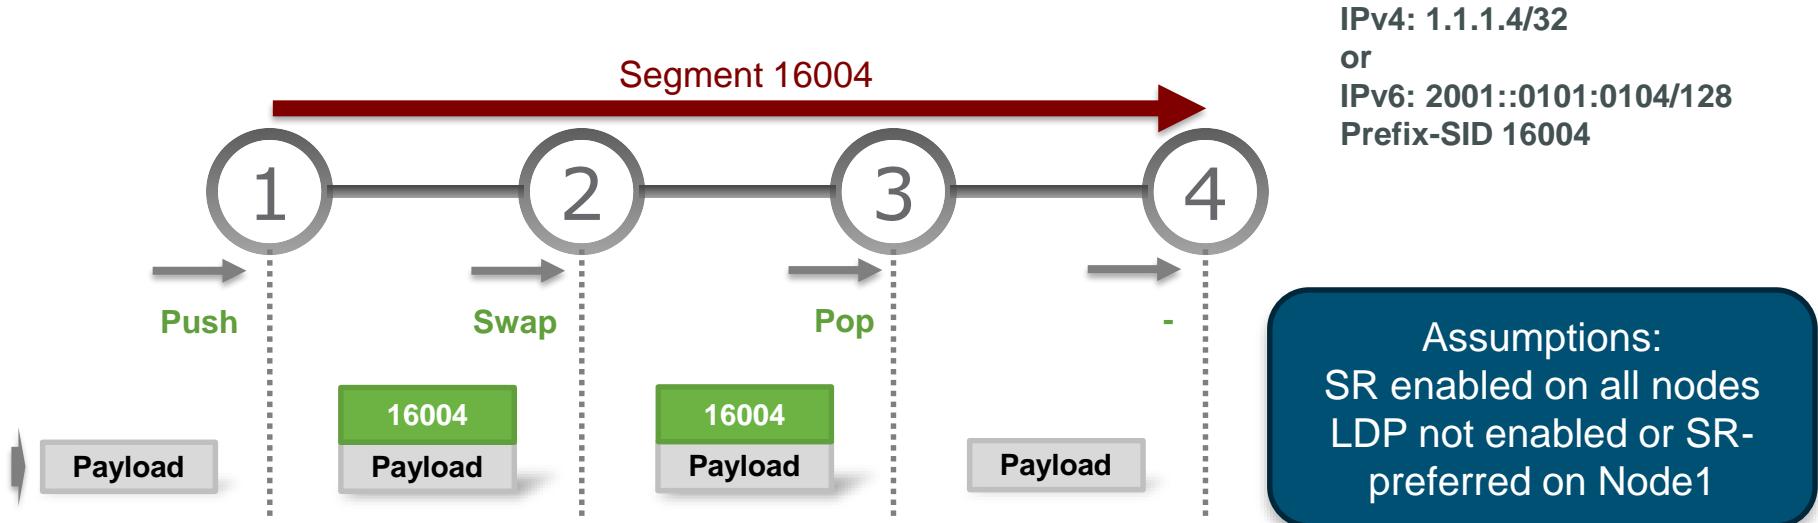


# Combining IGP Segments

- Signaled by ISIS/OSPF
- Steer traffic on any path through the network
- Path is specified by list of segments in packet header, a stack of labels
- No path is signaled
- No per-flow state is created



# MPLS Data Plane Operations



- Node4 advertises its loopback v4 or v6 address with attached prefix-SID 16004
  - IPv4 address: 1.1.1.4/32
  - IPv6 prefix: 2001::0101:0104/128
- Node4 requests default PHP functionality

# SR Label Forwarding Verification

```
RP/0/0/CPU0:xrvr-3#show mpls forwarding
```

Local Label	Outgoing Label	Prefix or ID	Outgoing Interface	Next Hop	Bytes Switched
16001	16001	SR Pfx (idx 1)	Gi0/0/0/0	99.2.3.2	0
16002	Exp-Null-v4	SR Pfx (idx 2)	Gi0/0/0/0	99.2.3.2	0
16004	Pop	SR Pfx (idx 4)	Gi0/0/0/1	99.3.4.4	0
16005	16005	SR Pfx (idx 5)	Gi0/0/0/1	99.3.4.4	0
16010	16010	SR Pfx (idx 10)	Gi0/0/0/0	99.2.3.2	0
	16010	SR Pfx (idx 10)	Gi0/0/0/1	99.3.4.4	0
24032	Pop	SR Adj (idx 1)	Gi0/0/0/0	99.2.3.2	0
24034	Pop	SR Adj (idx 3)	Gi0/0/0/1	99.3.4.4	0

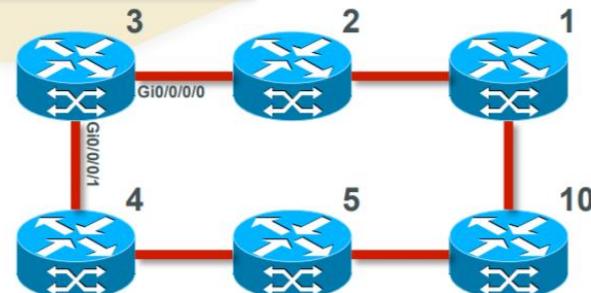


Remote prefix-SID

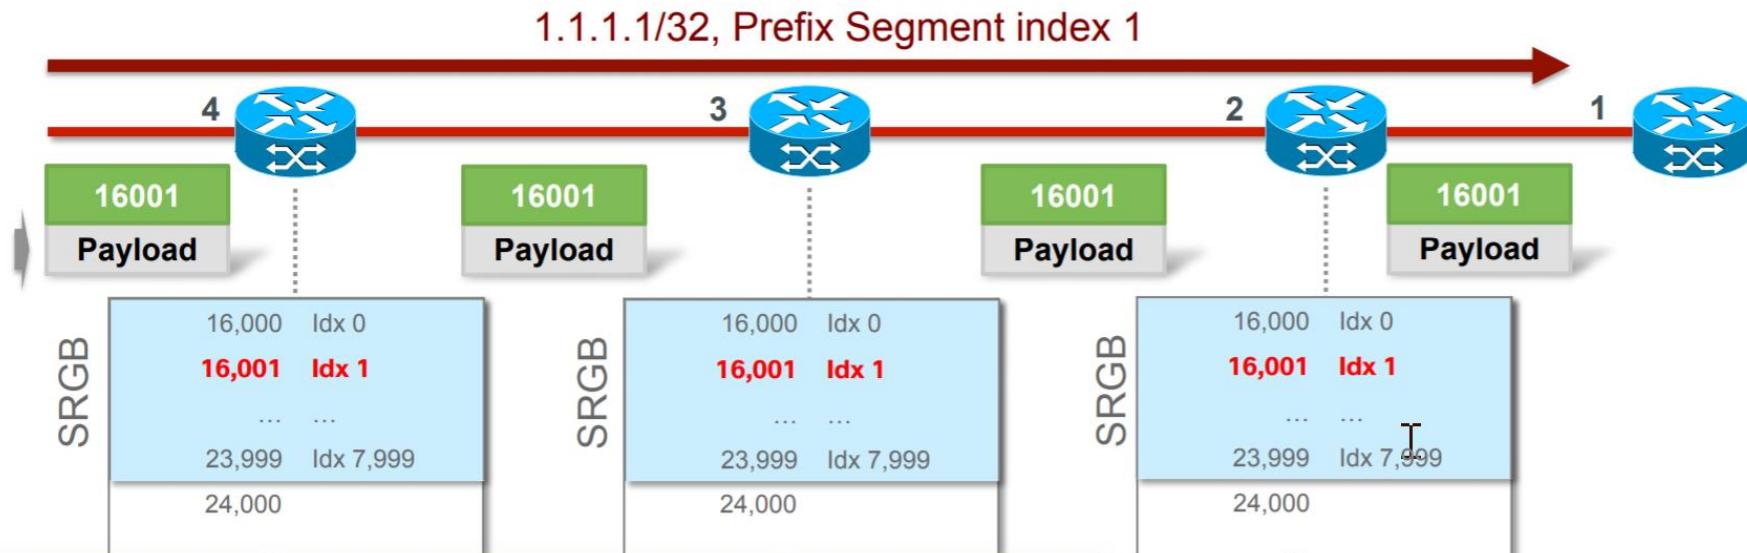
Neighbor prefix-SID  
Explicit-Null

Neighbor prefix-SID  
PHP on

Remote prefix-SIDs  
ECMP

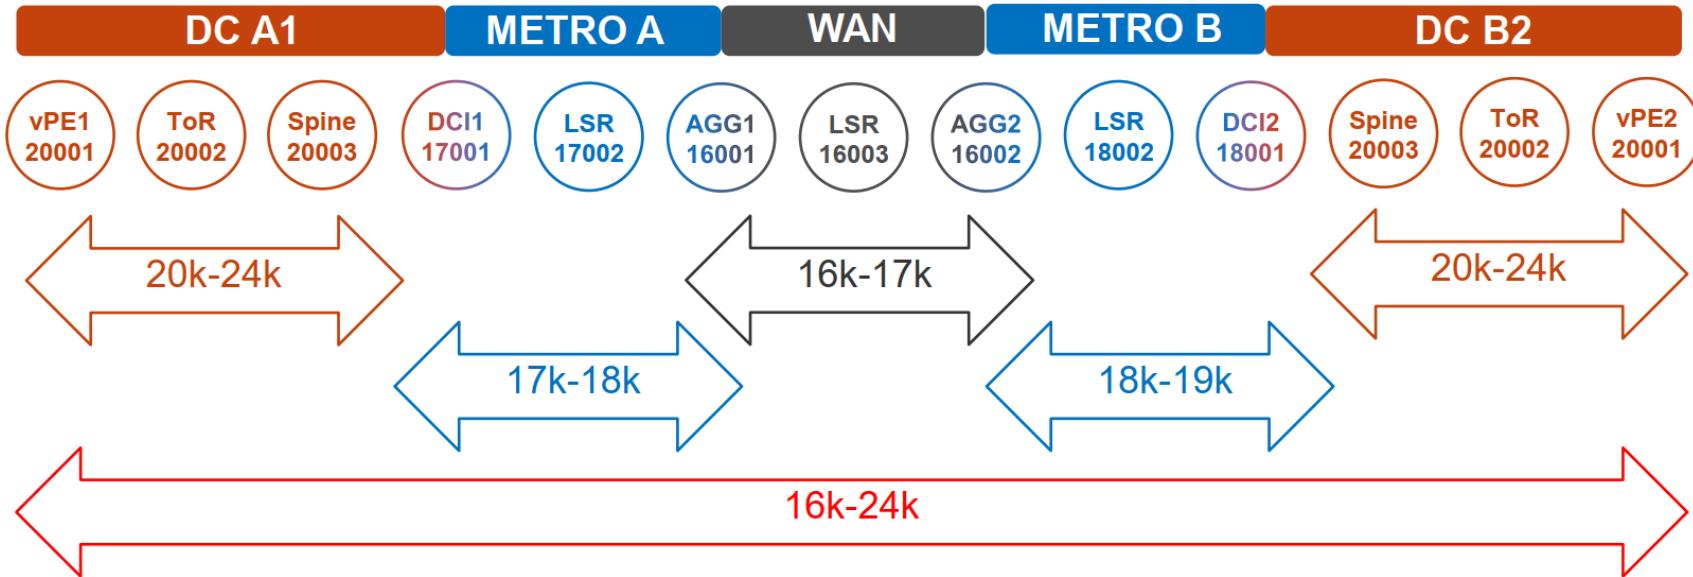


# Recommended SRGB



Same SRGB → prefix-SID has Global label value:  
Simple, predictable  
Much easier to troubleshoot  
Simplifies SDN programming

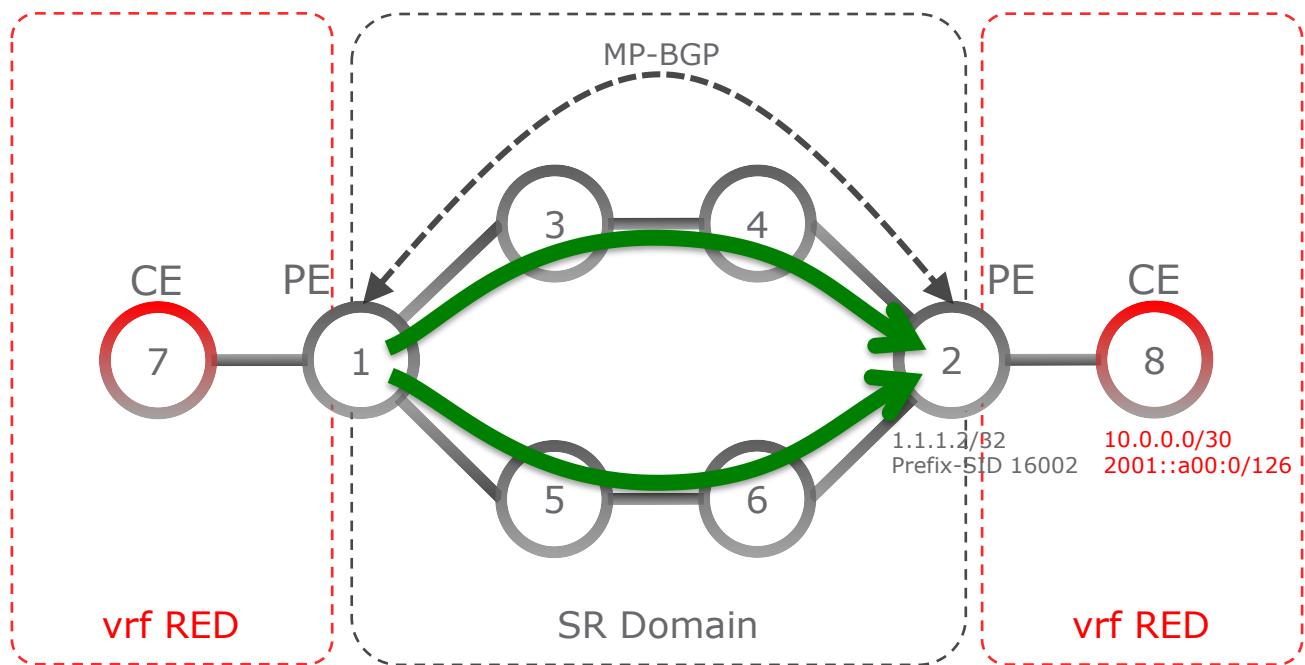
# SRGB and SID Allocation



- Homogenous end-to-end SRGB for simplicity
- Globally Unique Prefix SIDs for devices WAN and Metro domains
- Locally Unique Prefix SIDs for Datacenters

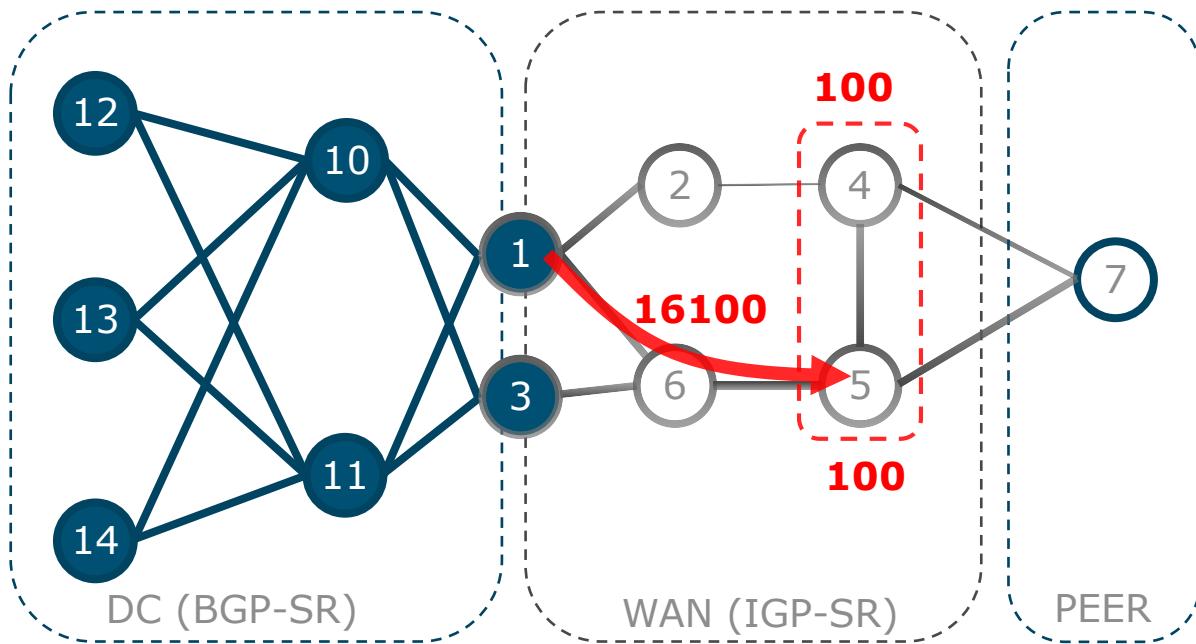
# Simple and Efficient Transport of MPLS services

- MPLS services ride on prefix segments
- Simple, one less protocol to operate (LDP)



# Anycast Prefix Segment ID (SID)

- Same prefix advertised by multiple nodes
- Traffic forwarded to one of Anycast prefix-SIDs based on best IGP path
- If primary node fails, traffic is auto re-routed to other node



# IGP with Segment Routing

```
router isis 1
  address-family ipv4 unicast
    metric-style wide
    segment-routing mpls
  !
  interface Loopback0
    address-family ipv4 unicast
      prefix-sid prefix-sid index 1
```

IGP-SR



TI-LFA



```
router ospf 1
  router-id 1.1.1.1
  segment-routing mpls
  segment-routing forwarding mpls
  area 0
  interface Loopback0
    passive enable
    prefix-sid index 1
```

# IGP with Segment Routing + TI-LFA

```
router isis 1
  address-family ipv4 unicast
    metric-style wide
    segment-routing mpls
  !
  interface Loopback0
    address-family ipv4 unicast
      prefix-sid prefix-sid index 1
```

IGP-SR



```
router ospf 1
  router-id 1.1.1.1
  segment-routing mpls
  segment-routing forwarding mpls
  area 0
  interface Loopback0
    passive enable
    prefix-sid index 1
```

```
router isis 1
  interface GigabitEthernet0/0/0/2
    address-family ipv4 unicast
      fast-reroute per-prefix
      fast-reroute per-prefix ti-lfa
```

TI-LFA



```
router ospf 1
  fast-reroute per-prefix
  fast-reroute per-prefix ti-lfa
```

SR/LDP Co-existence  
LDP->SR migration

# IGP with Segment Routing + TI-LFA + LDP->SR Migration

```
router isis 1
  address-family ipv4 unicast
    metric-style wide
    segment-routing mpls
  !
  interface Loopback0
    address-family ipv4 unicast
      prefix-sid prefix-sid index 1
```

```
router isis 1
  interface GigabitEthernet0/0/0/2
    address-family ipv4 unicast
      fast-reroute per-prefix
      fast-reroute per-prefix ti-lfa
```

```
router isis 1
  address-family ipv4 unicast
    segment-routing mpls sr-prefer
```

IGP-SR

TI-LFA

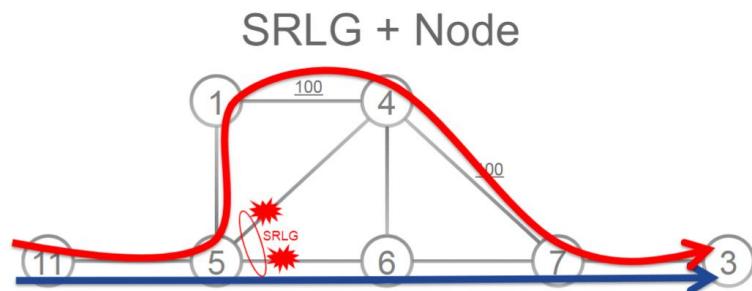
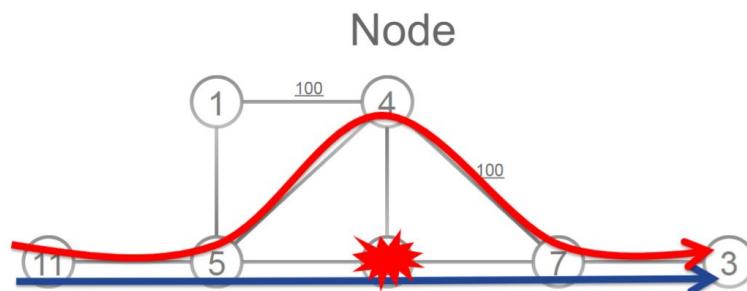
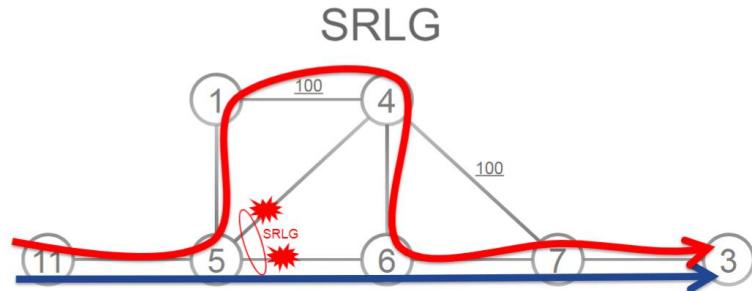
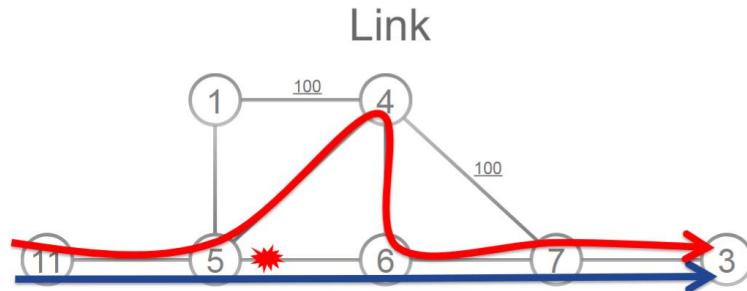
SR/LDP Co-existence  
LDP->SR migration

```
router ospf 1
  router-id 1.1.1.1
  segment-routing mpls
  segment-routing forwarding mpls
  area 0
  interface Loopback0
    passive enable
    prefix-sid index 1
```

```
router ospf 1
  fast-reroute per-prefix
  fast-reroute per-prefix ti-lfa
```

```
router ospf 1
  segment-routing mpls
  segment-routing sr-prefer
```

# TI-LFA Protecting Link, Node, SRLG

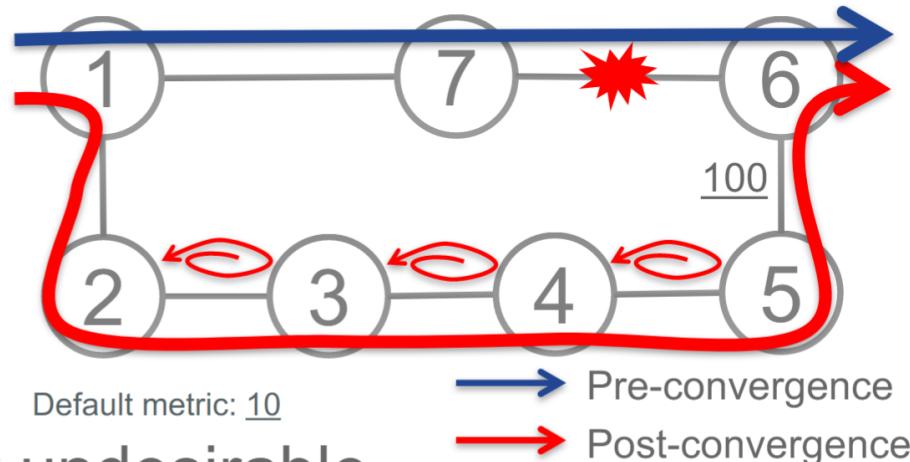


→ Pre-convergence  
→ TI-LFA/post-convergence



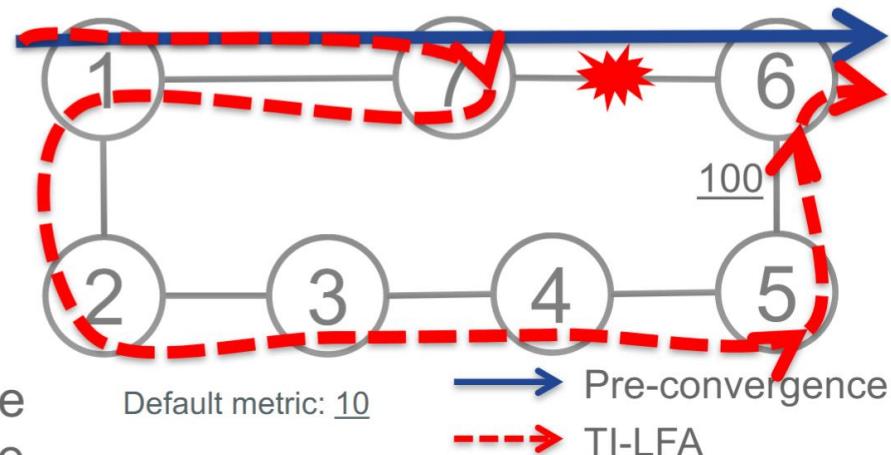
# What is MicroLoop ?

- Microloops are a day-one IP drawback
- IP hop-by-hop routing may induce microloop at any topology transition
  - Link up/down, metric up/down
- E.g. Microloops can occur after failure of link 6-7
- Microloops can increase packet loss, which is especially undesirable when FRR is used.



# SR MicroLoop Avoidance workflow

- TI-LFA protection kicks in on Node7, repairing the traffic to Node6 via Node5 and link Node5-Node6
- All nodes are notified of the topology change due to the failure
- E.g. Node1 computes the post-convergence SPT and detects possible microloops on the post-convergence paths for any destination, such as Node6
- If microloops are possible on the post-convergence path for a destination, then a SID-list is constructed to steer the traffic to that destination loop-free over the post-convergence path; in this example: <Prefix-SID(5), Adj-SID(5-6)> for destination Node6

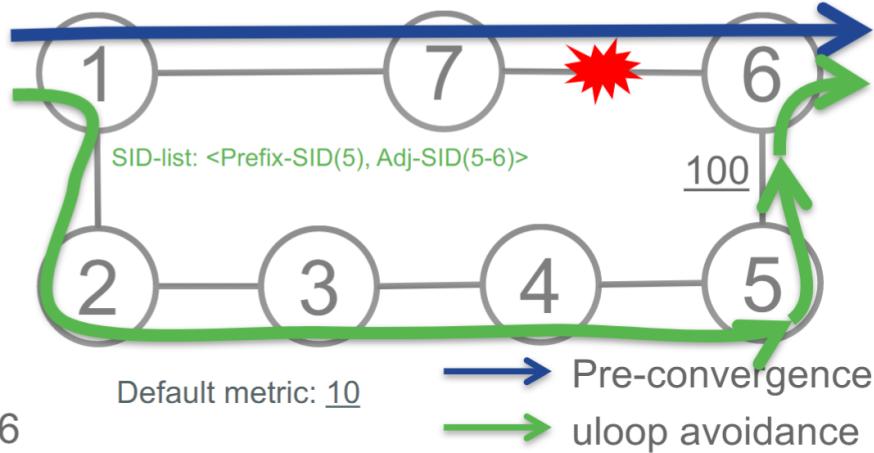


# SR Microloop Avoidance workflow

- IGP on Node1 updates the forwarding table and installs the SID-list imposition entries for those destinations with possible microloops, such as destination Node6

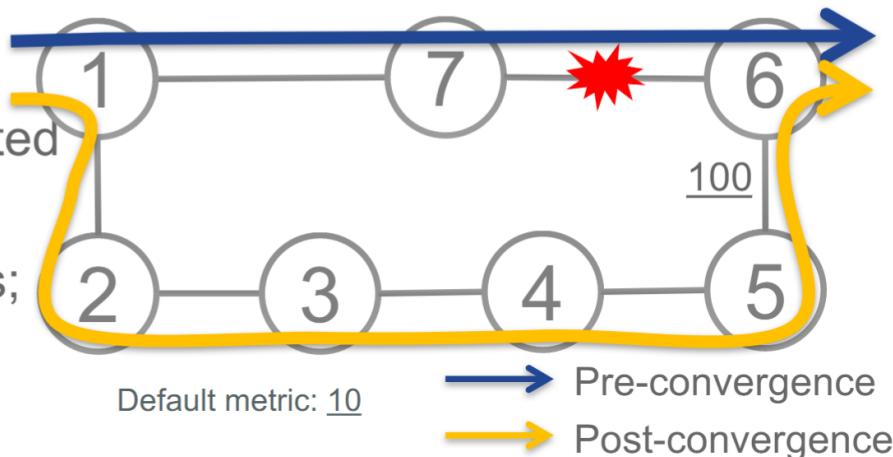
- Node1 imposes SID-list <Prefix-SID(5), Adj-SID(5-6)> on packets to Node6

- All nodes converge and update their forwarding tables, using SID-lists where needed



# SR Microloop Avoidance workflow

- Some time later, the new topology is applied and no more microloops are expected
- IGP updates the forwarding table, removing the microloop avoidance SID-lists; traffic now natively follows the post-convergence path
- Note: SR microloop avoidance is a local behavior, not all nodes need to implement it to get the benefits
  - E.g. if only Node1 has SR microloop avoidance, then e.g. traffic entering Node2 (not from Node1) to Node6 would still see microloops
  - When enabling SR microloop avoidance on Node2, then e.g. traffic entering Node3 (not from Node2) to Node6 would still see microloops, etc.
- There is incremental benefit for each node that has it implemented



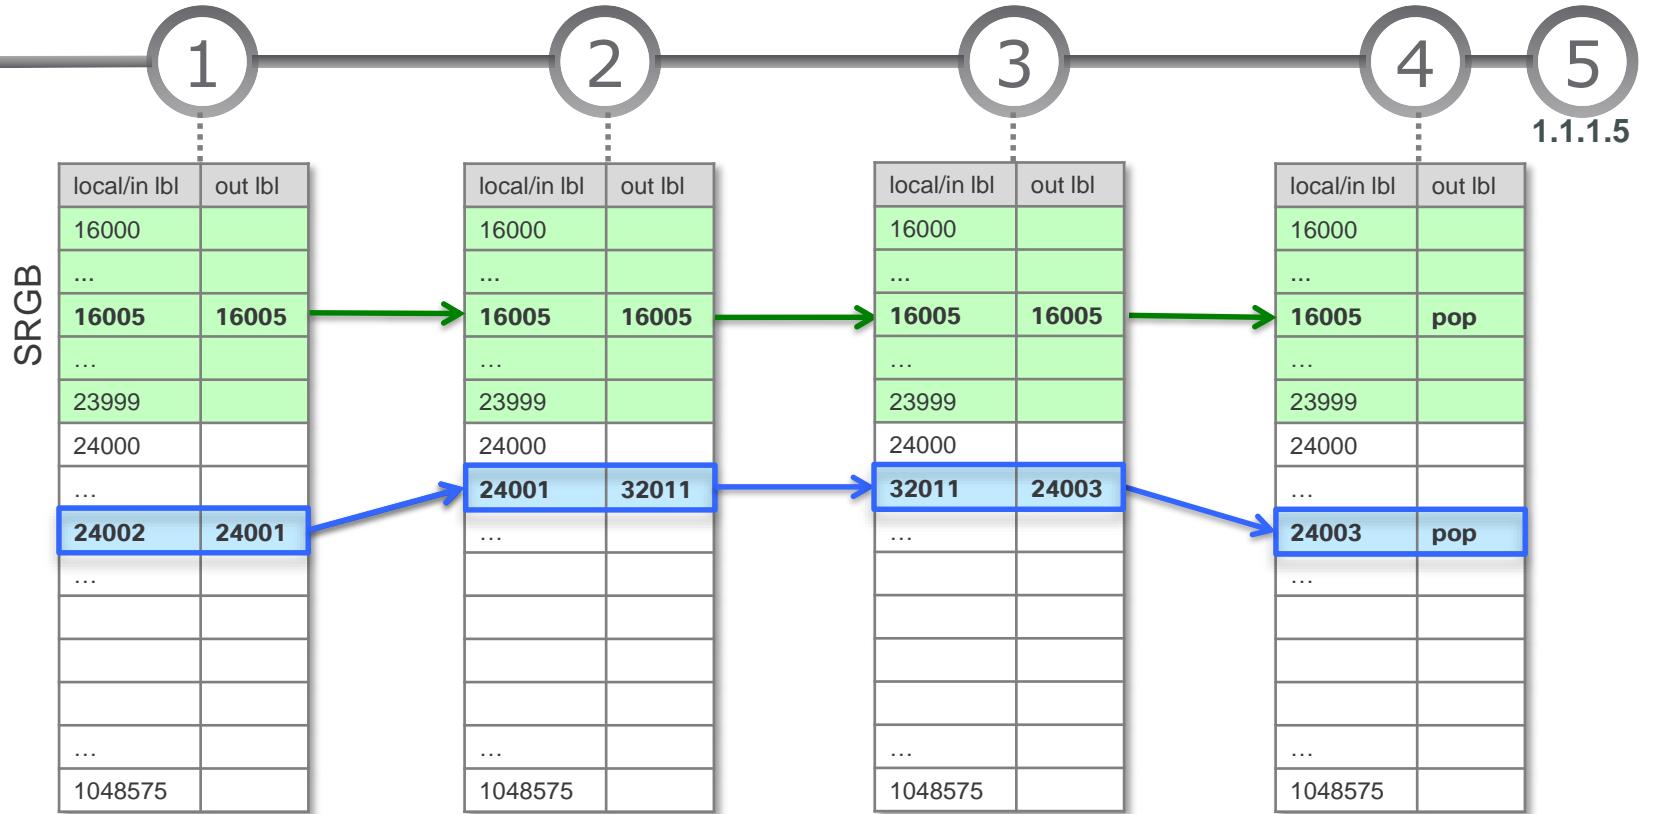
# MPLS-to-MPLS and MPLS-to-IP

All nodes: SR + LDP

SR

LDP

- LDP FEC to 1.1.1.5/32



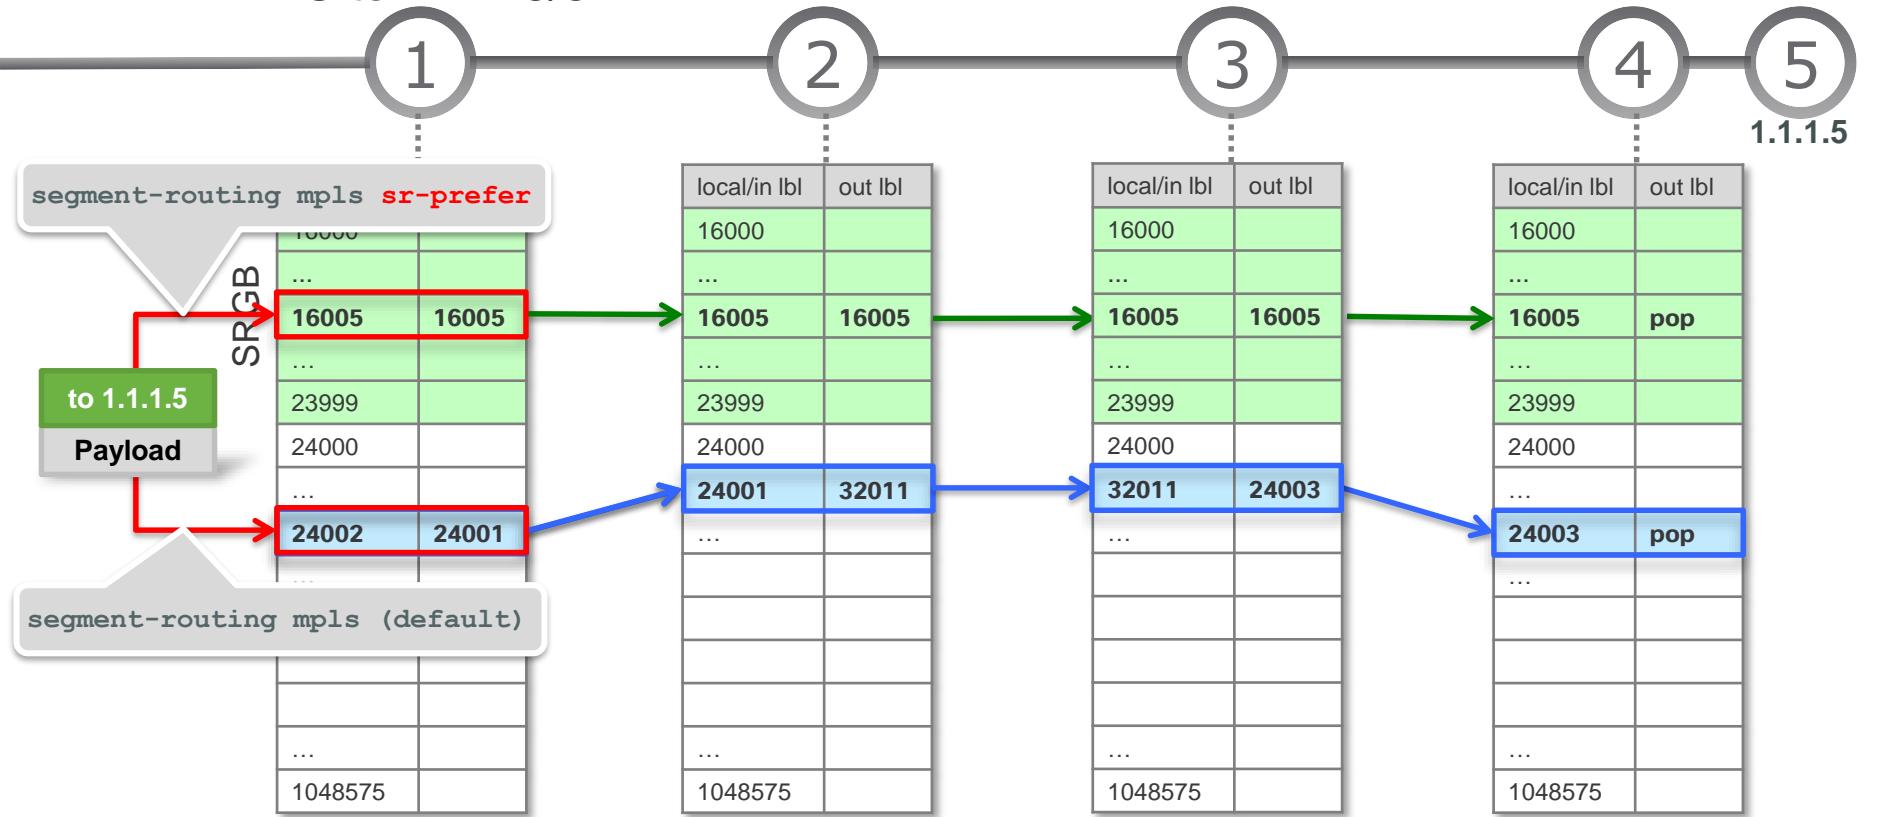
# MPLS-to-MPLS and MPLS-to-IP

All nodes: SR + LDP

SR

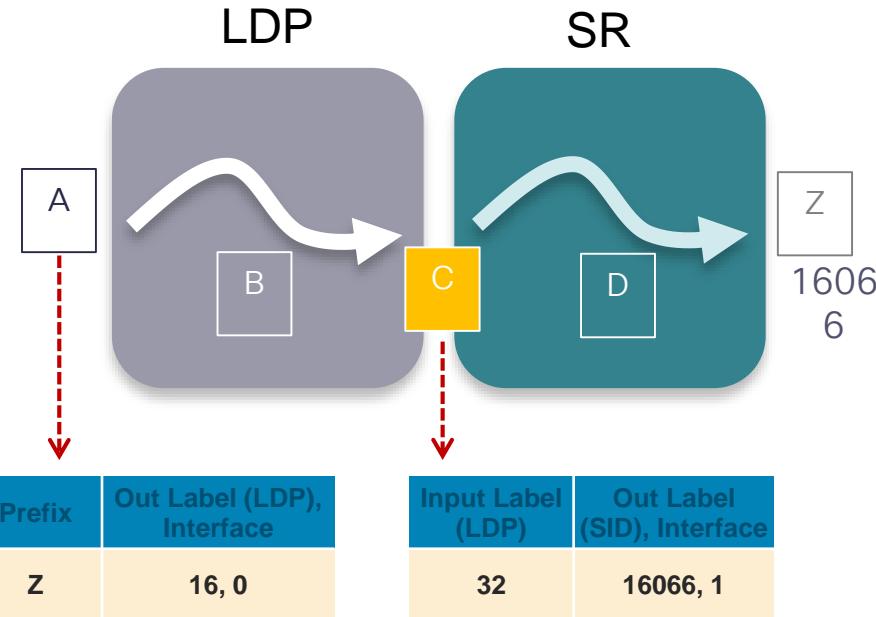
LDP

- LDP FEC to 1.1.1.5/32



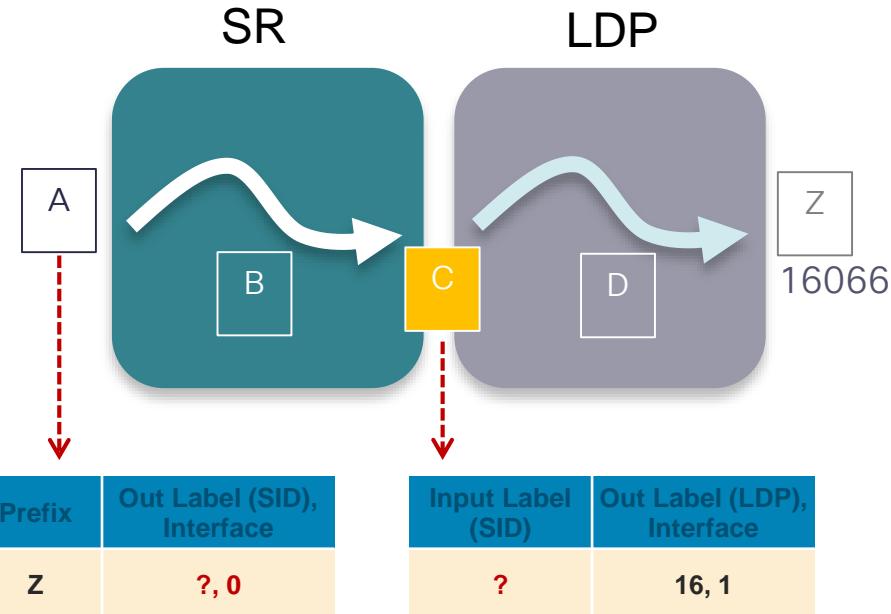
# LDP to SR

- When a node is LDP capable but its next-hop along the SPT to the destination is not LDP capable
  - no LDP outgoing label
- In this case, the LDP LSP is connected to the prefix segment
- C installs the following LDP-to-SR FIB entry:
  - incoming label: label bound by LDP for FEC Z
  - outgoing label: prefix segment bound to Z
  - outgoing interface: D
- This entry is derived automatically at the routing layer



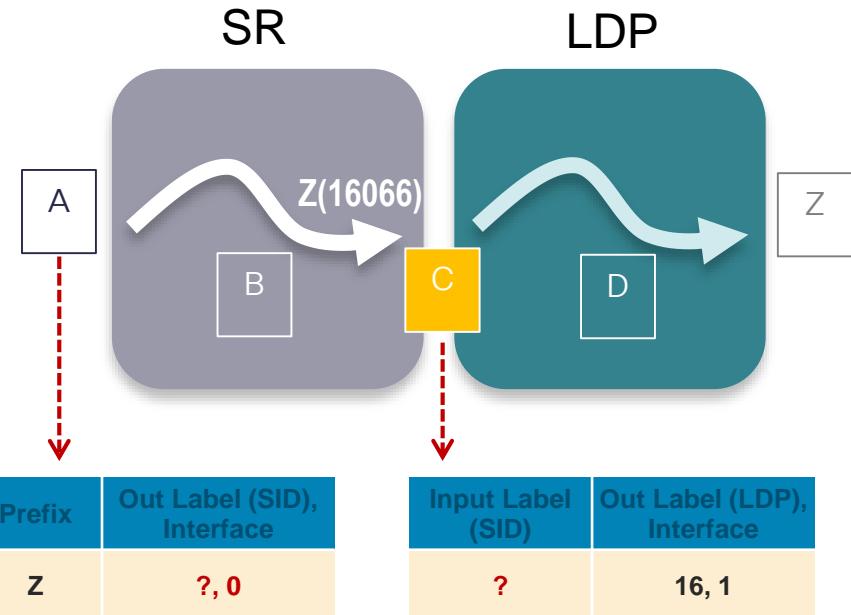
# SR to LDP

- When a node is SR capable but its next-hop along the SPT to the destination is not SR capable
  - no SR outgoing label available
- In this case, the prefix segment is connected to the LDP LSP
  - Any node on the SR/LDP border installs SR-to-LDP FIB entry(ies)



# Mapping Server

- A wants to send traffic to Z, but
  - Z is not SR-capable, Z does not advertise any prefix-SID  
→ which label does A have to use?
- The **Mapping Server** advertises the SID mappings for the non-SR routers
  - for example, it advertises that Z is 16068
- A and B install a normal SR prefix segment for 16066
- C realizes that its next hop along the SPT to Z is not SR capable hence C installs an SR-to-LDP FIB entry
  - incoming label: prefix-SID bound to Z (16066)
  - outgoing label: LDP binding from D for FEC Z
- A sends a frame to Z with a single label: 16066

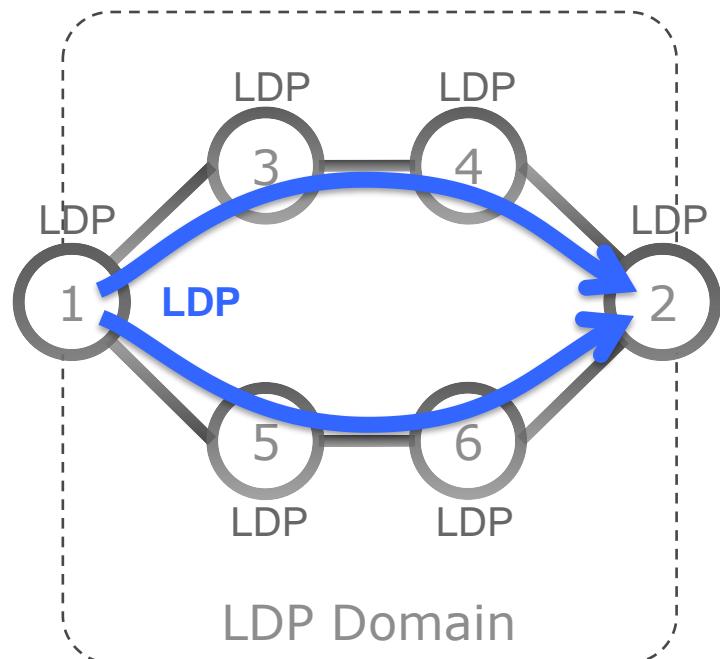


# Simplest migration LDP to SR

Assumptions:

- all the nodes can be upgraded to SR
- all the services can be upgraded to SR

- **Initial state:** All nodes run LDP, not SR

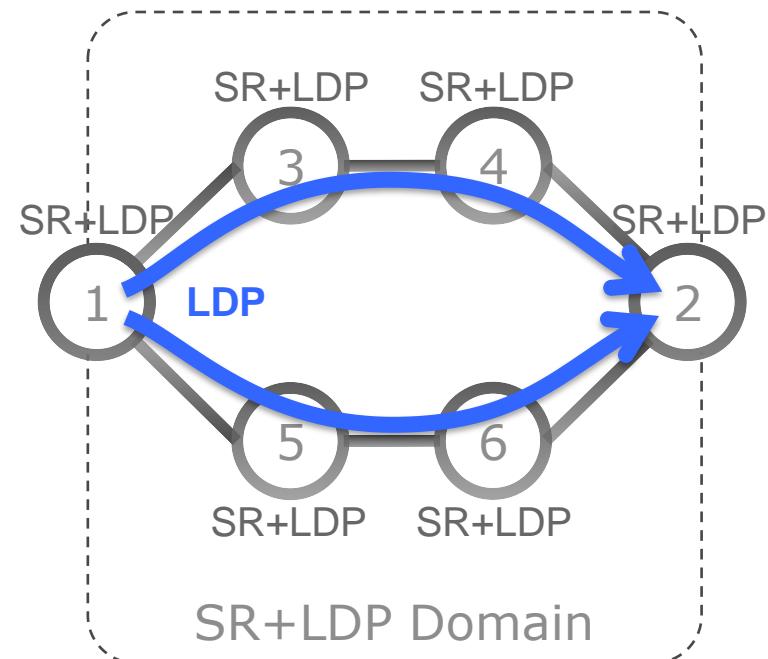


# Simplest migration LDP to SR

- **Initial state:** All nodes run LDP, not SR
- **Step1:** All nodes are upgraded to SR
  - In no particular order
  - leave default LDP label imposition preference

Assumptions:

- all the nodes can be upgraded to SR
- all the services can be upgraded to SR

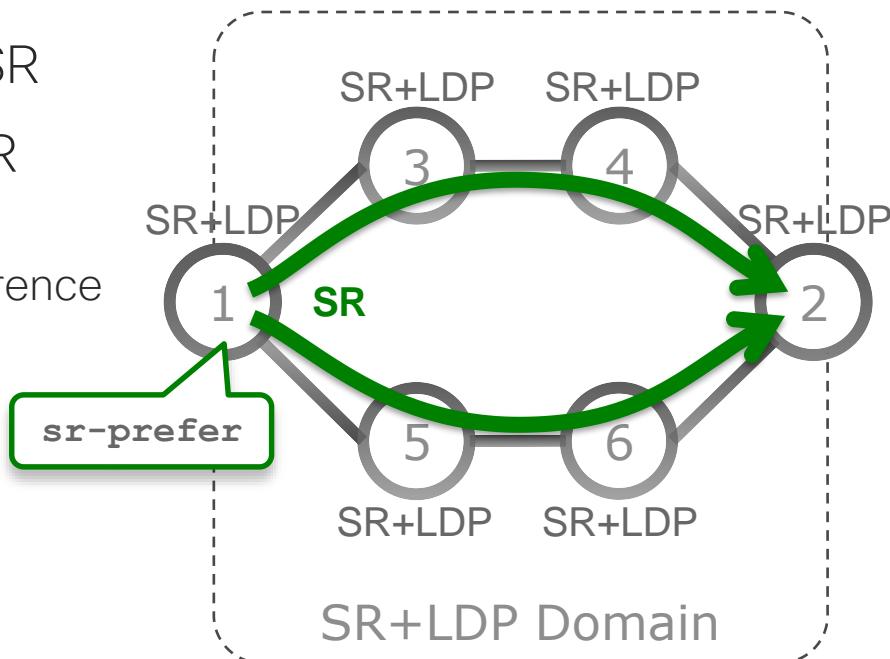


# Simplest migration LDP to SR

- **Initial state:** All nodes run LDP, not SR
- **Step1:** All nodes are upgraded to SR
  - In no particular order
  - leave default LDP label imposition preference
- **Step2:** All PEs are configured to prefer SR label imposition
  - In no particular order

Assumptions:

- all the nodes can be upgraded to SR
- all the services can be upgraded to SR

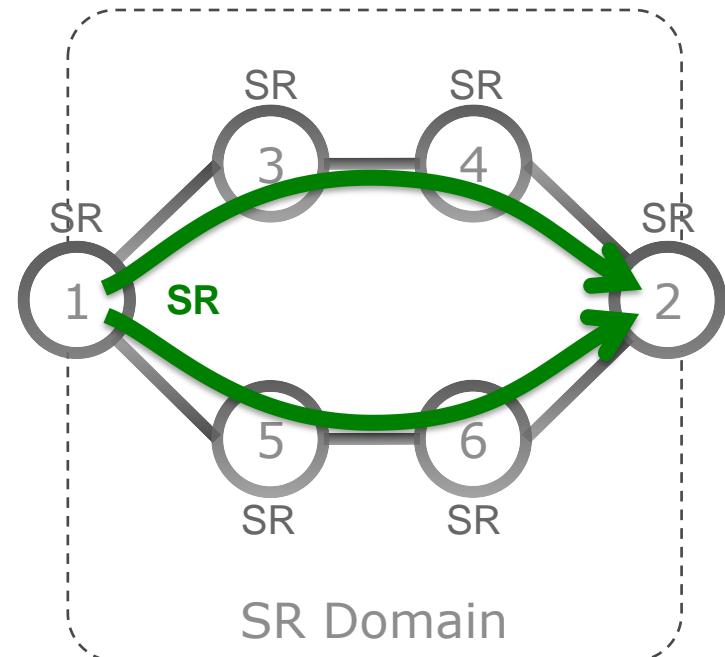


# Simplest migration LDP to SR

- **Initial state:** All nodes run LDP, not SR
- **Step1:** All nodes are upgraded to SR
  - In no particular order
  - leave default LDP label imposition preference
- **Step2:** All PEs are configured to prefer SR label imposition
  - In no particular order
- **Step3:** LDP is removed from the nodes in the network
  - In no particular order
- **Final state:** All nodes run SR, not LDP

Assumptions:

- all the nodes can be upgraded to SR
- all the services can be upgraded to SR





SR Policy Concept  
VPN SR TE based on  
Low Latency/Constraint/WECMP

# SR References

[www.segment-routing.net/tutorials/](http://www.segment-routing.net/tutorials/)

SR Traffic Engineering >

SR TE

# 3-tiered latency service

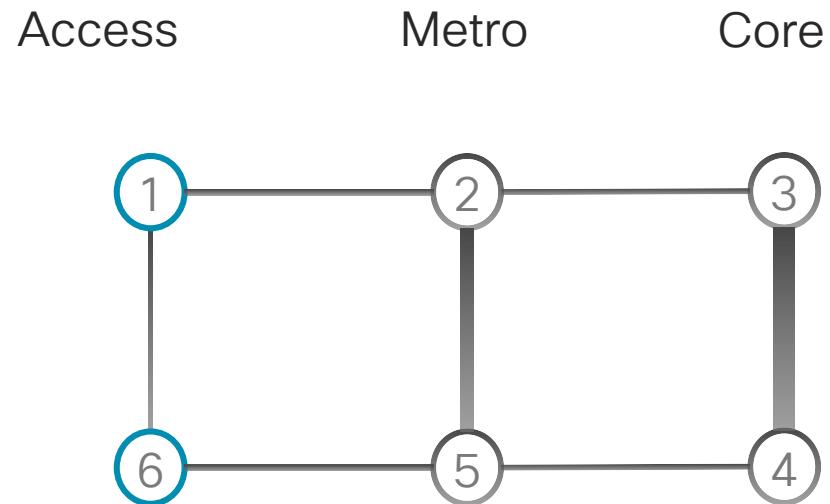
- Minimum Latency
  - ⇒ minimize the delay to the endpoint
- Minimum Cost
  - ⇒ minimize the isis metric to the endpoint
- Minimum Cost but with a bound on the delay
  - SRTE Policy
  - SR Native Algorithm
  - Minimize IGP metric
  - Cumulative delay along the path  $\leq$  delay bound



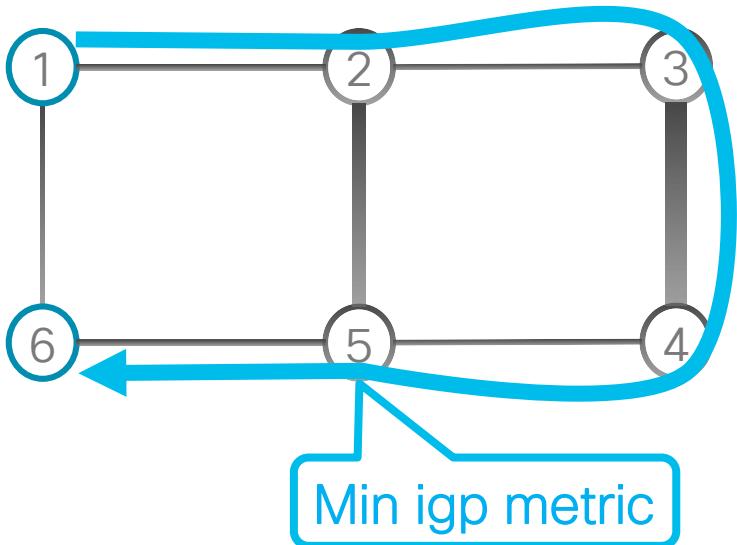
Another obvious business service that was never realized before SR

# Business Relevance: Cost vs Latency

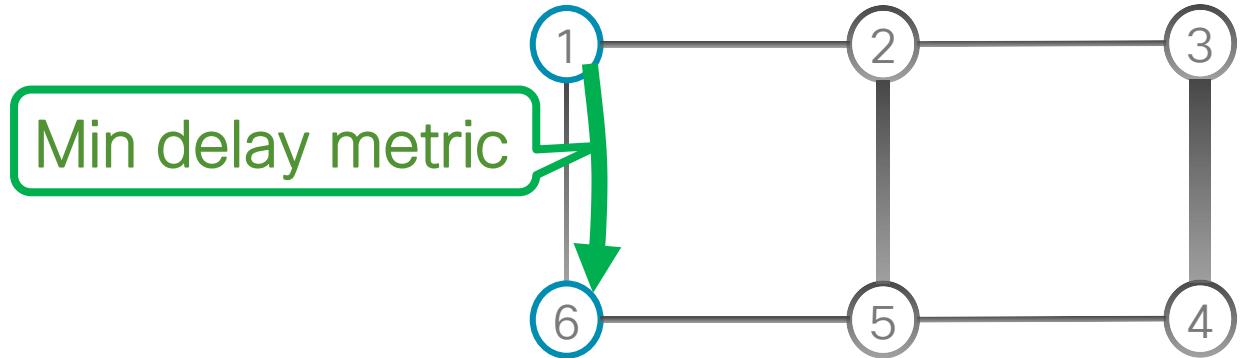
- Core
  - Cheaper
  - Longer delay
- Access Shortcut
  - Most expensive
  - Lowest latency
- Metro Shortcut
  - In-between



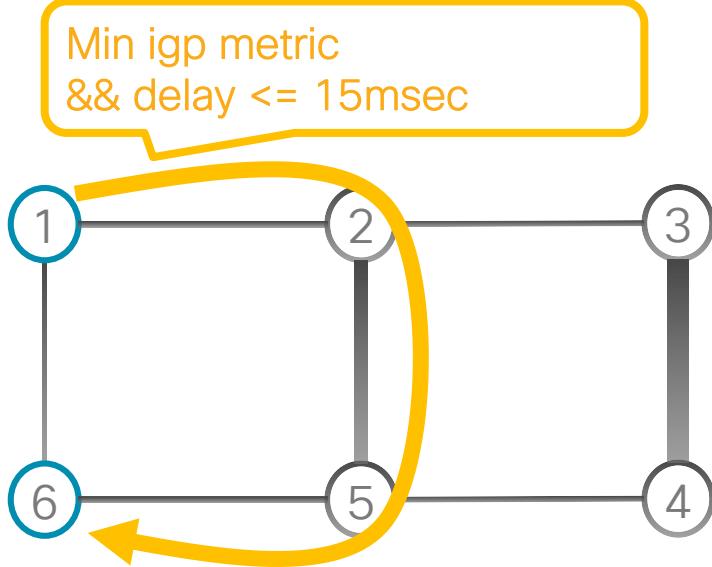
# Bulk of the traffic - Best-Effort



# Ultra Low Latency – 5G

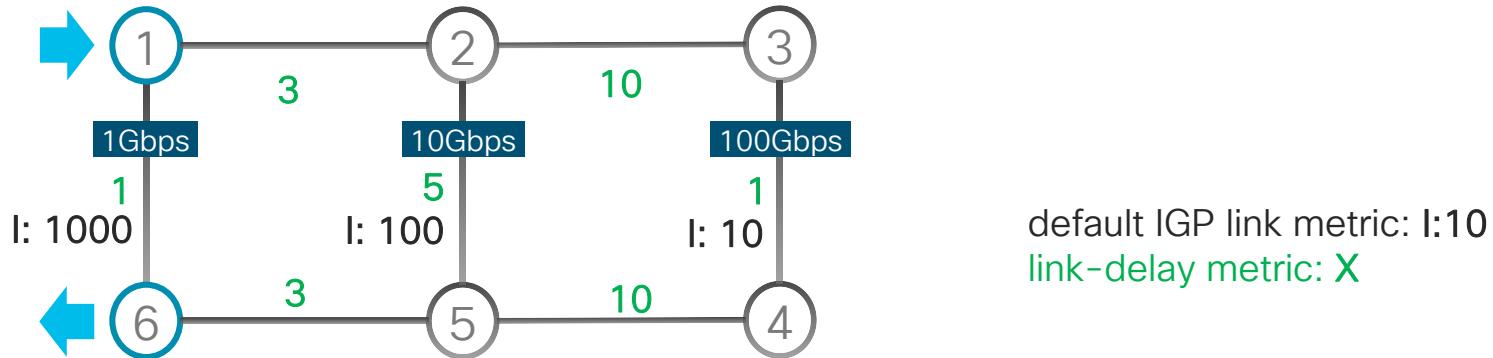


# Business traffic with constraint



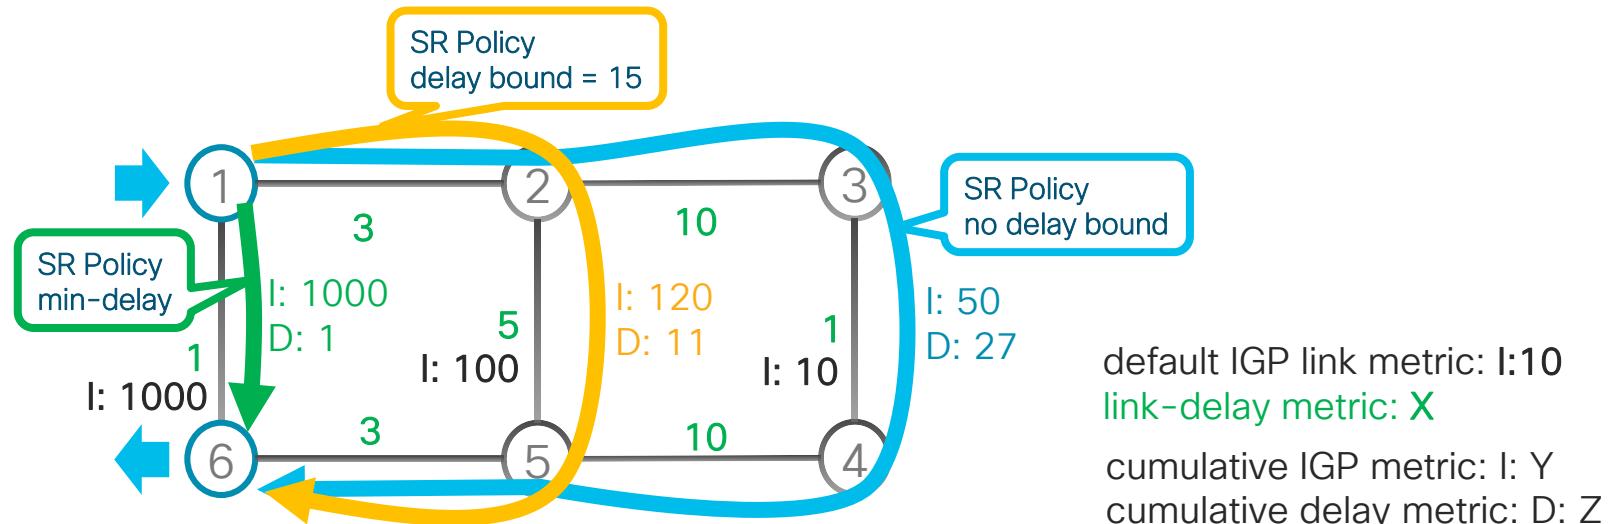
# Use-case: min IGP metric with bound on delay

- SRTE returns the shortest IGP path that has a cumulative delay < bound
- Goal: improve network utilization by moving traffic closer to the core while maintaining a delay SLA

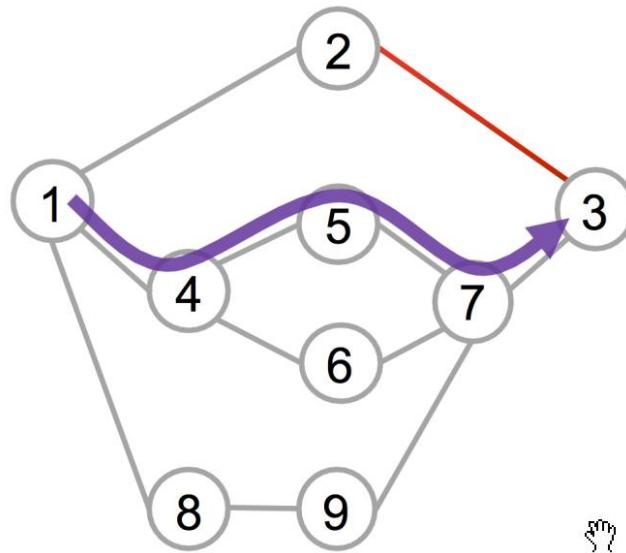


# Use-case: min IGP metric with bound on delay

- SRTE returns the shortest IGP path that has a cumulative delay < bound
- Goal: improve network utilization by moving traffic closer to the core while maintaining a delay SLA

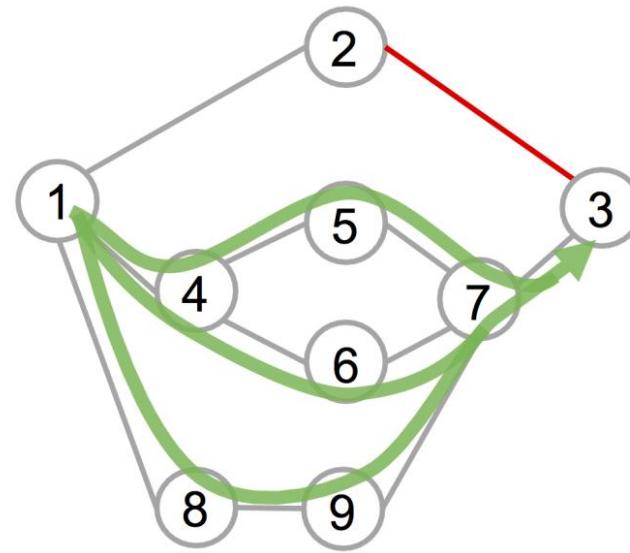


# RSVP-TE Optimization vs SR Optimization



Pre-SRTE is circuit-based  
CSPF => non-ECMP path

Re-using this for SRTE is not good  
SID List: <24014, 24045, 24057, 24073>  
Poor ECMP, big SID list, ATM-optimized

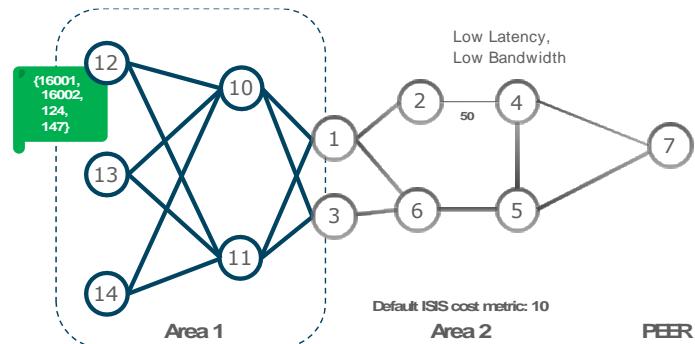


SR-native TE is needed  
Recognized Innovation - SIGCOMM 2015

SID List: <16007, 16003>  
ECMP, Small SID list, IP-optimized

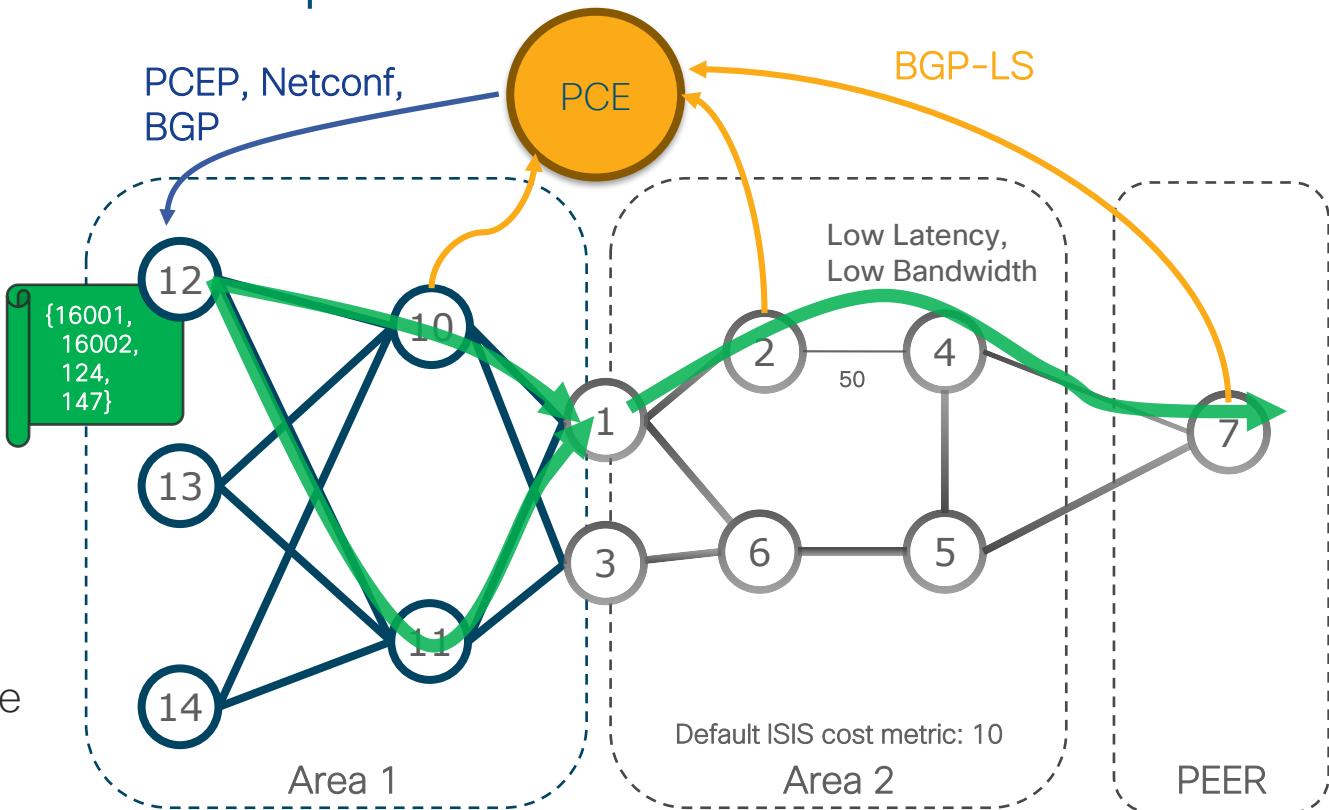
# Distributed vs Centralised Path Computation

- There are two possibilities to compute the dynamic path for an SR Policy:
  - Head-end computes the path itself (“distributed”)
  - Head-end requests SR PCE to compute the path (“centralised”)
- By default, dynamic paths are computed by the head-end
- Head-end uses an external PCE when local computation is not possible
  - An external PCE is required if more information is needed than is available on a head-end; e.g. multi-area/domain paths, or disjoint paths from different head-ends



# Centralized Path Computation Element

- The head-end is configured to use the PCE for dynamic path computation
- The PCE computes that the green path can be encoded as
  - 16001
  - 16002
  - 124
  - 147
- The PCE programs the path to the head-end



# SRTE

- Simple, Automated and Scalable
  - No core state: **state in the packet header**
  - No tunnel interface: “**SR Policy**”
  - No head-end a-priori configuration: **on-demand** policy instantiation
  - No head-end a-priori steering: **automated** steering
- Multi-Domain
  - SR PCE for compute
  - Binding-SID (BSID) for scale
- Lots of Functionality
  - Designed with **lead operators** along their use-cases

I

# SR Policy Identification

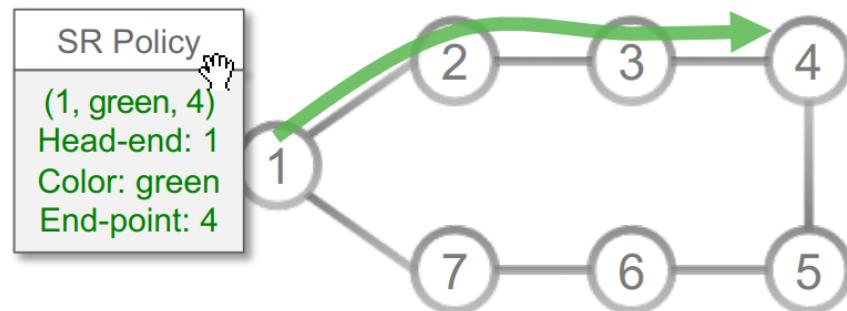
- An SR Policy is uniquely identified by a tuple  
**(head-end, color, end-point)**

Head-end: where the SR Policy is instantiated (*implemented*)

Color: a numerical value to differentiate multiple SRTE Policies between the same pair of nodes

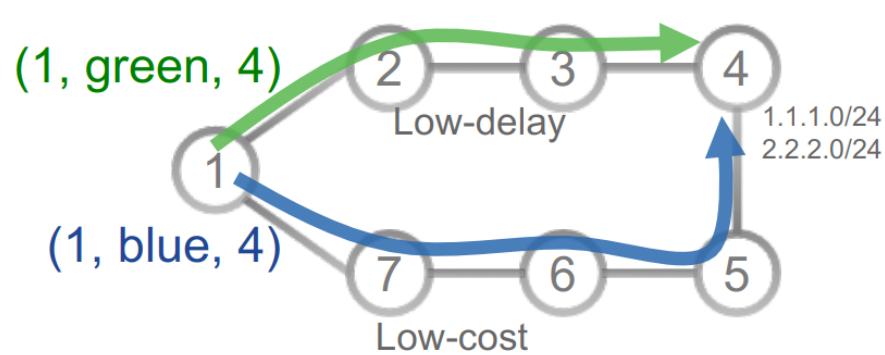
End-point: the destination of the SR Policy

- At a given head-end, an SR Policy is uniquely identified by a tuple **(color, end-point)**



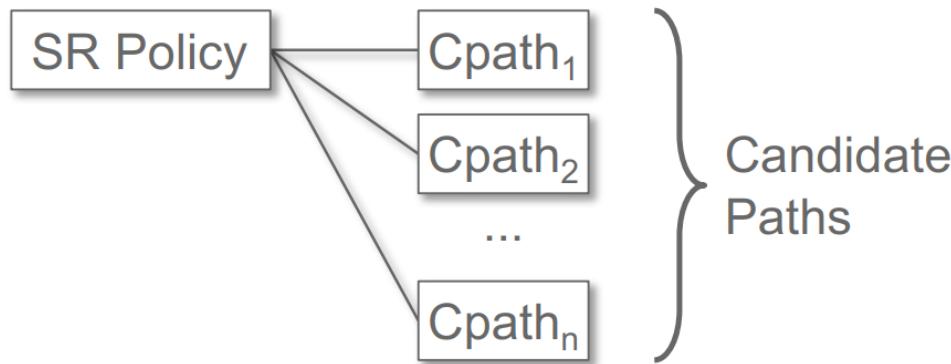
# SR Policy Color

- Each SR Policy has a color
  - Color can be used to indicate a certain treatment (SLA, policy) provided by an SR Policy
- Only one SR Policy with a given color C can exist between a given node pair (head-end (H), end-point (E))
  - In other words: each SR Policy triplet (H, C, E) is unique
- Example:
  - Low-cost=“blue”, Low-delay=“green”
  - steer traffic to 1.1.1.0/24 via Node4 into Low-cost SR Policy (1, blue, 4)
  - steer traffic to 2.2.2.0/24 via Node4 into Low-delay SR Policy (1, green, 4)



# SR Policy – Candidate Paths

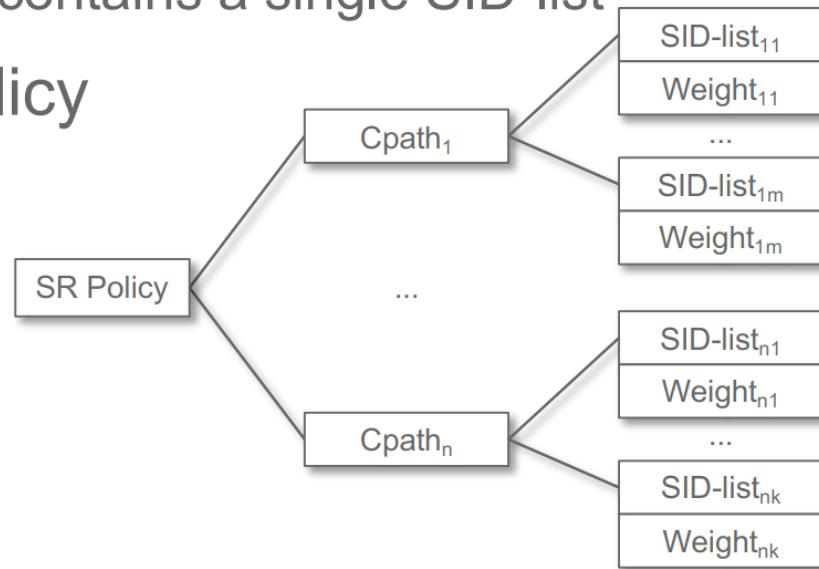
- An SR Policy consists of one or more candidate paths (Cpaths)



- An SR Policy instantiates one single path in RIB/FIB
  - the selected\* path, which is the preferred valid candidate path
- A candidate path is either dynamic or explicit

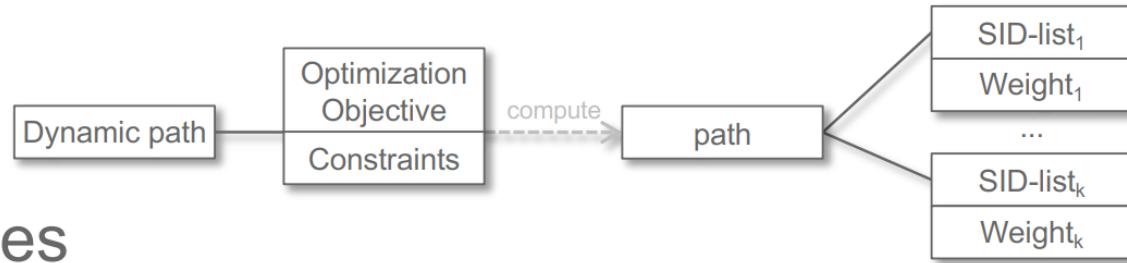
# SR Policy – Candidate Path

- A candidate path is a single segment list (SID-list) or a set of weighted\* SID-lists
  - Typically, an SR Policy path only contains a single SID-list
- Traffic steered into an SR Policy path is load-shared over all SID-lists of the path



SID = Segment ID

# Dynamic Path



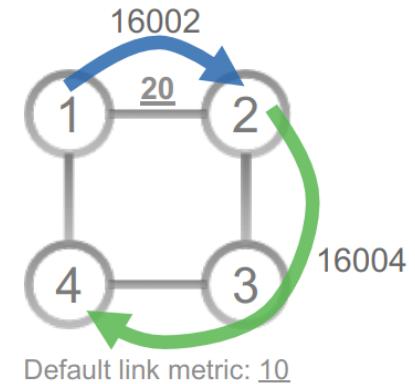
- A dynamic path expresses an **optimization objective** and **a set of constraints**
- The **head-end computes** a solution to the optimization problem as a SID-list or a set of SID-lists
- When the head-end does not have enough topological information (e.g. multi-domain problem), the head-end **may delegate the computation to a PCE**
- Whenever the network situation changes, the path is **recomputed**

# Topological path to SID-list

- After the path is computed, the SID-list that expresses the desired path is derived
- High-level algorithm:
  1. Node = head-end
  2. Find an IGP prefix-SID that leads as far down the desired path as possible (without using any link not included in the desired path)
  3. If no such prefix-SID exists, use the Adj-SID to the first neighbor along the path
  4. Node = the farthest node that is reached; goto 2.

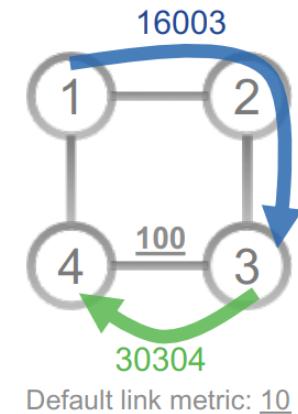
# Topological path to SID-list – Example 1

- Desired topological path = 1→2→3→4
- SID-list = <16002, 16004>
  - 16002 brings the packet from 1 to 2 (shortest path from Node1 to Node2)
  - 16004 brings the packet from 2 to 4 via 3 (shortest path from Node2 to Node4)



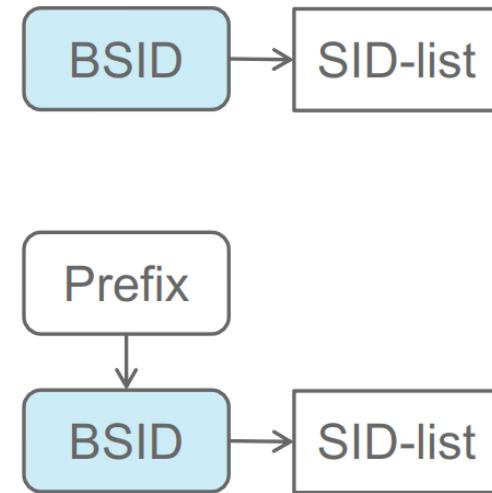
# Topological path to SID-list – Example 2

- Desired topological path = 1→2→3→4
- SID-list = <16003, 30304>
  - 16003 brings the packet from 1 to 3 (shortest path from Node1 to Node3)
  - 30304 brings the packet from 3 to 4 using the Adjacency-SID



# Binding-SID (BSID) is fundamental

- The BSID of the SR Policy selected path is installed in the forwarding table
- **Remote steering**
  - A packet arriving on the SR Policy head-end with the BSID as Active Segment (top of label stack) is steered into the SR Policy associated with the BSID
- **Local steering**
  - A packet that matches a forwarding entry that resolves on the BSID of an SR Policy is steered into that SR Policy



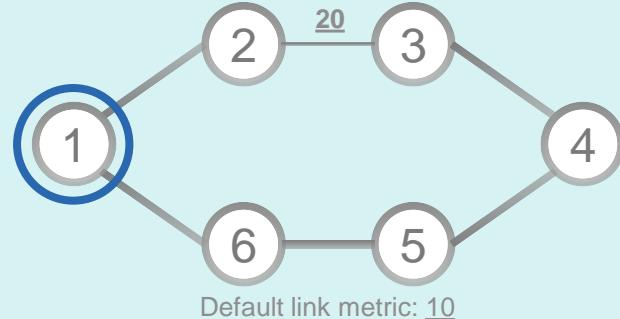
# SR Policy - Configuration Example

On Node1:

```
segment-routing
traffic-eng
policy POLICY1
color 20 end-point ipv4 1.1.1.4
binding-sid mpls 1000
candidate-paths
preference 100
dynamic
metric type te
constraints
affinity
exclude-any color red
!
preference 200
explicit segment-list SIDLIST1
!
segment-list name SIDLIST1
index 10 mpls label 16002
index 20 mpls label 30203
index 30 mpls label 16004
```

SRTE

SR Policy



# SR Policy – Configuration Example

On Node1:

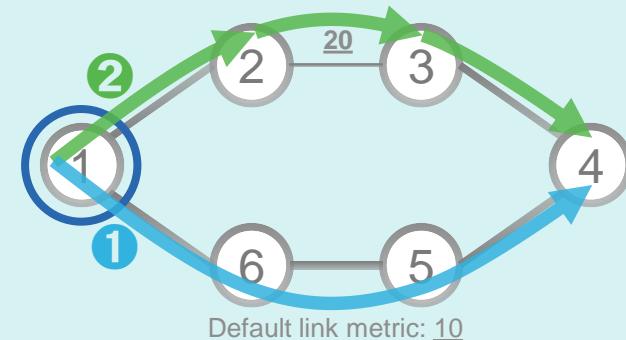
```
segment-routing
traffic-eng
policy POLICY1
color 20 end-point ipv4 1.1.1.4
binding-sid mpls 1000
candidate-paths
  1 preference 100
    dynamic
      metric type te
    constraints
      affinity
        exclude-any color red
  !
  2 preference 200
    explicit segment-list SIDLIST1
```

User-defined name

Color and End-point

Binding-SID

Local Candidate Paths



```
segment-list name SIDLIST1
index 10 mpls label 16002
index 20 mpls label 30203
index 30 mpls label 16004
```

# SR Policy – configuration example

On Node1:

```
segment-routing
  traffic-eng
    policy POLICY1
      color 20 end-point ipv4 1.1.1.4
      binding-sid mpls 1000
      candidate-paths
        1 preference 100
          dynamic
          metric type te
          constraints
          affinity
          exclude-any color red
        2 preference 200
          explicit segment-list SIDLIST1
      !
      segment-list name SIDLIST1
        index 10 mpls label 16002
        index 20 mpls label 30203
        index 30 mpls label 16004
```

Path preference 100

Dynamic path

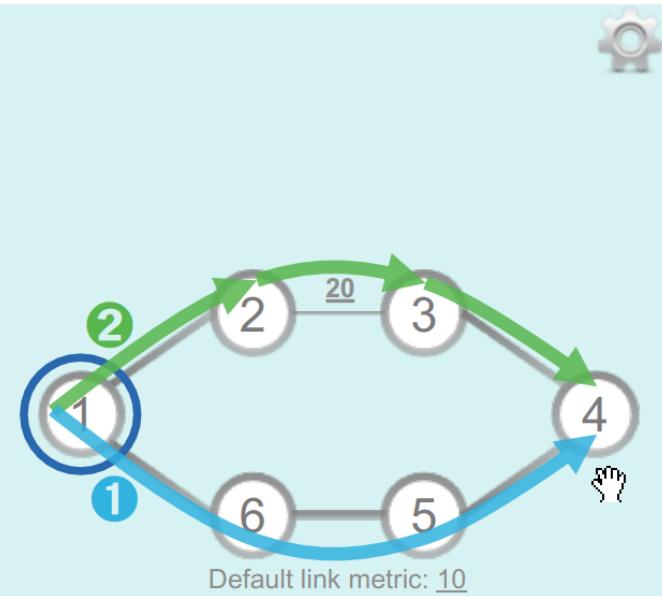
Opt. Obj.: TE metric

Constraint

Path preference 200

Explicit SID-list1

SID-list1



```
segment-routing
  traffic-eng
    affinity-map
    color red bit-position 0
```

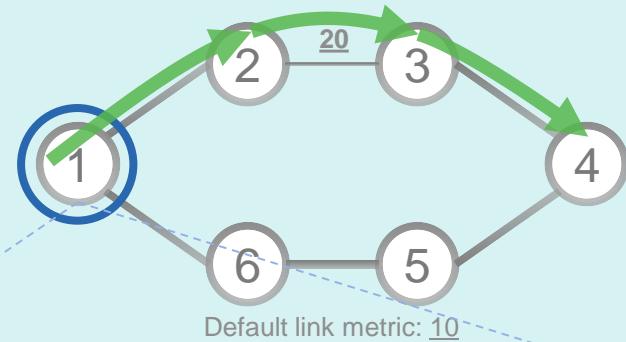
# SR Policy – Configuration Example

On Node1:

```
segment-routing
  traffic-eng
    policy POLICY1
      color 20 end-point ipv4 1.1.1.4
      binding-sid mpls 1000
      candidate-paths
        preference 100
          dynamic
            metric type te
        constraints
          affinity
          exclude-any red
      !
      preference 200
      explicit segment-list SIDLIST1
    !
    segment-list name SIDLIST1
      index 10 mpls label 16002
      index 20 mpls label 30203
      index 30 mpls label 16004
```

Selected Path:

- Valid Path
- Highest Pref value



FIB @ head-end Node1

Incoming label: 1000

Action: pop and push <16002, 30203, 16004>

# SR Policy – configuration example

On Node1:

```
segment-routing
  traffic-eng
    policy POLICY1
      color 20 end-point ipv4 1.1.1.4
      binding-sid mpls 1000
    candidate-paths
      1 preference 100
        dynamic
        metric type te
        constraints
          affinity
            exclude-any color red
      !
      2 preference 200
        explicit segment-list SIDLIST1
    !
    segment-list name SIDLIST1
      index 10 mpls label 16002
      index 20 mpls label 30203
      index 30 mpls label 16004
```

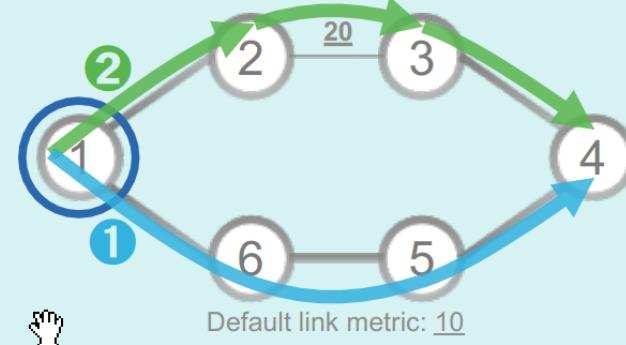
User-defined name

Color and End-point

Binding-SID

SR Policy ID:  
(20,1.1.1.4)

Local Candidate Paths



```
segment-routing
  traffic-eng
    affinity-map
      color red bit-position 0
```

# SR Policy – configuration example

On Node1:

```
segment-routing
  traffic-eng
    policy POLICY1
      color 20 end-point ipv4 1.1.1.4
      binding-sid mpls 1000
      candidate-paths
        preference 100
          dynamic mpls
          metric
            type te
          affinity
            exclude-any red
        !
        preference 200
          explicit segment-list SIDLIST1
        !
segment-list name SIDLIST1
  index 10 mpls label 16002
  index 20 mpls label 30203
  index 30 mpls label 16004
```

Node1 may receive other candidate paths for SR Policy (20, 1.1.1.4) from other sources, some examples:

Source of path is not considered for path selection

Selected Path:  
• Valid Path  
• Highest Pref value

Other candidate paths received for SR Policy (20, 1.1.1.4)

Path received via BGP signaling

preference 150  
binding-sid mpls 1000  
weight 1, SID-list <16002, 16005>  
weight 2, SID-list <16004, 16008>

Path received via PCEP signaling

preference 120  
binding-sid mpls 1000  
SID-list <16002, 16005>

Path received via NETCONF signaling

preference 50  
binding-sid mpls 1000  
SID-list <16002, 16005>

# WECMP example

On Node1:

```
segment-routing
  traffic-eng
    policy POLICY1
      color 20 end-point ipv4 1.1.1.4
      binding-sid mpls 1000
      candidate-paths
        preference 200
        explicit segment-list SIDLIST1
          weight 1
        !
        explicit segment-list SIDLIST2
          weight 4
        !
      segment-list name SIDLIST1
        index 10 mpls label 16002
        index 20 mpls label 30203
        index 30 mpls label 16004
      !
      segment-list name SIDLIST2
        index 10 address ipv4 1.1.1.4
```

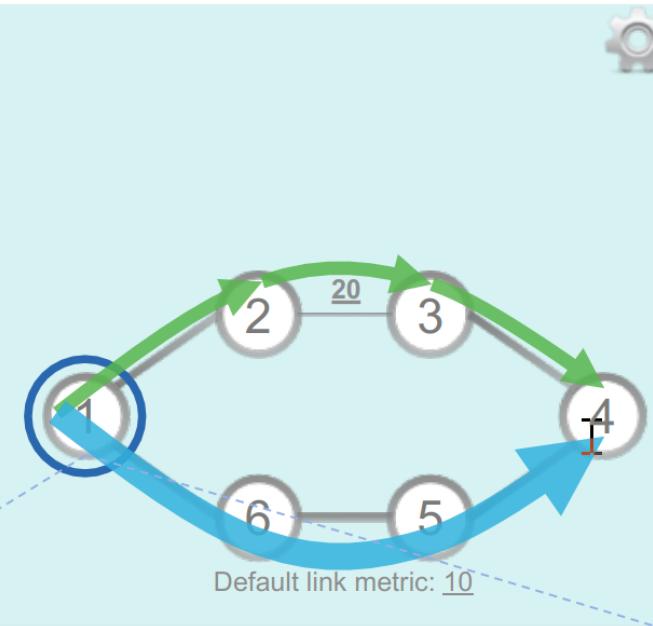
Path preference  
200

Explicit SID-list1,  
Weight 1

Explicit SID-list2,  
Weight 4

SID-list1

SID-list2

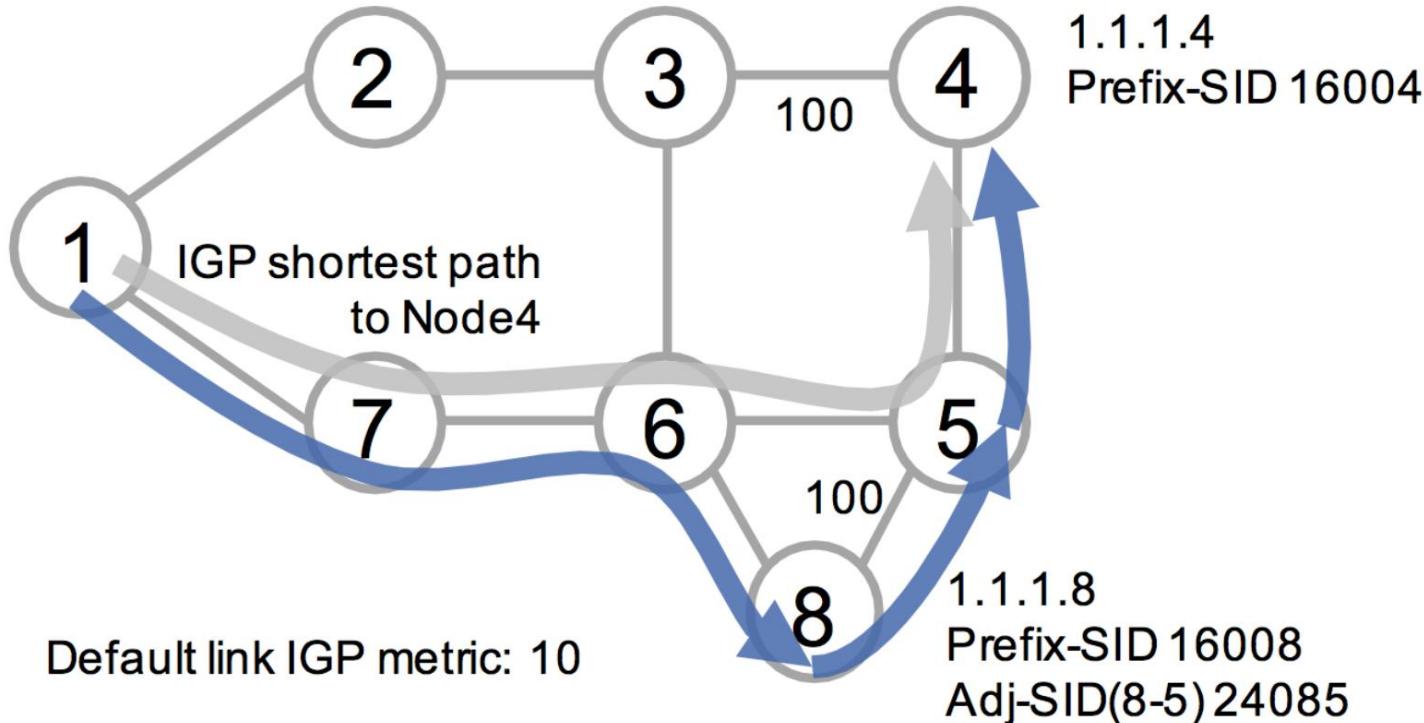


FIB @ head-end Node1

Incoming label: 1000

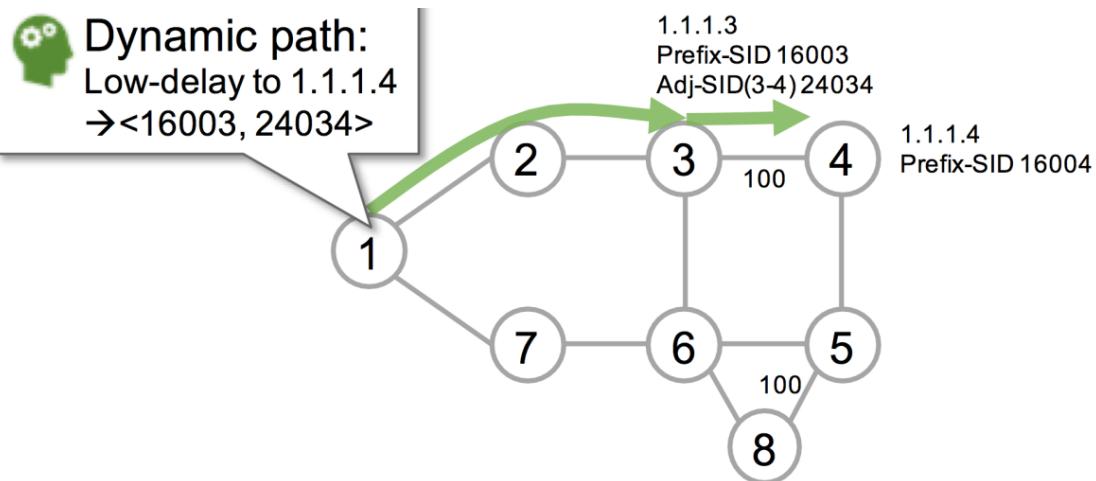
Action:pop and push <16002, 30203, 16004> (20%)  
push <16004> (80%)

# SR Policy : Dynamic Path & Explicit Path



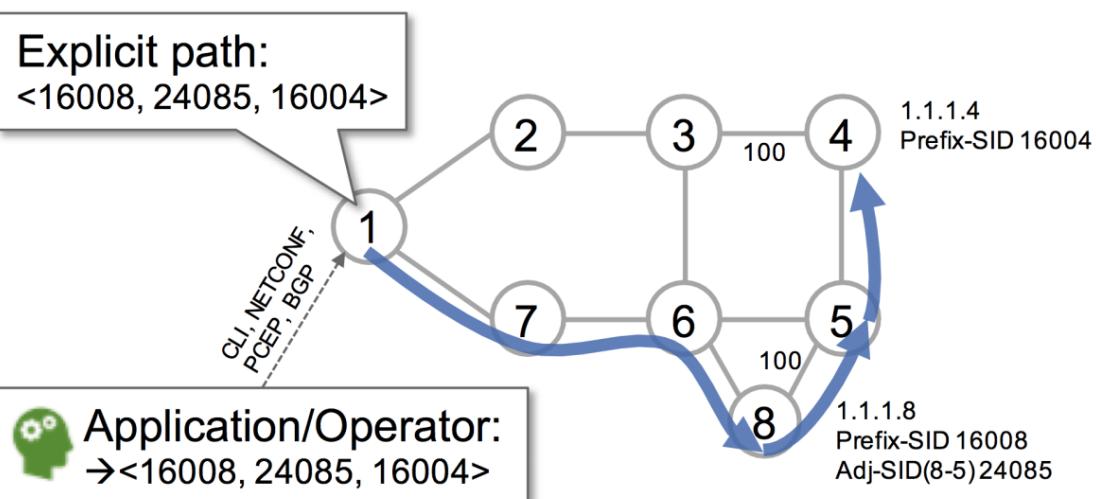
# Dynamic Path

(a)

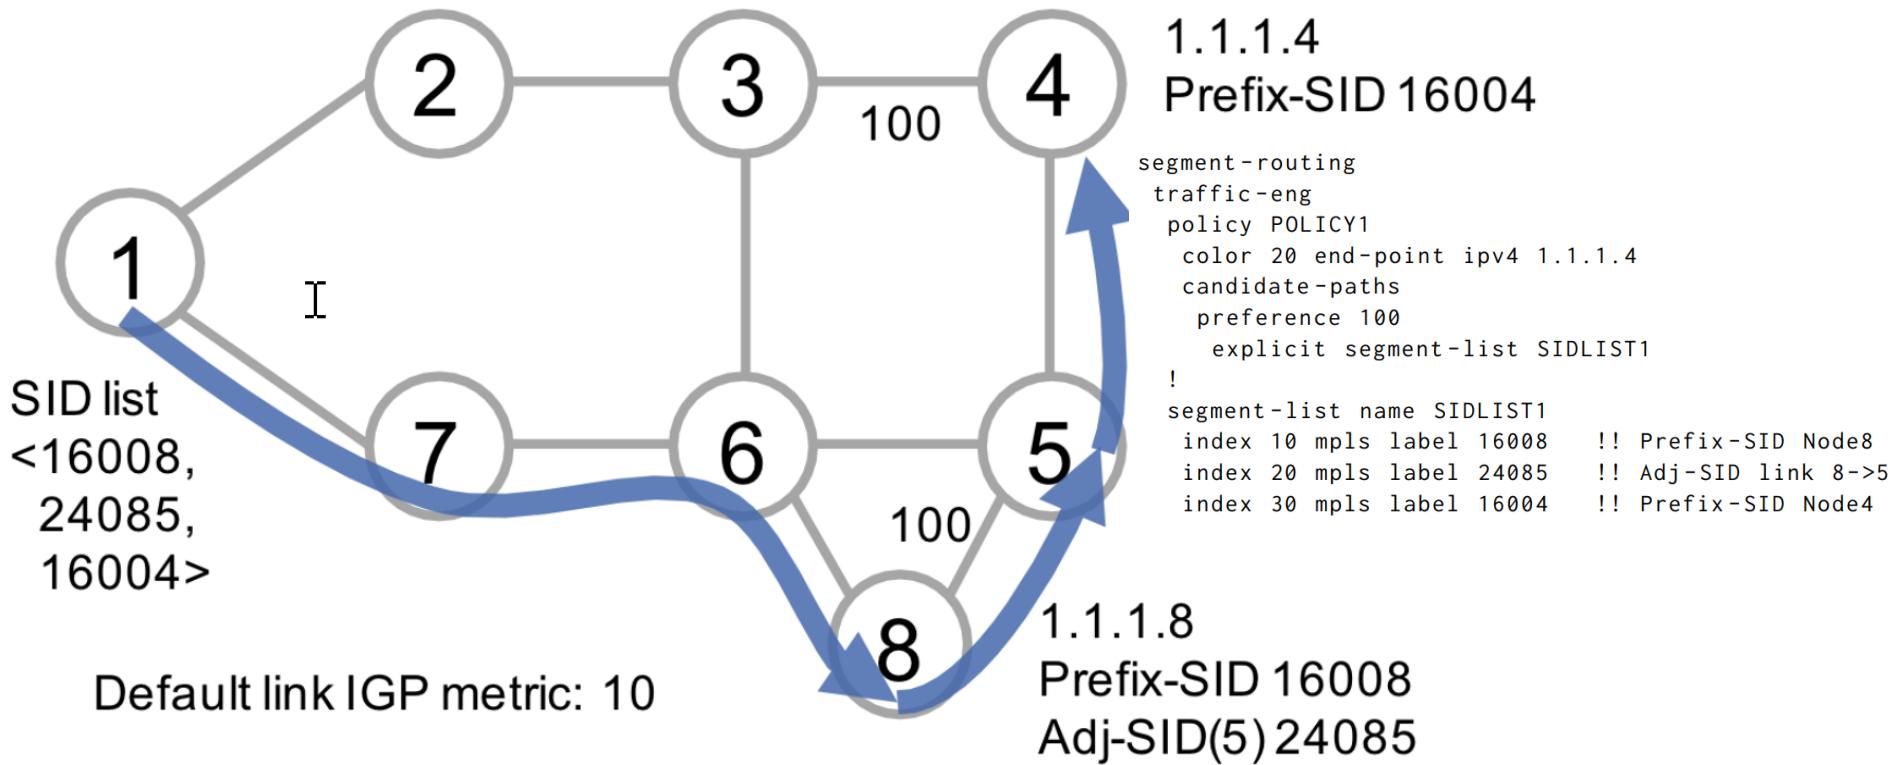


# Explicit Path

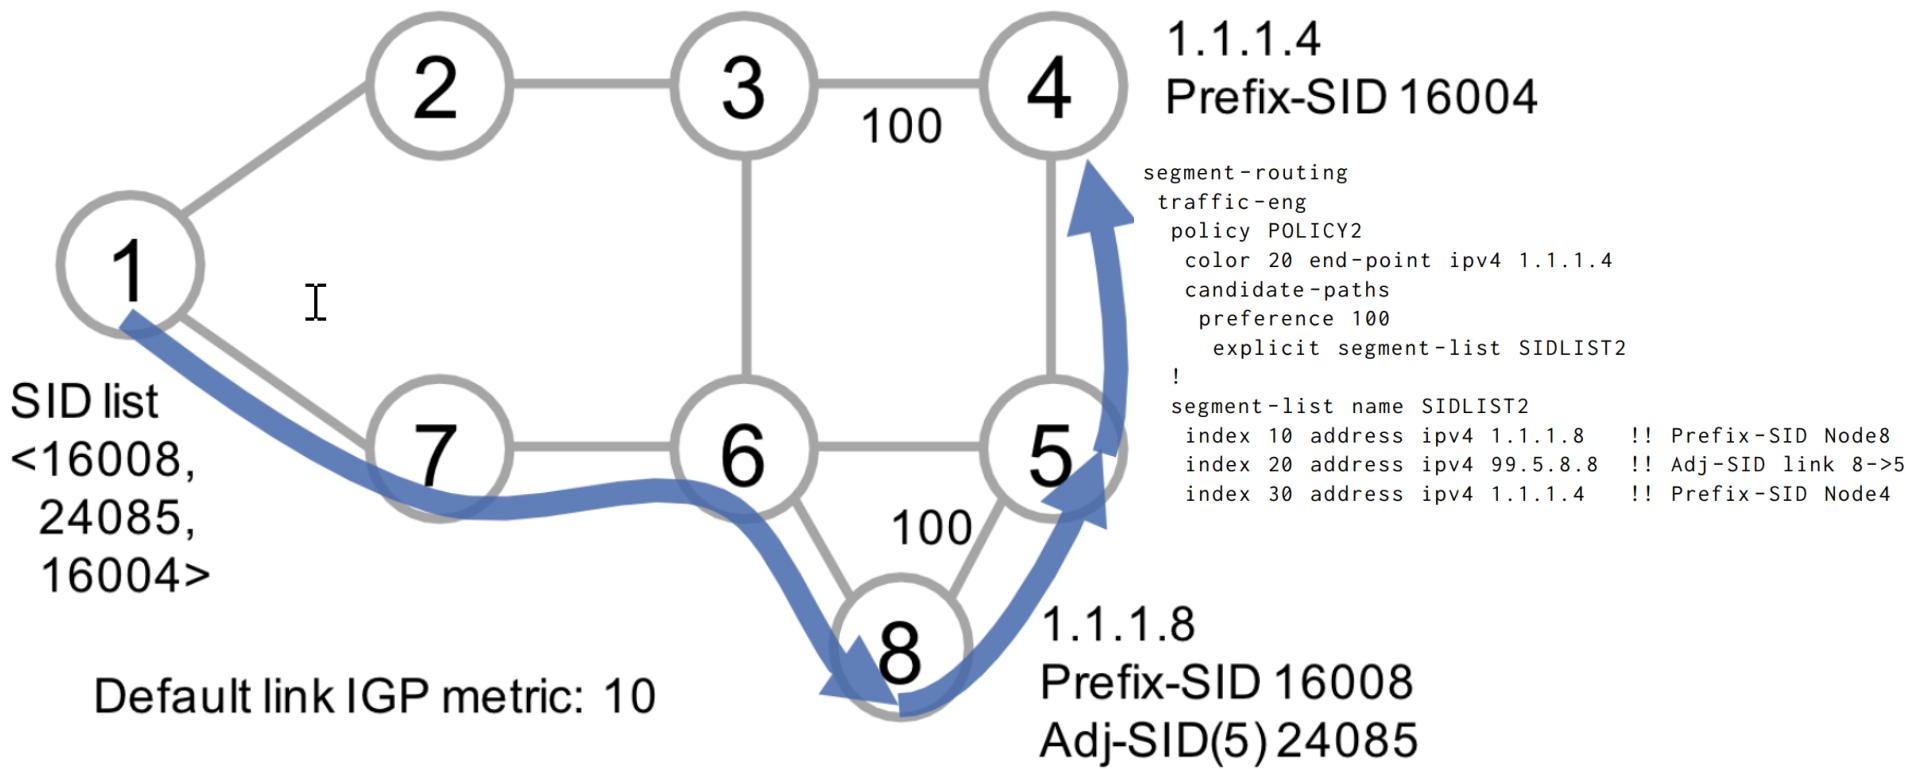
(b)



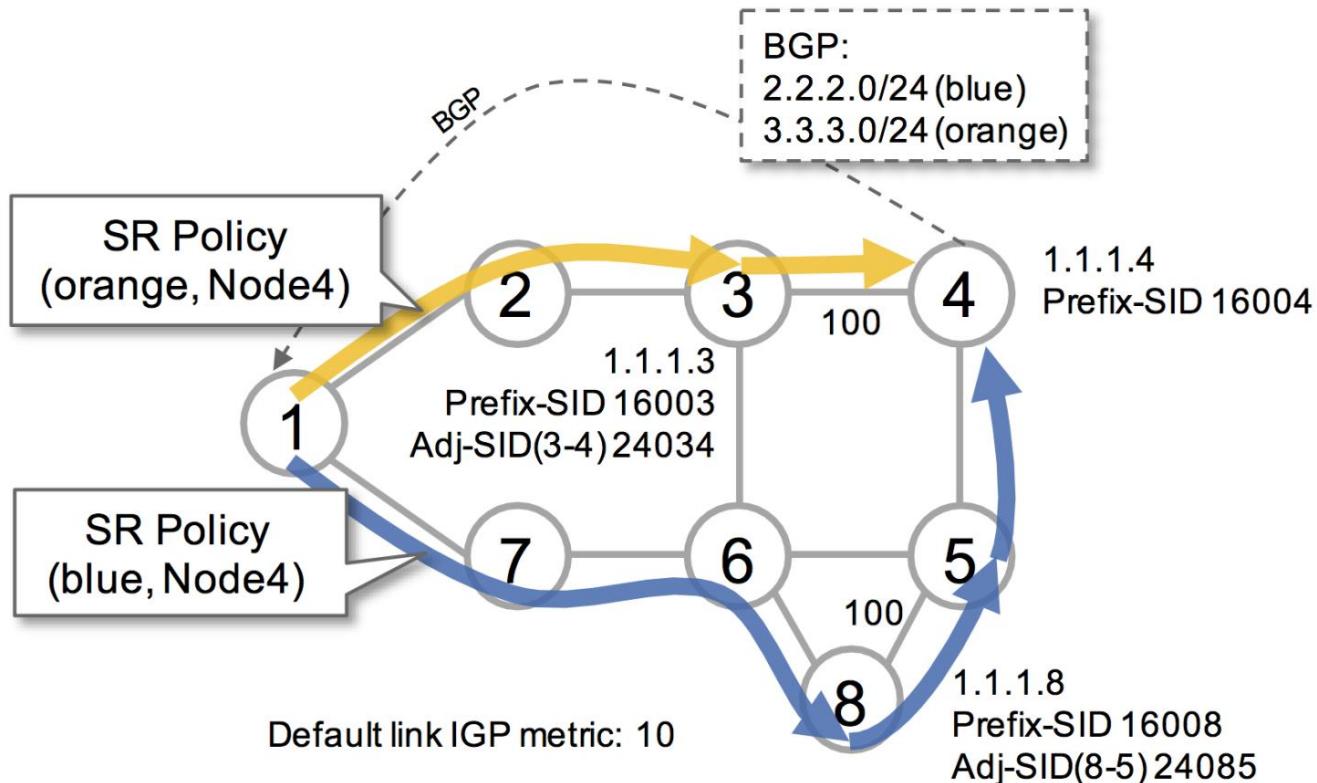
# SR Policy : Explicit Path SID List with MPLS Label



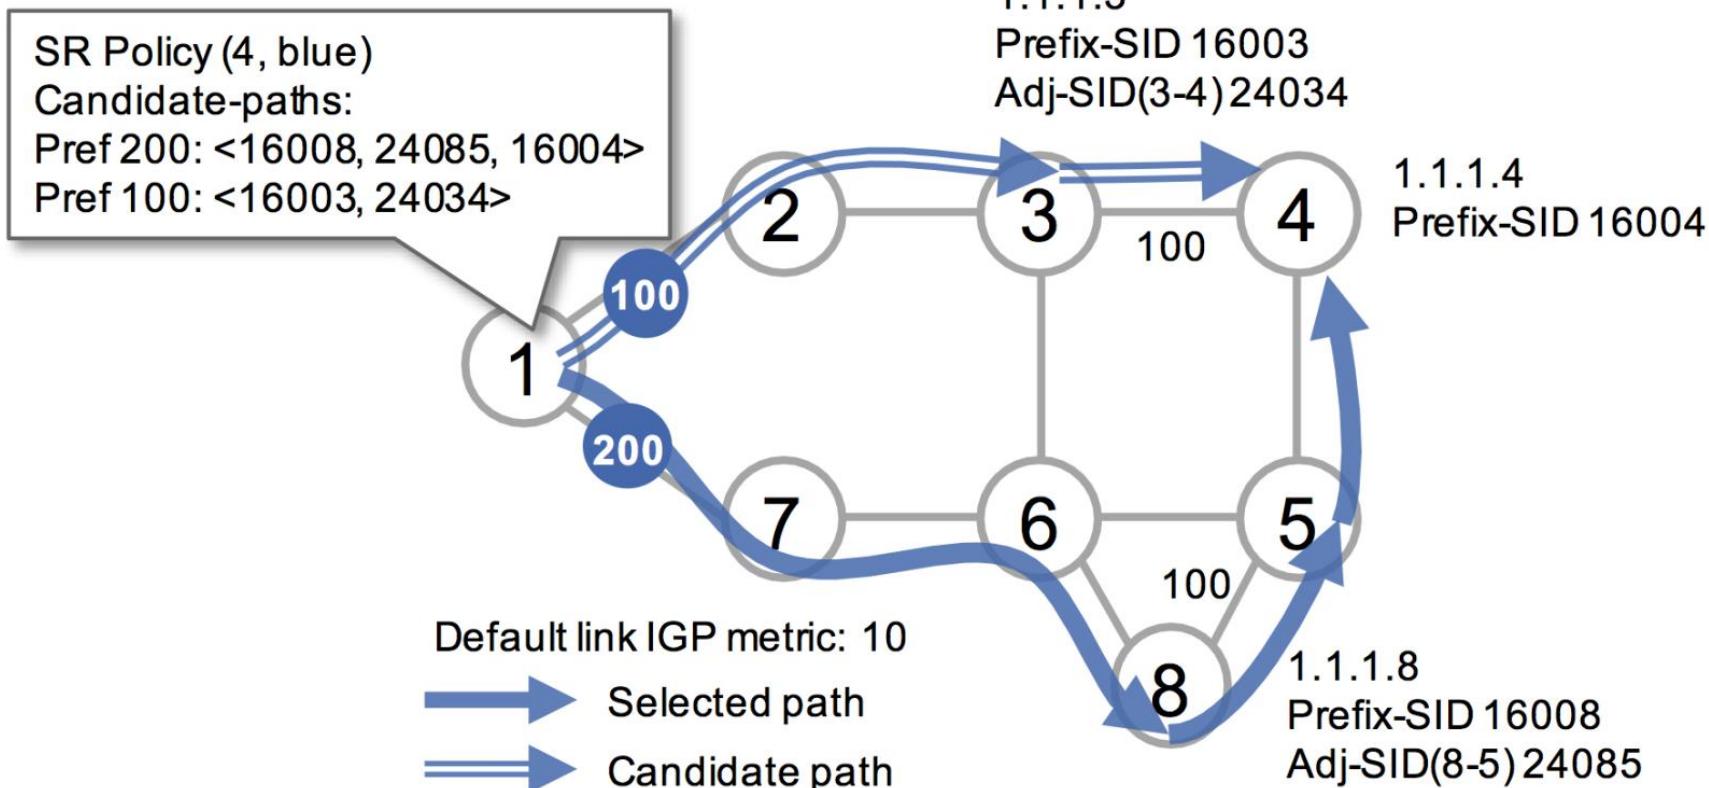
# SR Policy : Explicit Path SID List with IP Address



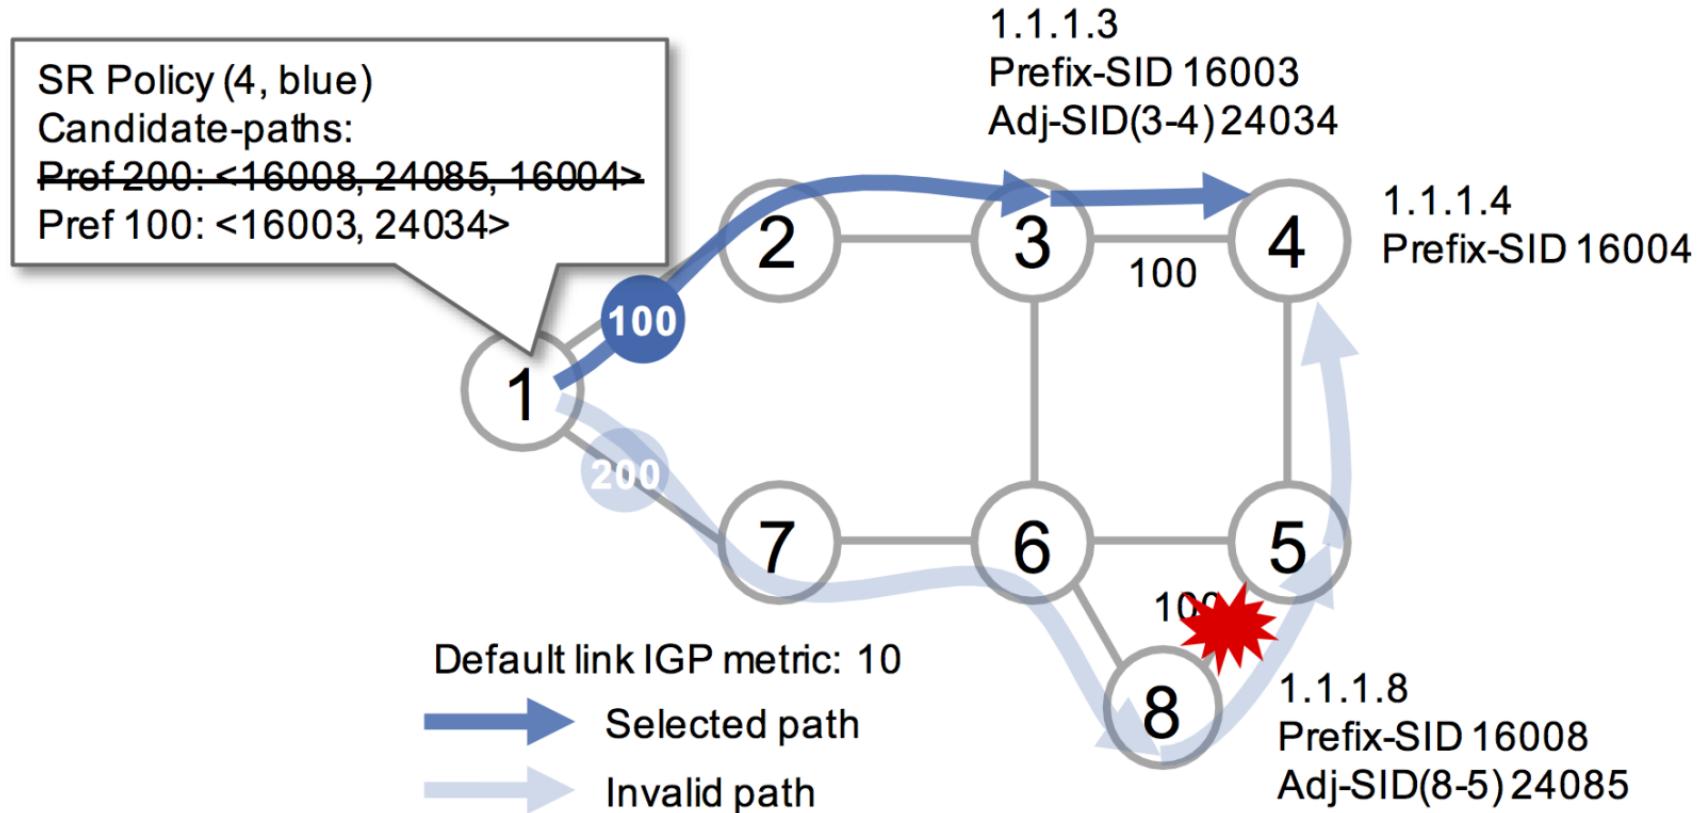
# SR Policy with Single Explicit Candidate Path



# SR Policy with Two Explicit Candidate Paths

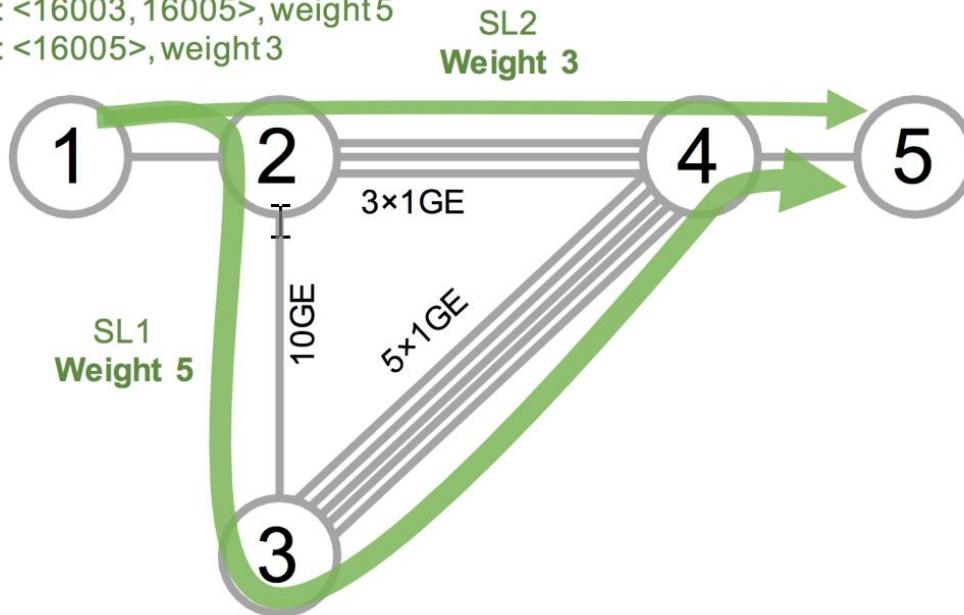


# SR Policy with Two Explicit Candidate Paths



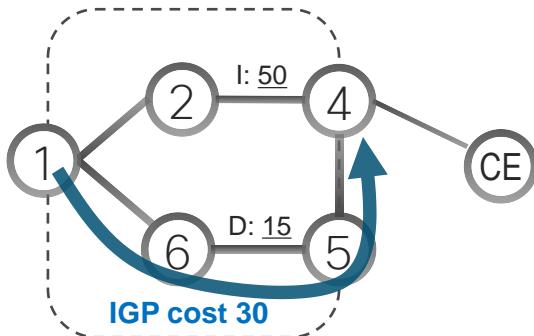
# Use-case : Weighted ECMP SR Policy

SR Policy WECMP  
SL1: <16003, 16005>, weight 5  
SL2: <16005>, weight 3

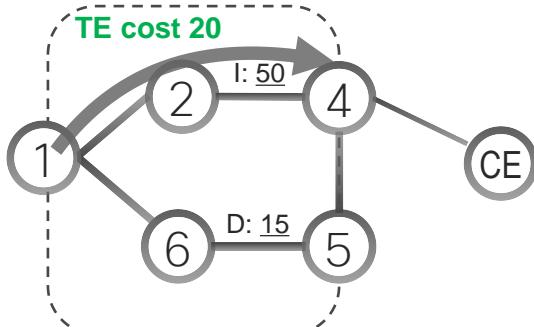


segment - routing  
traffic -eng  
segment - list SL1  
index 10 mpls label 16003  
index 20 mpls label 16005  
!  
segment - list SL2  
index 10 mpls label 16005  
!  
policy WECMP  
color 20 end - point ipv4 1.1.1.5  
candidate - paths  
preference 100  
explicit segment - list SL1  
**weight 5**  
!  
explicit segment - list SL2  
**weight 3**

# Different VPNs need different underlay SLA



**Basic VPN should use lowest cost underlay path**

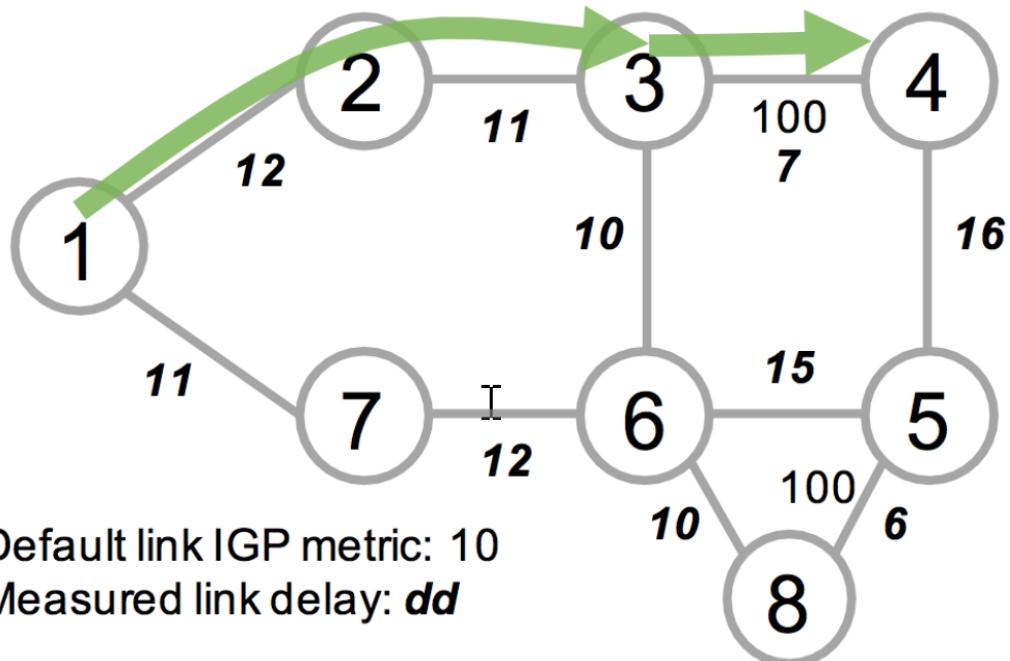


**Premium VPN should use lowest delay path**

**Objective:**  
operationalize this service for simplicity, scale and performance

# SR Policy with Dynamic Path + Low Delay Intent

1.1.1.3  
Prefix-SID 16003  
Adj-SID(3-4) 24034



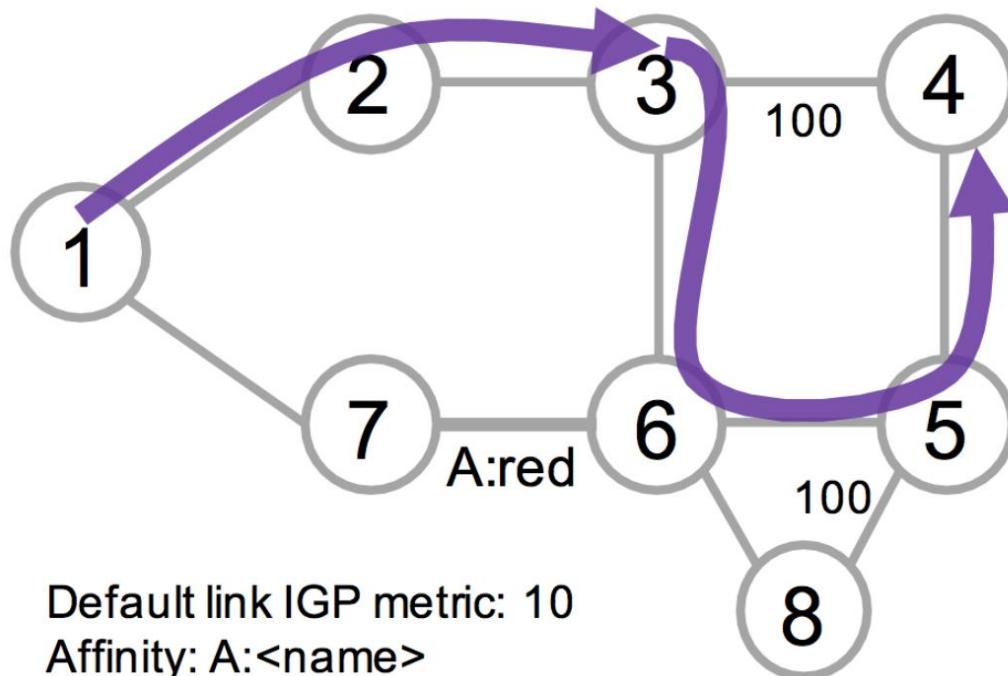
1.1.1.4  
Prefix-SID 16004

```
segment-routing
traffic-eng
policy POLICY2
color 30 end-point ipv4 1.1.1.4
candidate-paths
preference 100
dynamic mpls
metric
type delay
```

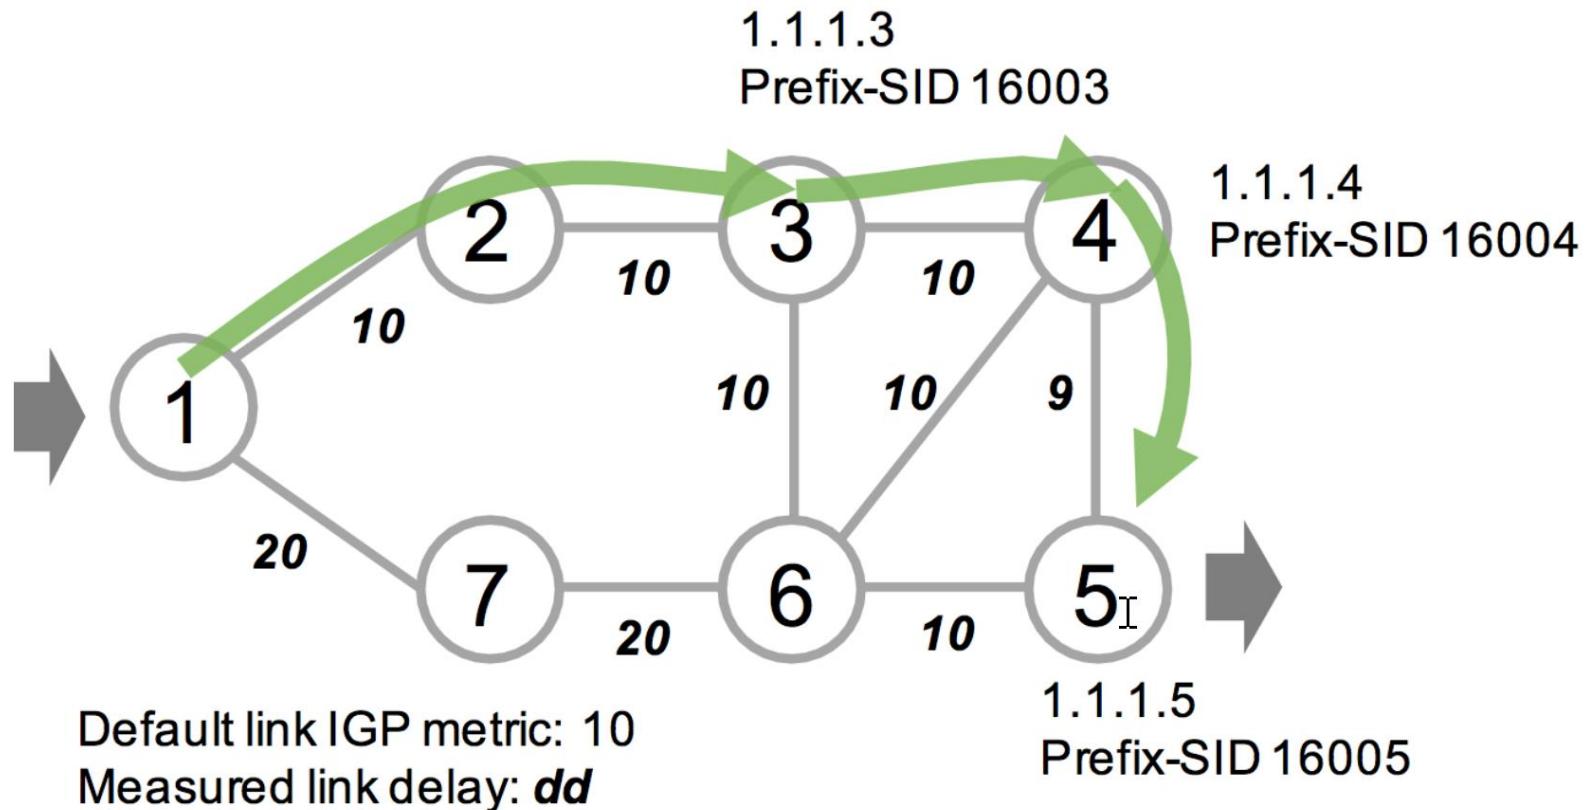
# SR Policy with Dynamic Path + Avoiding Specific Link

1.1.1.3  
Prefix-SID 16003  
Adj-SID(3-4) 24034

1.1.1.4  
Prefix-SID 16004



# Use-Case : Low Delay Margin for ECMP



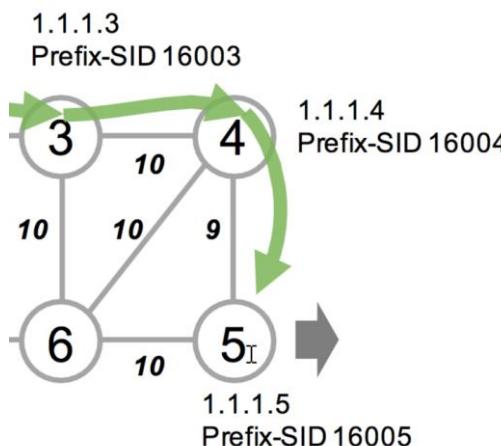
# Use-Case : Low Delay Margin for ECMP

```
segment-routing
  traffic-eng
    policy LOW-DELAY
      color 20 end-point ipv4 1.1.1.5
      candidate-paths
        preference 100
        dynamic
          metric
            type delay
```

No ECMP

R3 – R4 – R5 (delay = 19)

Lowest Delay Metric = 19



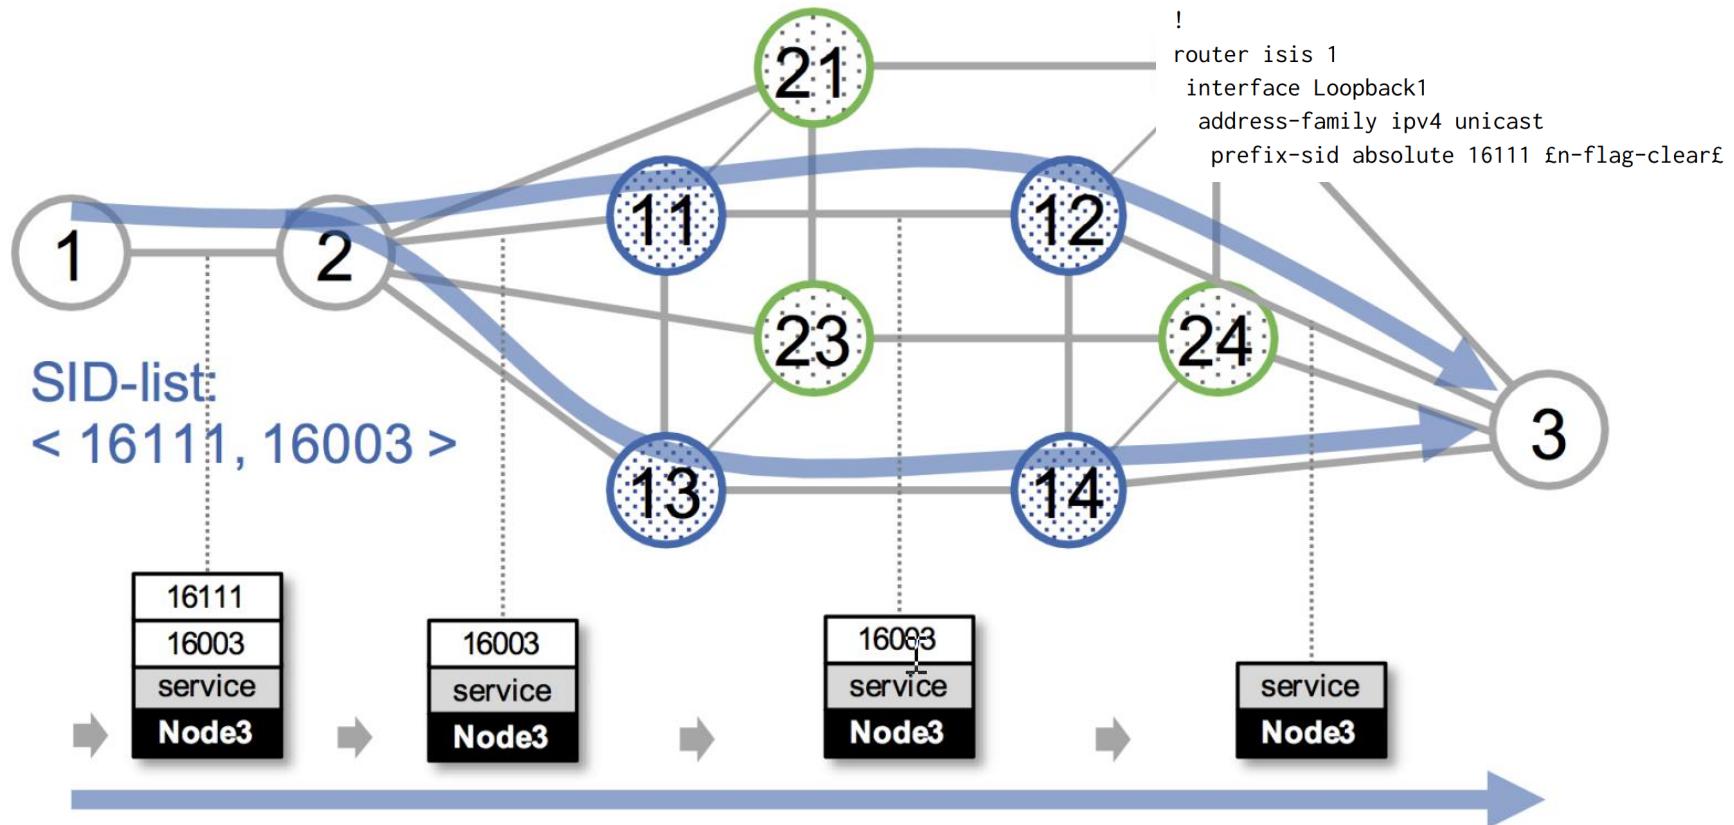
```
segment-routing
  traffic-eng
    policy LOW-DELAY
      color 20 end-point ipv4 1.1.1.5
      candidate-paths
        preference 100
        dynamic
          metric
            type delay
            margin absolute 2000
```

ECMP

R3 – R4 – R5 (delay = 19)  
R3 – R6 – R5 (delay = 20)

Margin = 2 ms (2000 us)  
Lowest + 2 = 19 + 2 = 21 ms

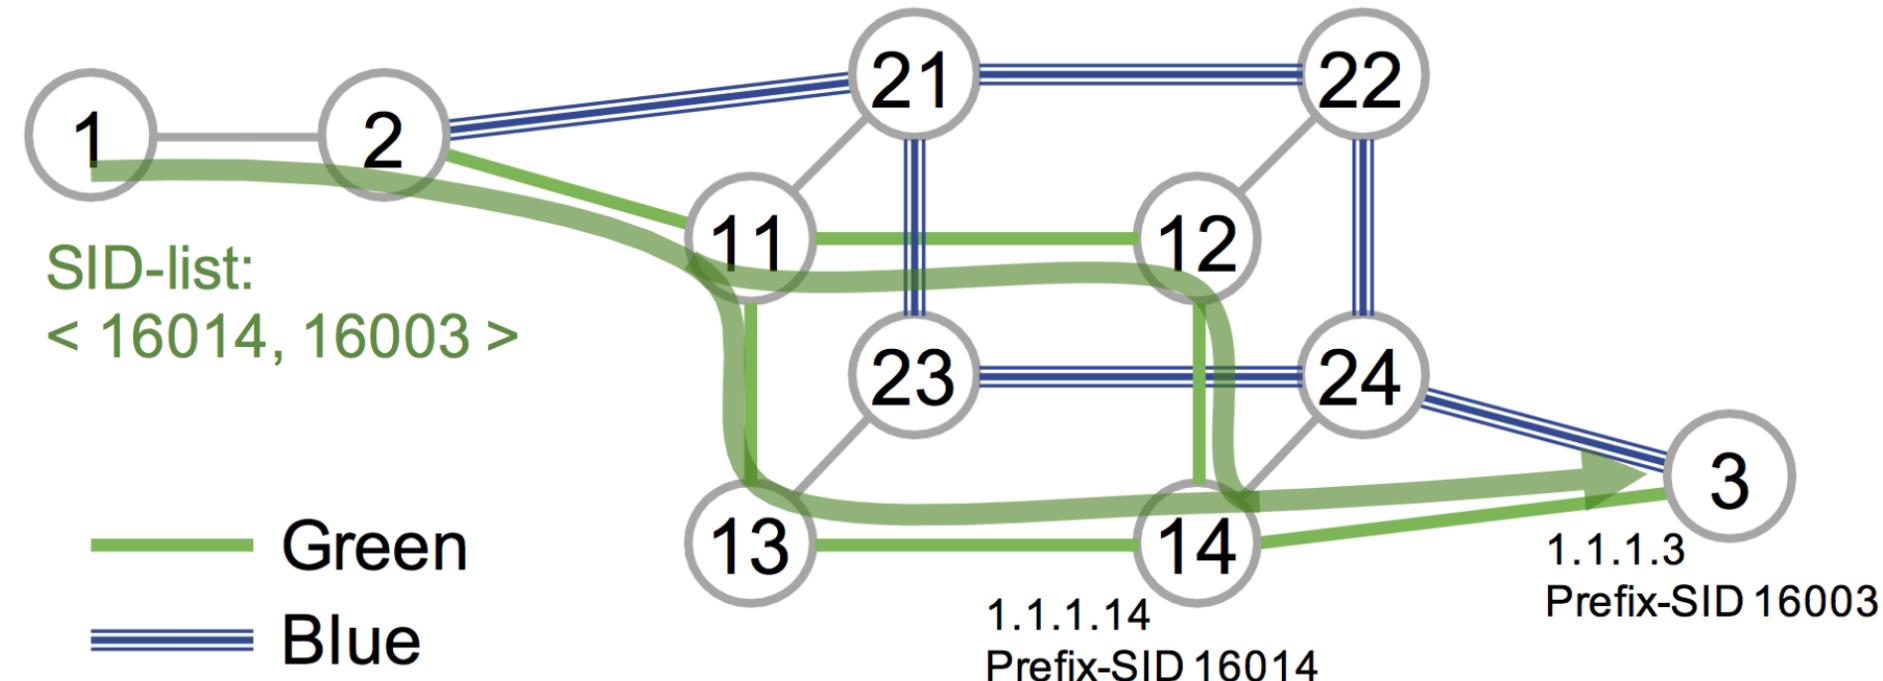
# Dual Plane Disjoint with Anycast-SID



# Dual Plane Disjoint with Anycast-SID

```
interface Loopback1
description blue plane anycast address
ipv4 address 1.1.1.111/32
!
router isis 1
interface Loopback1
address-family ipv4 unicast
prefix-sid absolute 16111 fn-flag-clea
segment-routing
traffic-eng
policy BLUE
color 10 end-point ipv4 1.1.1.3
candidate-paths
preference 100
explicit segment-list SIDLIST1
!
policy GREEN
color 20 end-point ipv4 1.1.1.3
candidate-paths
preference 100
explicit segment-list SIDLIST2
!
segment-list name SIDLIST1
index 10 address ipv4 1.1.1.111 !! blue plane anycast
index 20 address ipv4 1.1.1.3
!
segment-list name SIDLIST2
index 10 address ipv4 1.1.1.222 !! green plane anycast
index 20 address ipv4 1.1.1.3
```

# Use-Case : Dual Plane Affinity Resource Avoidance



# Use-Case : Dual Plane Affinity Resource Avoidance

```
segment-routing
traffic-eng
affinity-map
  name GREEN bit-position 0
  name BLUE bit-position 2
!
interface Gi0/0/0/0
  !! link to Node11, in plane Green
  affinity name GREEN
!
interface Gi0/0/0/1
  !! link to Node21, in plane Blue
  affinity name BLUE
```

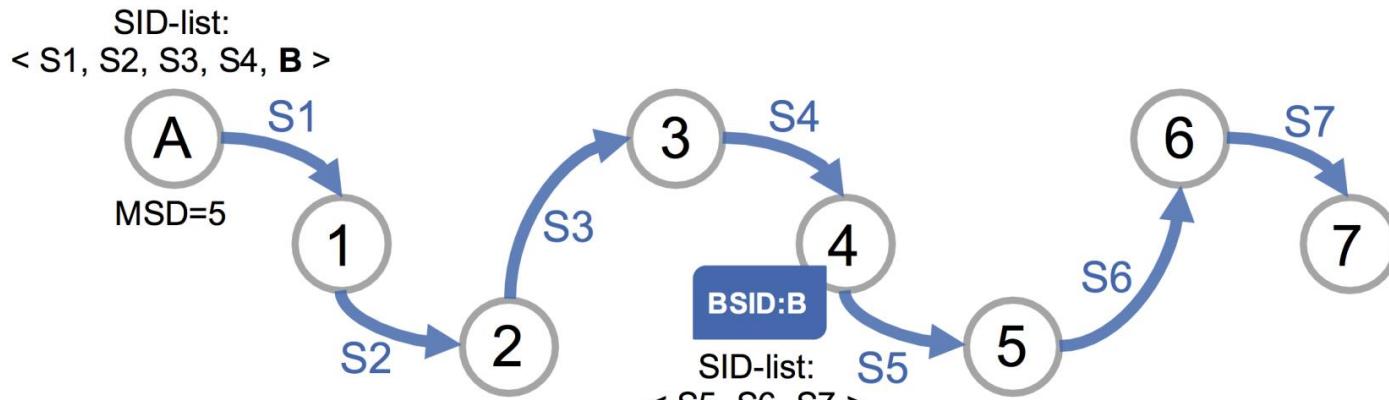
All Nodes & consistents

```
segment-routing
traffic-eng
policy VIA-PLANE-GREEN
  color 20 end-point ipv4 1.1.1.3
  candidate-paths
    preference 100
    dynamic
    metric
    type igrp
  !
  constraints
    affinity
    exclude-any
      name BLUE
  !
  policy VIA-PLANE-BLUE
    color 30 end-point ipv4 1.1.1.3
    candidate-paths
      preference 100
      dynamic
      metric
      type igrp
    !
    constraints
      affinity
      exclude-any
        name GREEN
```

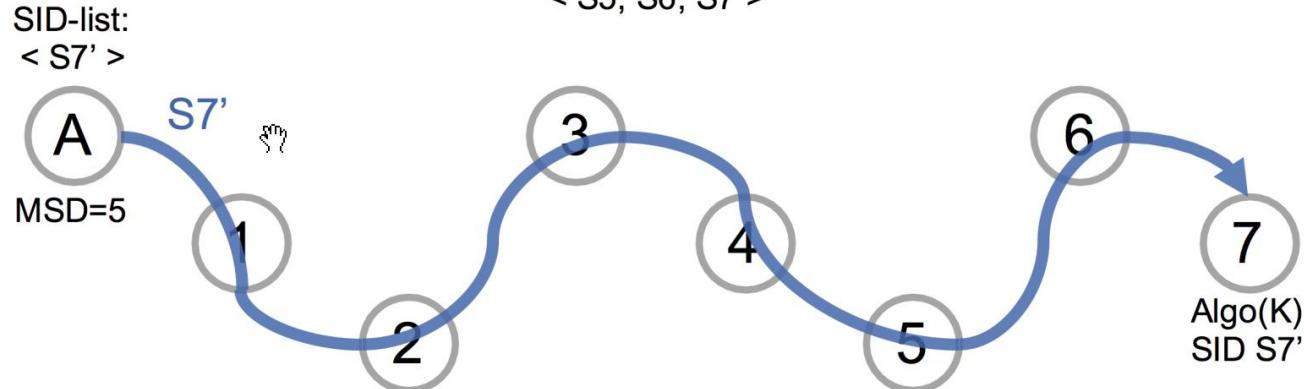
Head End - R2

# How many SIDs ?

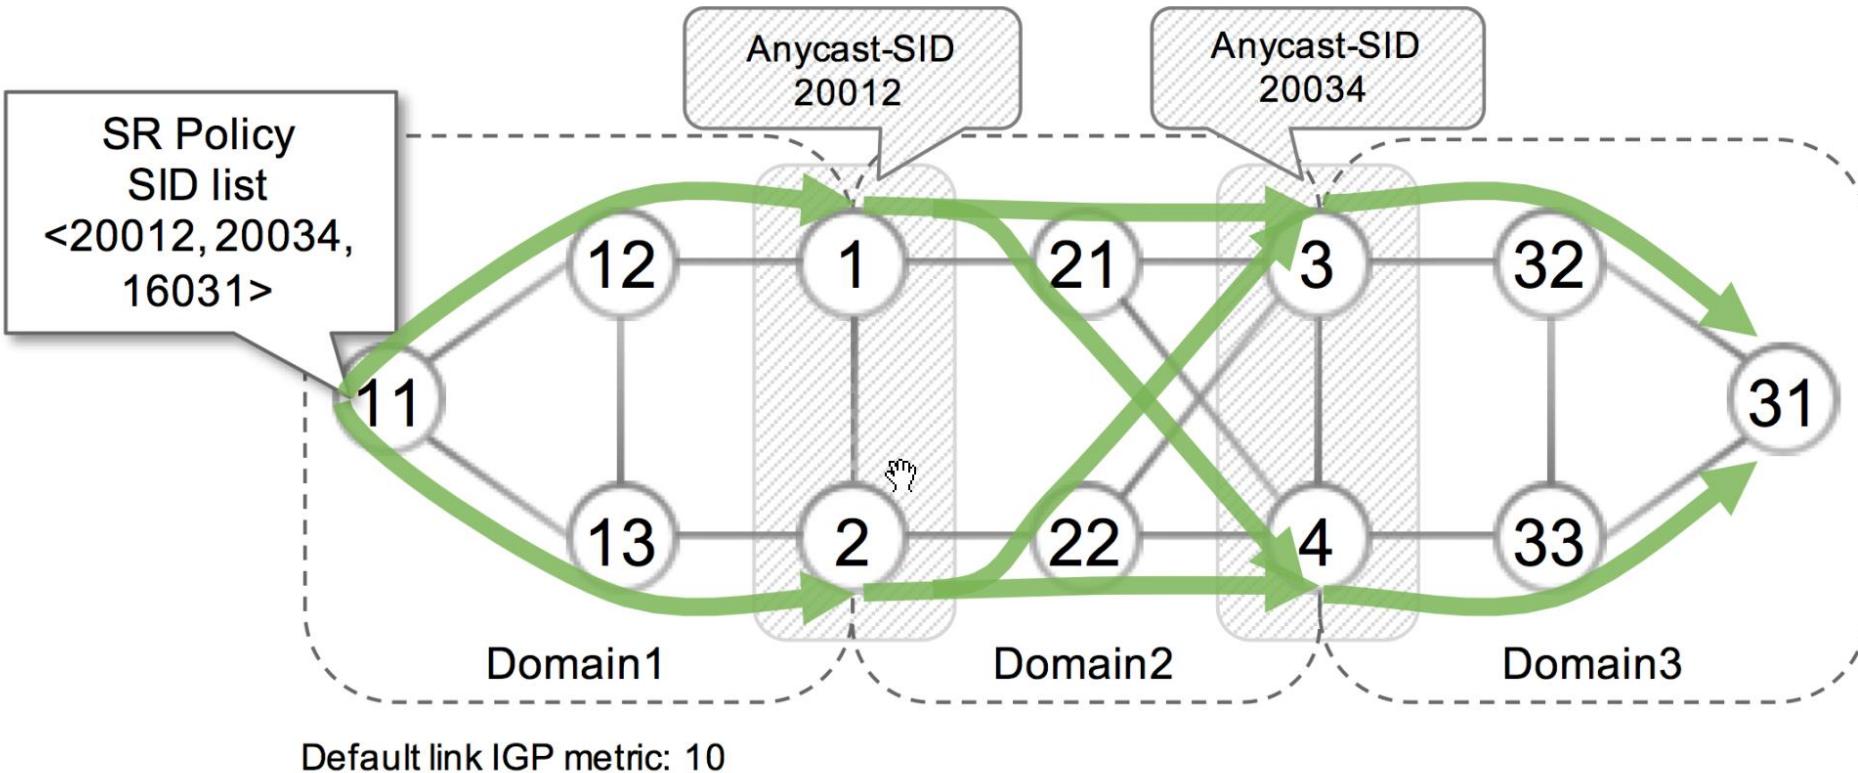
(a)



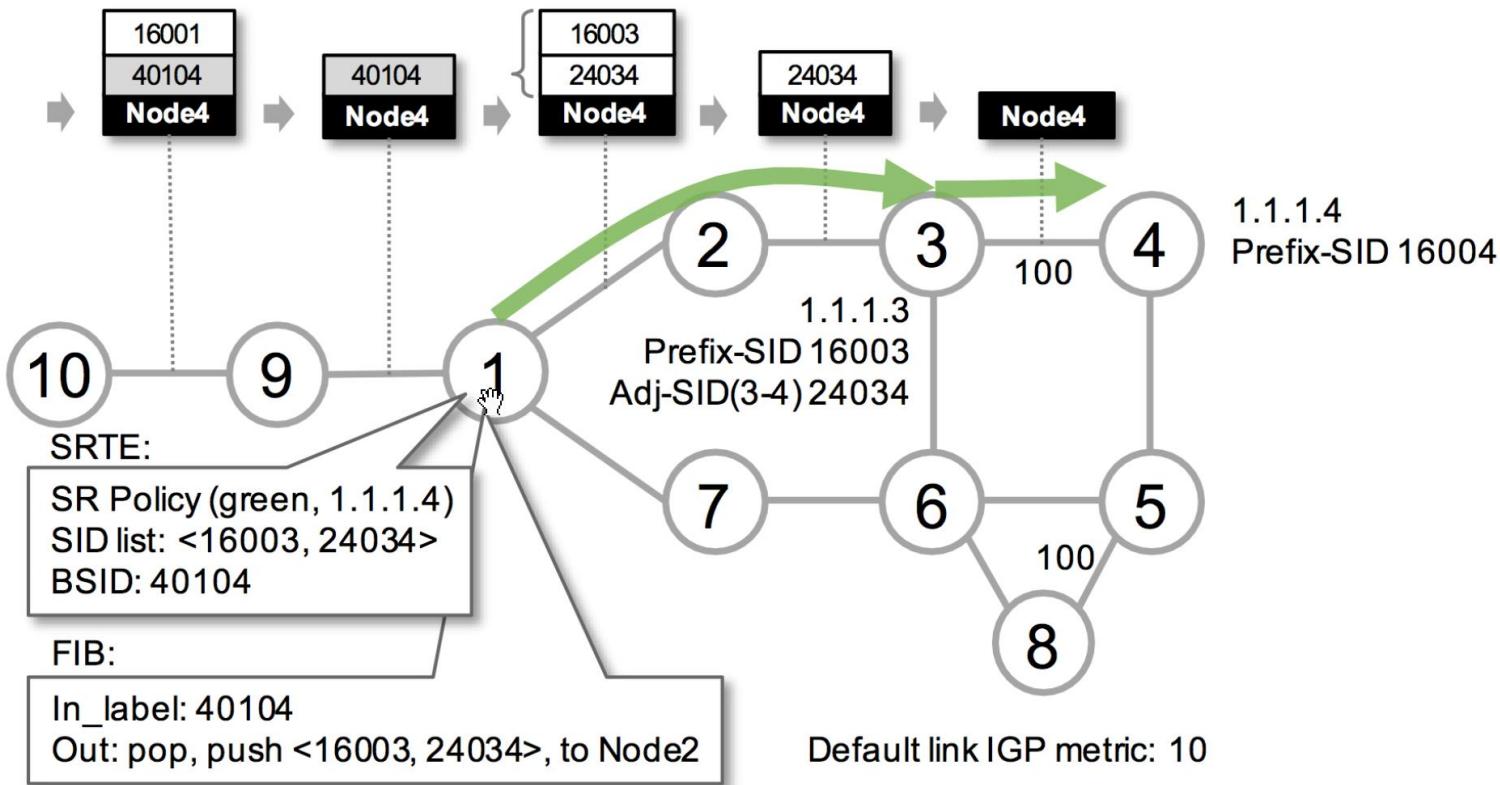
(b)



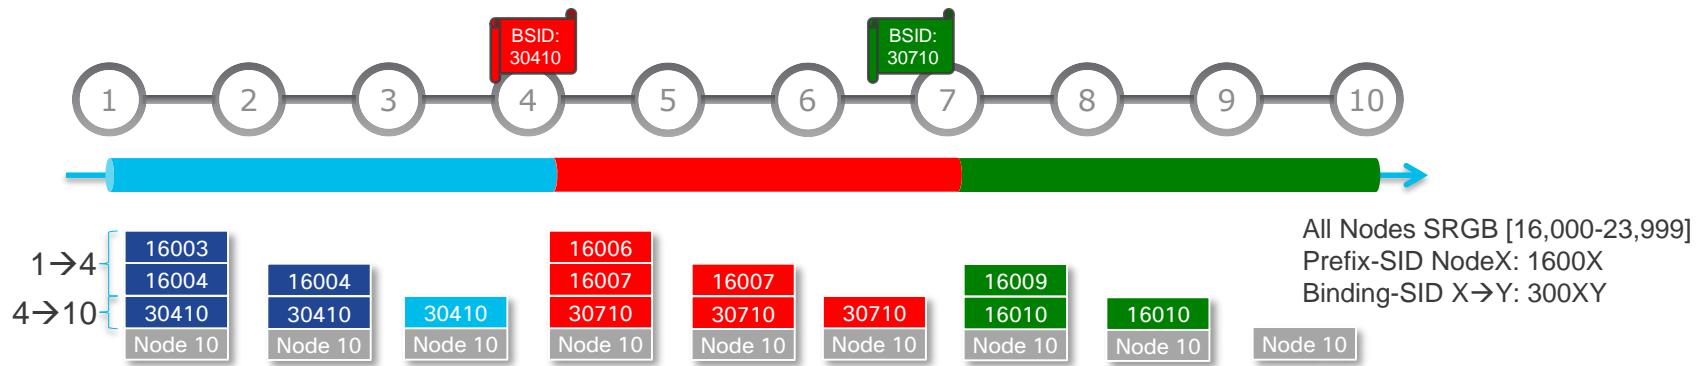
# How many SIDs ?



# Binding SID

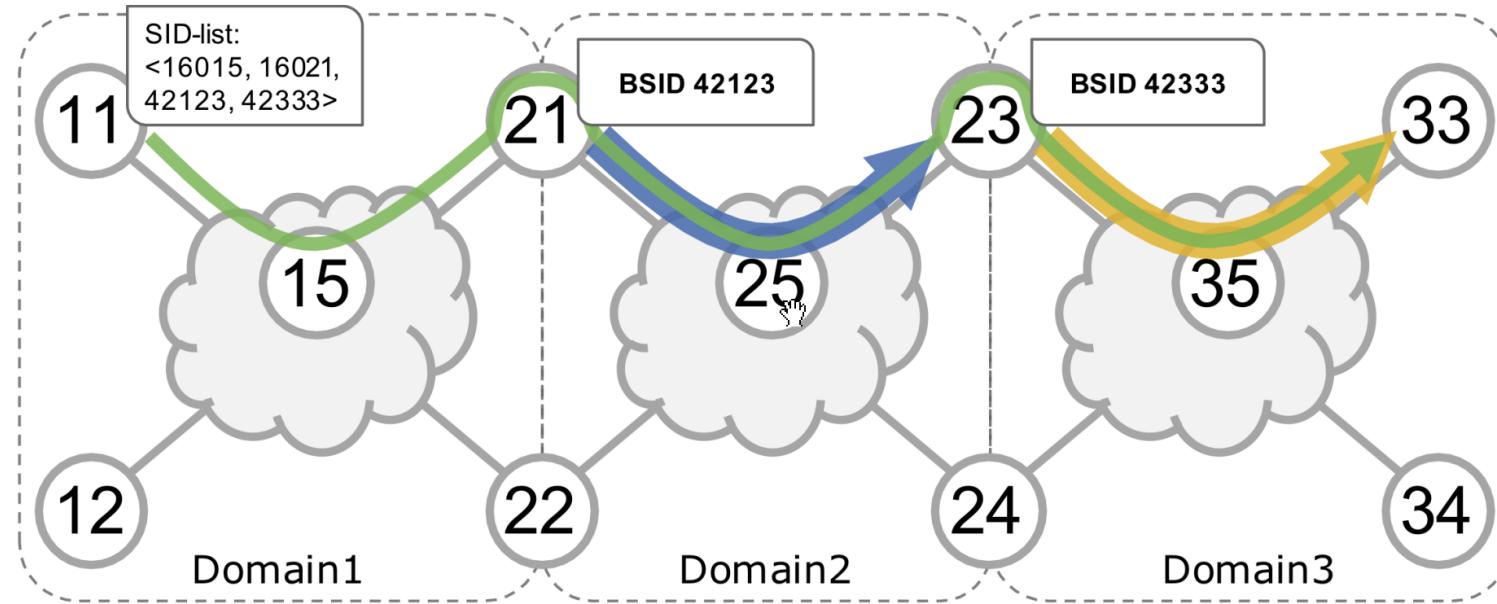


# Binding-SID – Stitching – Illustration



- Assume Node1 can't push 8 labels to go to Node10
  - “compress” label stack by **stitching** SRTE Policies:
    - Node1 pushes:
      - 2 labels to go to Node4
      - Binding-SID to go to Node10
    - Node4 pops Binding-SID and pushes:
      - 2 labels to go to Node7
      - Binding-SID to go to Node10
    - Node7 pops Binding-SID and pushes 2 labels to go to Node10

# Binding SID



11→21	{	16015
21→23	{	16021
23→33	{	42123
		42333
		Node33

16021
42123
42333
Node33

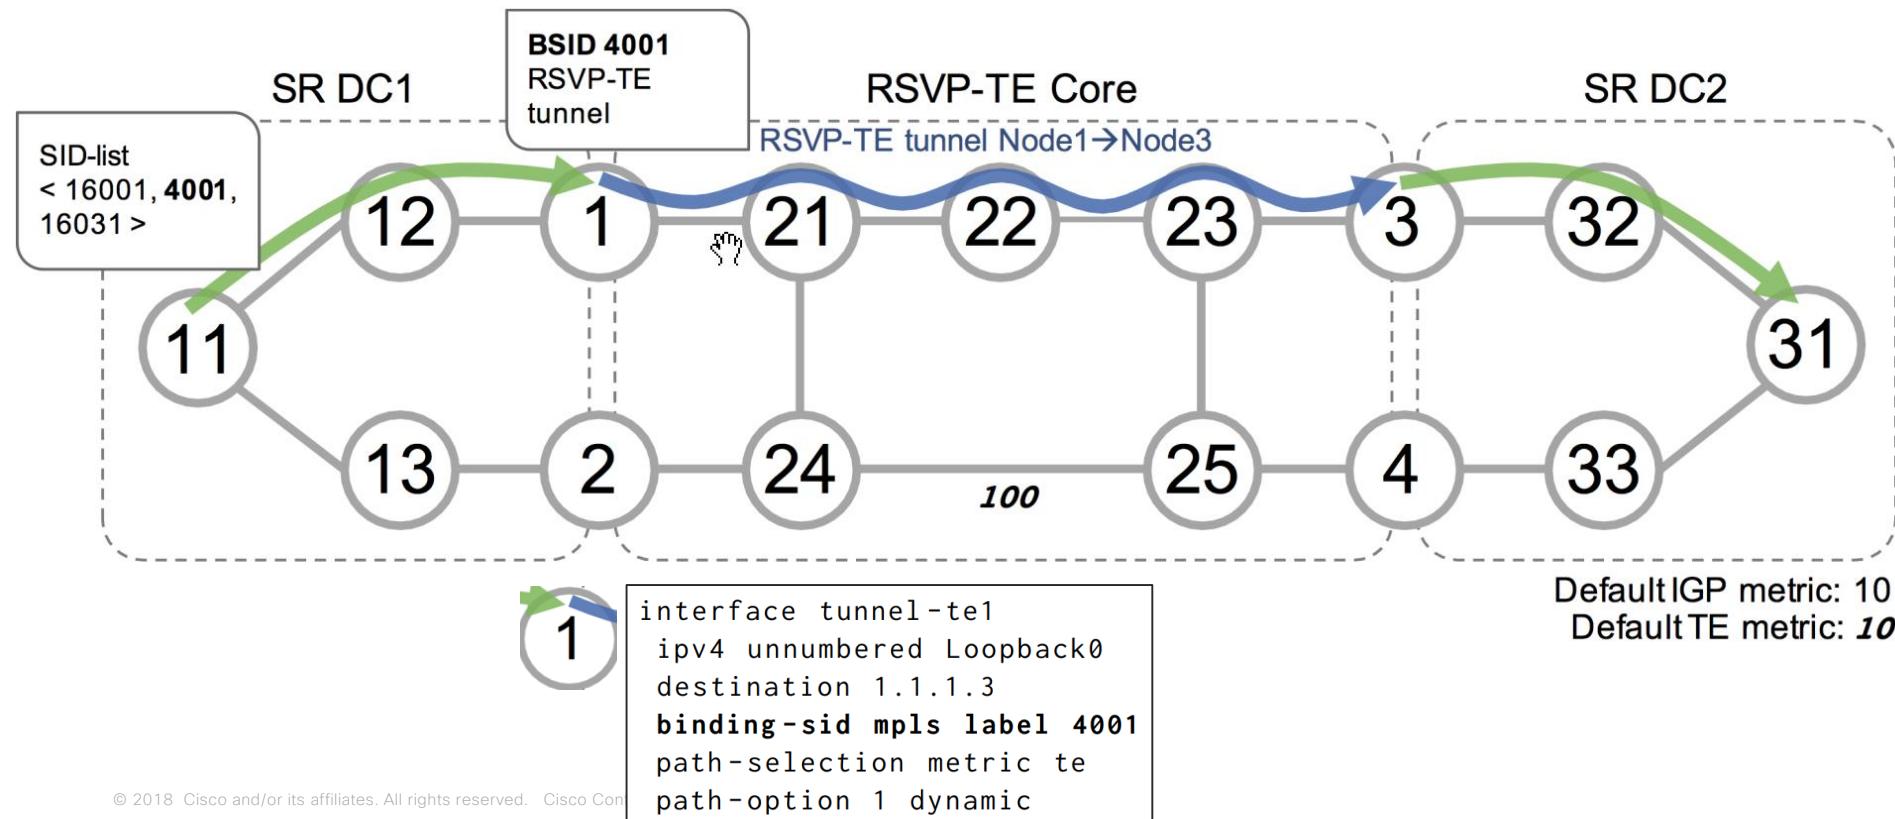
16025
16023
42333
Node33

16023
42333
Node33

16035
16033
Node33

16033
Node33

# Binding SID for RSVP-TE Tunnel





# Segment Routing Automated Steering

Intent-based

Automated  
Steering

L3VPN

BGP-based

BGP

BGP

BGP

2.2.2.0/24, vrf Acme  
**nexthop 1.1.1.4**  
**color green (30)**

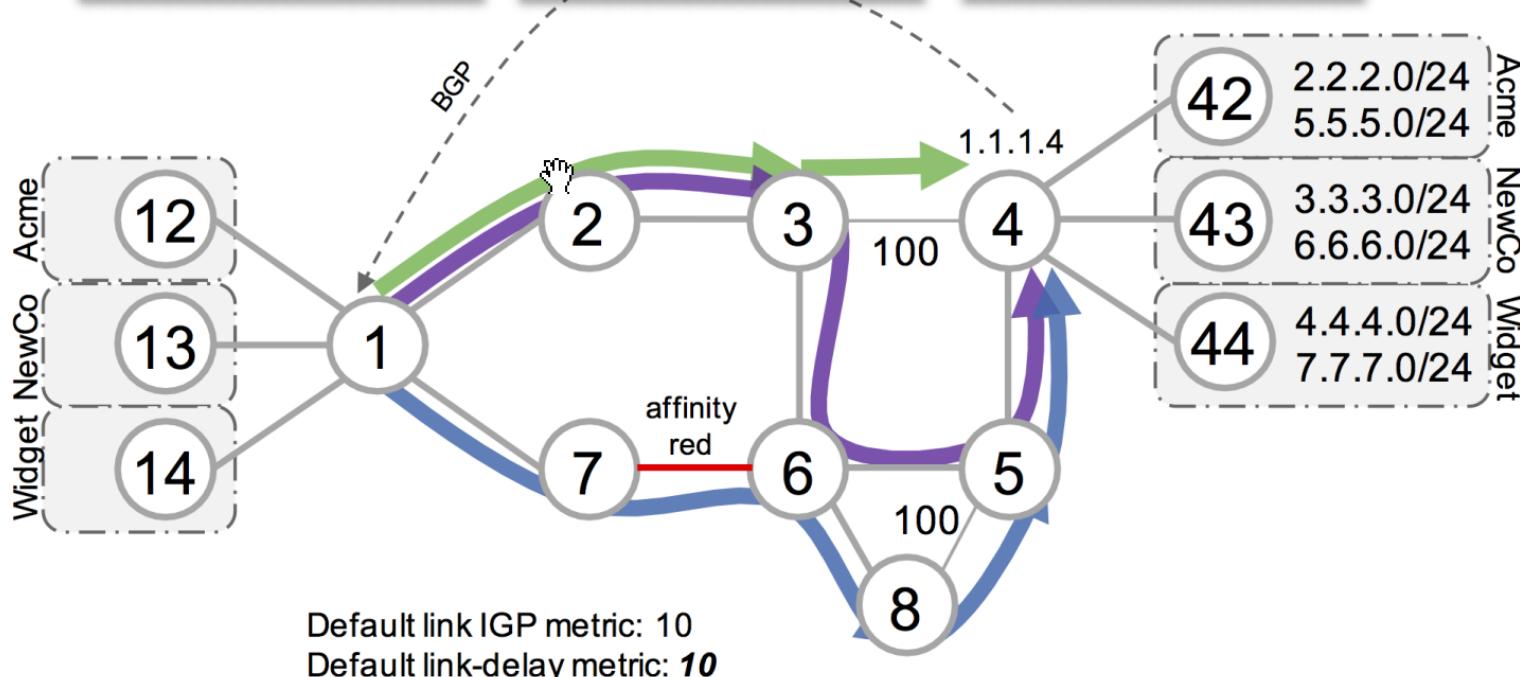
5.5.5.0/24, vrf Acme  
**nexthop 1.1.1.4**  
**color green (30)**

3.3.3.0/24, vrf NewCo  
**nexthop 1.1.1.4**  
**color blue (20)**

6.6.6.0/24, vrf NewCo  
**nexthop 1.1.1.4**  
**color green (30)**

4.4.4.0/24, vrf Widget  
**nexthop 1.1.1.4**  
**color purple (40)**

7.7.7.0/24, vrf Widget  
**nexthop 1.1.1.4**  
**(no color)**



# Intent-based Automated Steering L3VPN

1

```
segment - routing
traffic -eng
affinity -map
name red bit - position 0
!
segment - list name BLUE_PATH
index 10 mpls label 16008
index 20 mpls label 24085
index 30 mpls label 16004
!
policy BLUE
color 20 end - point ipv4 1.1.1.4
candidate - paths
preference 100
explicit segment - list BLUE_PATH
!
```

1

```
policy GREEN
color 30 end - point ipv4 1.1.1.4
candidate - paths
preference 100
dynamic
metric
type delay
!
policy PURPLE
color 40 end - point ipv4 1.1.1.4
candidate - paths
preference 100
dynamic
metric
type igrp
constraints
affinity
exclude -any
name red
```

4

```
extcommunity -set opaque BLUE
20
end -set
!
extcommunity -set opaque GREEN
30
end -set
!
extcommunity -set opaque PURPLE
40
end -set

route - policy VRF - COLOR
if destination in (3.3.3.0/24) then
set extcommunity color BLUE
endif
if destination in (6.6.6.0/24) then
set extcommunity color GREEN
endif
pass
end - policy
!
route - policy GLOBAL - COLOR
if destination in (8.8.8.0/24) then
set extcommunity color PURPLE
endif
pass
end - policy
```

4

```
router bgp 1
bgp router -id 1.1.1.4
address - family ipv4 unicast
address - family vpngv4 unicast
!
neighbor 1.1.1.1
remote -as 1
update - source Loopback0
address - family ipv4 unicast
route - policy GLOBAL - COLOR out
!
address - family vpngv4 unicast
!
vrf NewCo
rd auto
address - family ipv4 unicast
!
neighbor 99.4.43.43
remote -as 2
description to CE43
address - family ipv4 unicast
route - policy VRF - COLOR out
```

Intent-based

Automated  
Steering

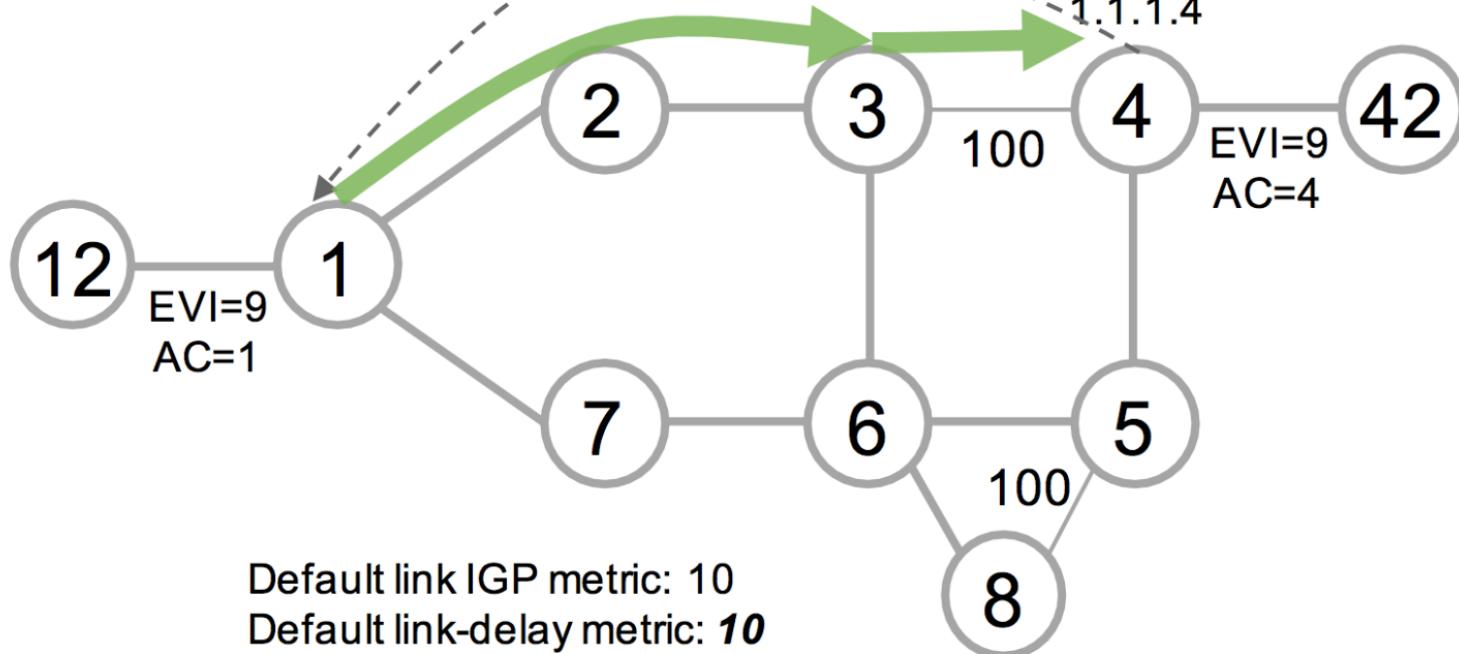
EVPN

BGP-based

BGP

EVI 9, AC 4

**nexthop 1.1.1.4  
color green (30)**



# Intent-based Automated Steering EVPN VPWS

1

```
segment - routing
traffic -eng
policy GREEN
color 30 end - point ipv4 1.1.1.4
candidate - paths
preference 100
dynamic mpls
metric
type delay
```

```
interface GigabitEthernet0 /0/0/0.100
l2transport
encapsulation dot1q 100
!
l2vpn
xconnect group evpn - vpws
p2p EVI9
interface GigabitEthernet0 /0/0/0.100
neighbor evpn evi 9 target 4 source 1
```

1

```
router bgp 1
bgp router -id 1.1.1.1
address - family l2vpn evpn
!
neighbor 1.1.1.4
remote -as 1
update - source Loopback0
!
address - family l2vpn evpn
```

4

```
extcommunity -set opaque GREEN
# color green identifies low - delay
30
end -set
!
route - policy evpn_vpws_policy
if rd in (1.1.1.4:9) then
set extcommunity color GREEN
endif
end - policy
!
router bgp 1
bgp router -id 1.1.1.4
address - family l2vpn evpn
!
neighbor 1.1.1.1
remote -as 1
update - source Loopback0
!
address - family l2vpn evpn
route - policy evpn_vpws_policy out
```

4

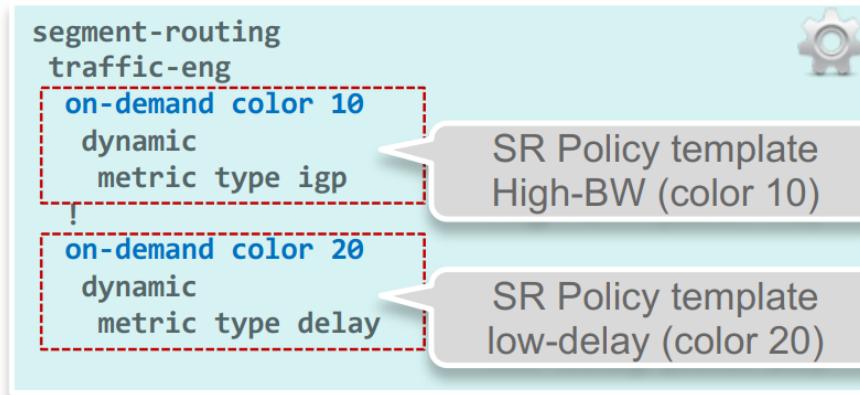
```
interface GigabitEthernet0 /0/0/0.100
l2transport
encapsulation dot1q 100
!
l2vpn
xconnect group evpn - vpws
p2p EVI9
interface GigabitEthernet0 /0/0/0.100
neighbor evpn evi 9 target 1 source 4
```



# Segment Routing On Demand NextHop

# On-demand SR Policy

- Configure an SR Policy template for each color for which on-demand SR Policy instantiation is desired
- An example with two color templates configured:
  - color 10 for high bandwidth (optimize IGP metric)
  - color 20 for low-delay (optimize link-delay metric)



# On-Demand Nexthop reachability

5

```
router bgp 1
neighbor 1.1.1.10
address-family vpnv4 unicast
!
```

```
segment-routing
traffic-eng
```

```
on-demand color 10
dynamic
pcep
metric
type igp
```

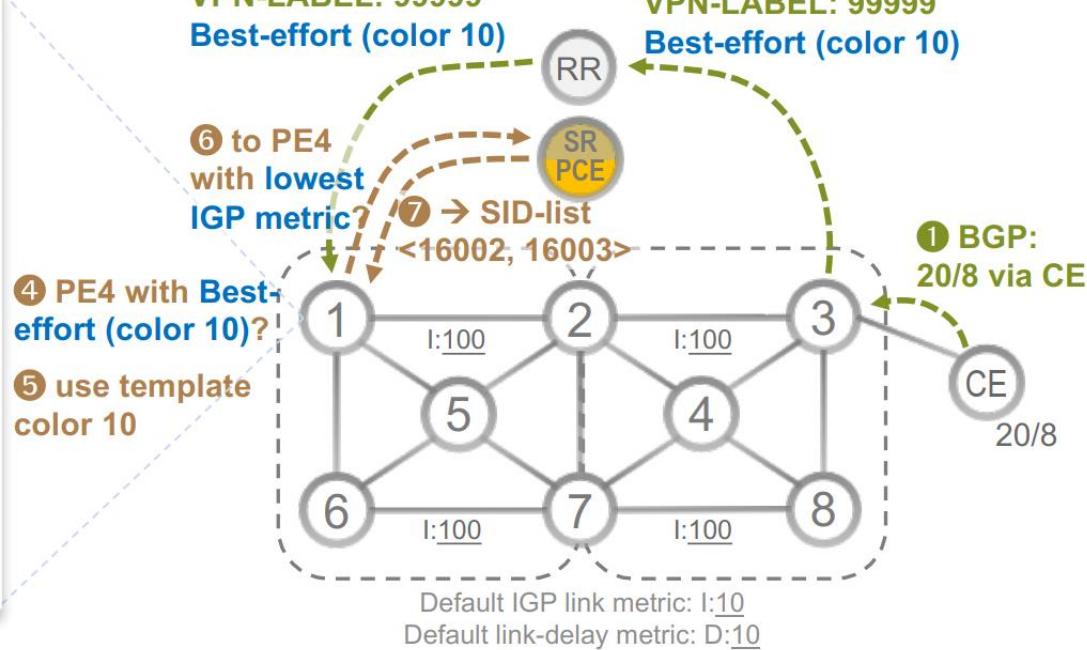
```
on-demand color 20
dynamic
pcep
metric
type delay
```

SR Policy template  
Best-effort (color 10)

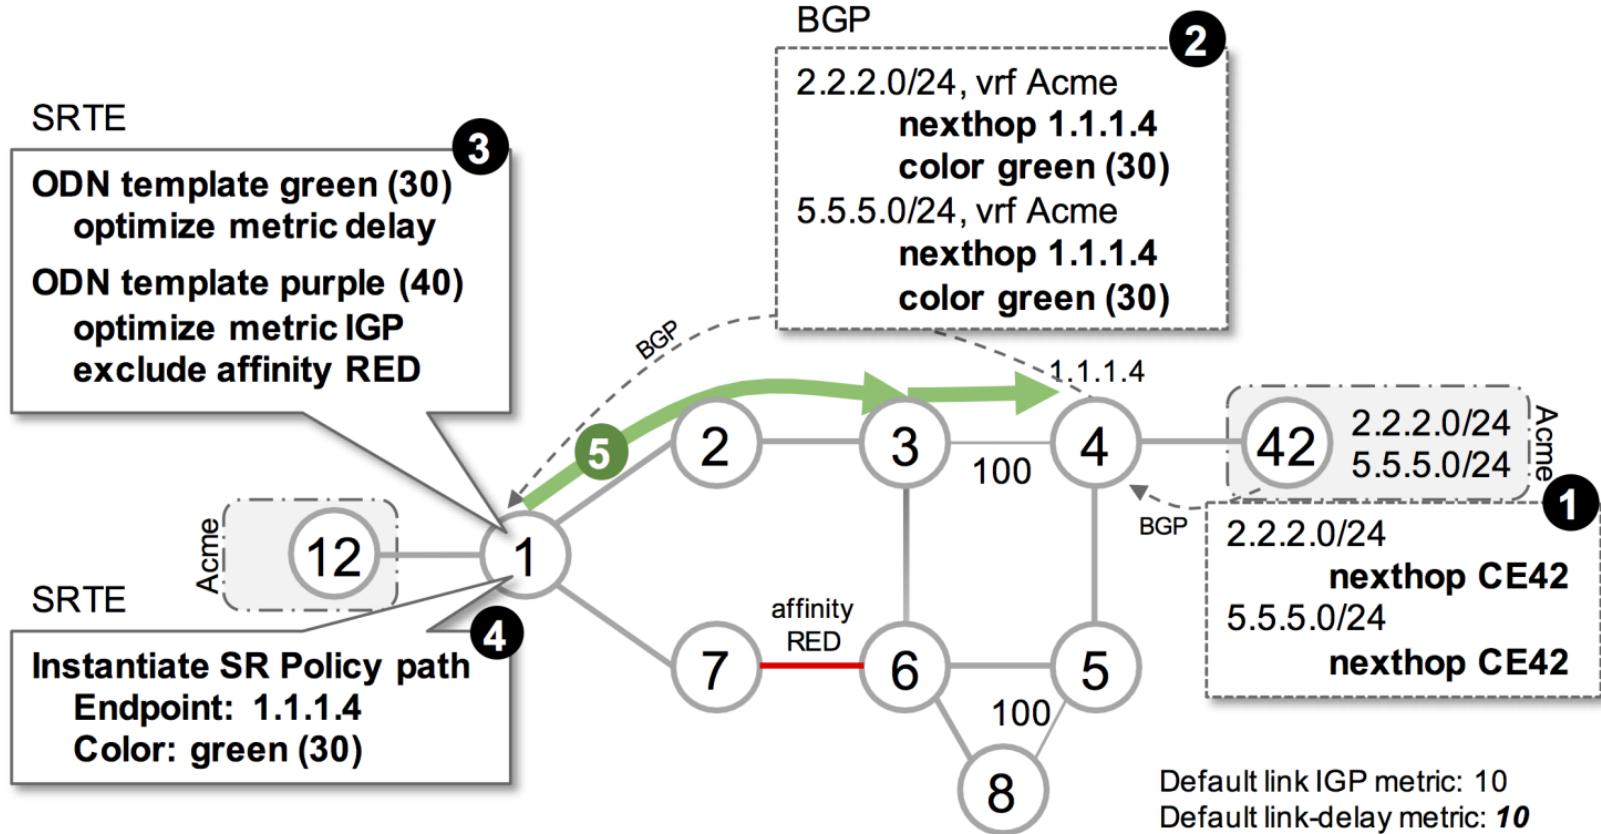


③ BGP: 20/8 via PE3
VPN-LABEL: 99999
Best-effort (color 10)

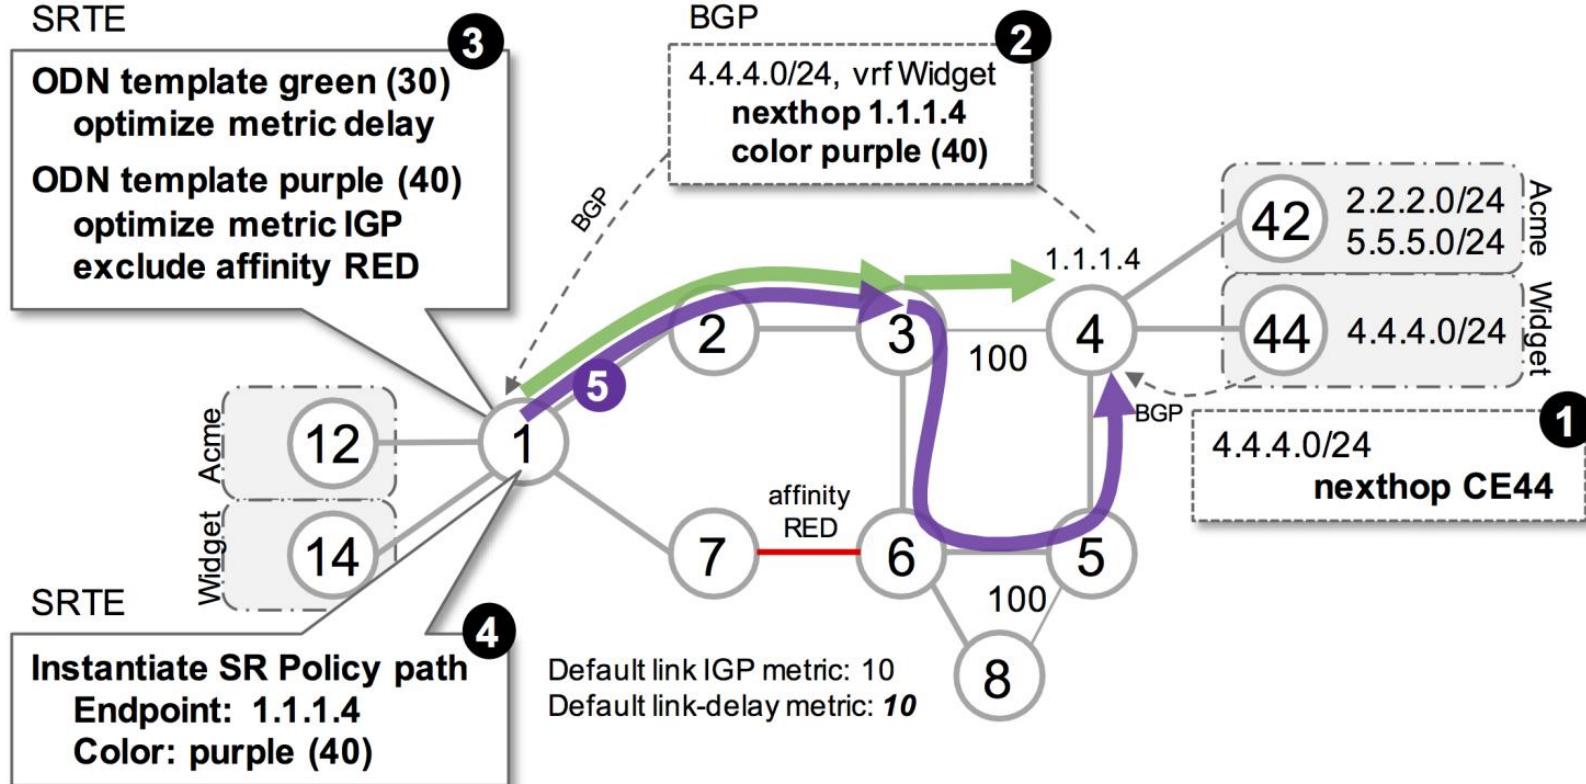
② BGP: 20/8 via PE3
VPN-LABEL: 99999
Best-effort (color 10)



# ODN On Demand NextHop



# ODN On Demand NextHop



# Intent-based ODN L3VPN

1

```
segment - routing
traffic -eng
affinity -map
    name RED bit - position 1
!
!! green ODN template
on - demand color 30
    candidate - paths
        preference 100
        dynamic
        metric
        type delay
!
!! purple ODN template
on - demand color 40
    candidate - paths
        preference 100
        constraints
        affinity
        !! RED is defined in affinity -map above
            exclude -any
            name RED
        dynamic
        metric
        type igrp
```

4

```
extcommunity -set opaque BLUE
20
end -set
!
extcommunity -set opaque GREEN
30
end -set
!
extcommunity -set opaque PURPLE
40
end -set

route - policy VRF - COLOR
if destination in (3.3.3.0/24) then
set extcommunity color BLUE
endif
if destination in (6.6.6.0/24) then
set extcommunity color GREEN
endif
pass
end - policy
!
route - policy GLOBAL - COLOR
if destination in (8.8.8.0/24) then
set extcommunity color PURPLE
endif
pass
end - policy
```

4

```
router bgp 1
bgp router -id 1.1.1.4
address - family ipv4 unicast
address - family vpngv4 unicast
!
neighbor 1.1.1.1
remote -as 1
update - source Loopback0
address - family ipv4 unicast
route - policy GLOBAL - COLOR out
!
address - family vpngv4 unicast
!
vrf NewCo
rd auto
address - family ipv4 unicast
!
neighbor 99.4.43.43
remote -as 2
description to CE43
address - family ipv4 unicast
route - policy VRF - COLOR out
```

# Automated Steering

1

```
segment - routing
traffic -eng
affinity -map
name red bit - position 0
!
segment - list name BLUE_PATH
index 10 mpls label 16008
index 20 mpls label 24085
index 30 mpls label 16004
!
policy BLUE
color 20 end - point ipv4 1.1.1.4
candidate - paths
preference 100
explicit segment - list BLUE_PATH
!
```

1

```
policy GREEN
color 30 end - point ipv4 1.1.1.4
candidate - paths
    preference 100
    dynamic
    metric
    type delay
!
policy PURPLE
color 40 end - point ipv4 1.1.1.4
candidate - paths
    preference 100
    dynamic
    metric
    type igp
constraints
    affinity
    exclude -any
    name red
```

vs

# ODN

1

```
segment - routing
traffic -eng
affinity -map
name RED bit - position 1
!
```

**!! green ODN template**  
**on - demand color 30**

```
candidate - paths
preference 100
dynamic
metric
type delay
!
```

**!! purple ODN template**  
**on - demand color 40**

```
candidate - paths
preference 100
constraints
affinity
!! RED is defined in affinity -map above
exclude -any
name RED
dynamic
metric
type igp
```

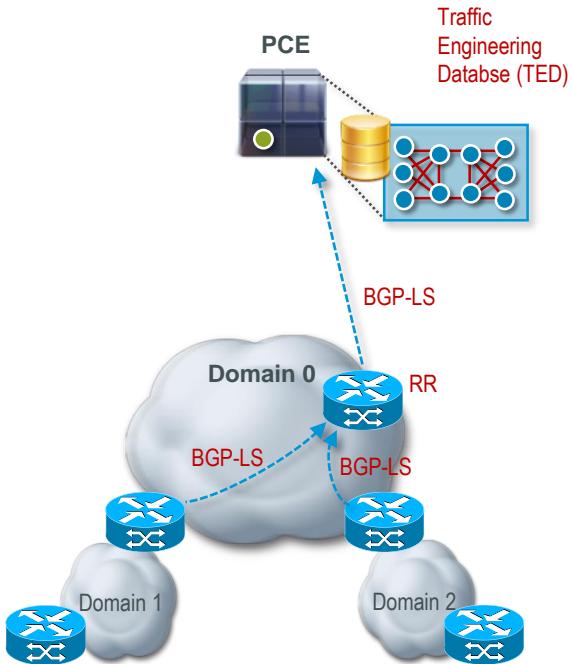


# SR PCE - BGP LS + PCEP

## Use-Case Seamless MPLS - BGP LU

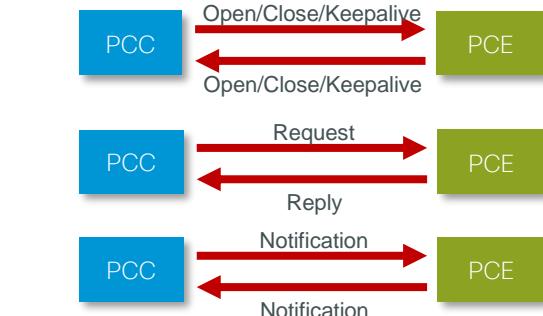
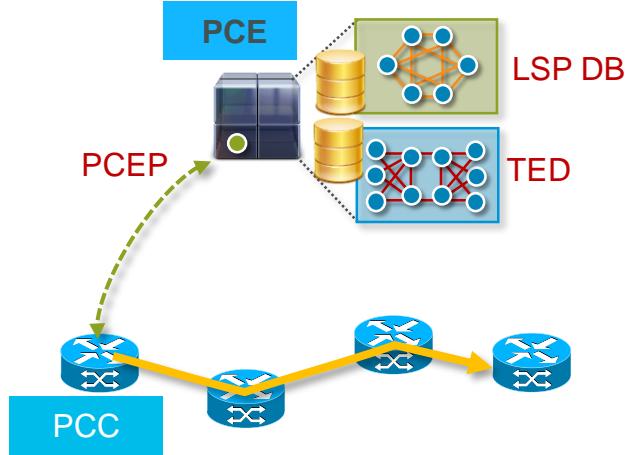
# BGP-LS Overview

- Optimal Path Computation for Multi-area TE
- Solution is BGP, not IGP.
- BGP-LS is an address-family
  - afi=16388, safi=71
- Defined to carry IGP link-state database via BGP
  - Supports both IS-IS and OSPF
  - Delivers topology information to outside agents



# PCEP Architectural Introduction

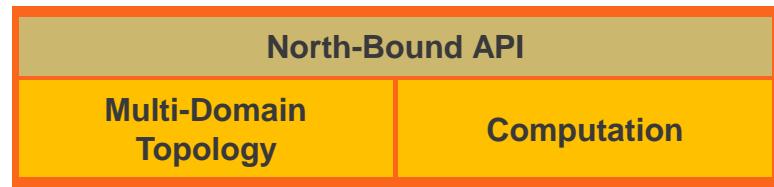
- Path computation
- Large, multi-domain and multi-layer networks
- Path computation element (PCE)
  - Computes network paths (topology, paths, etc.)
  - Stores TE topology database (synchronized with network)
  - May initiate path creation
  - Stateful – stores path database included resources used (synchronized with network)
- Path computation client (PCC)
  - May send path computation requests to PCE
  - May send path state updates to PCE
- Used between head-end router (PCC) and PCE to:
  - Request/receive path from PCE subject to constraints
  - State synchronization between PCE and router
  - Hybrid CSPF



# Segment Routing - Path Computation Element (SR-PCE)

- An IOS XR-powered Stateful Path Computation Element (PCE)
- Multi-Domain topology Collection
  - Real-time reactive feed
- Computation
  - Native SR-TE algorithms backed by extensive scientific research

**SR PCE**

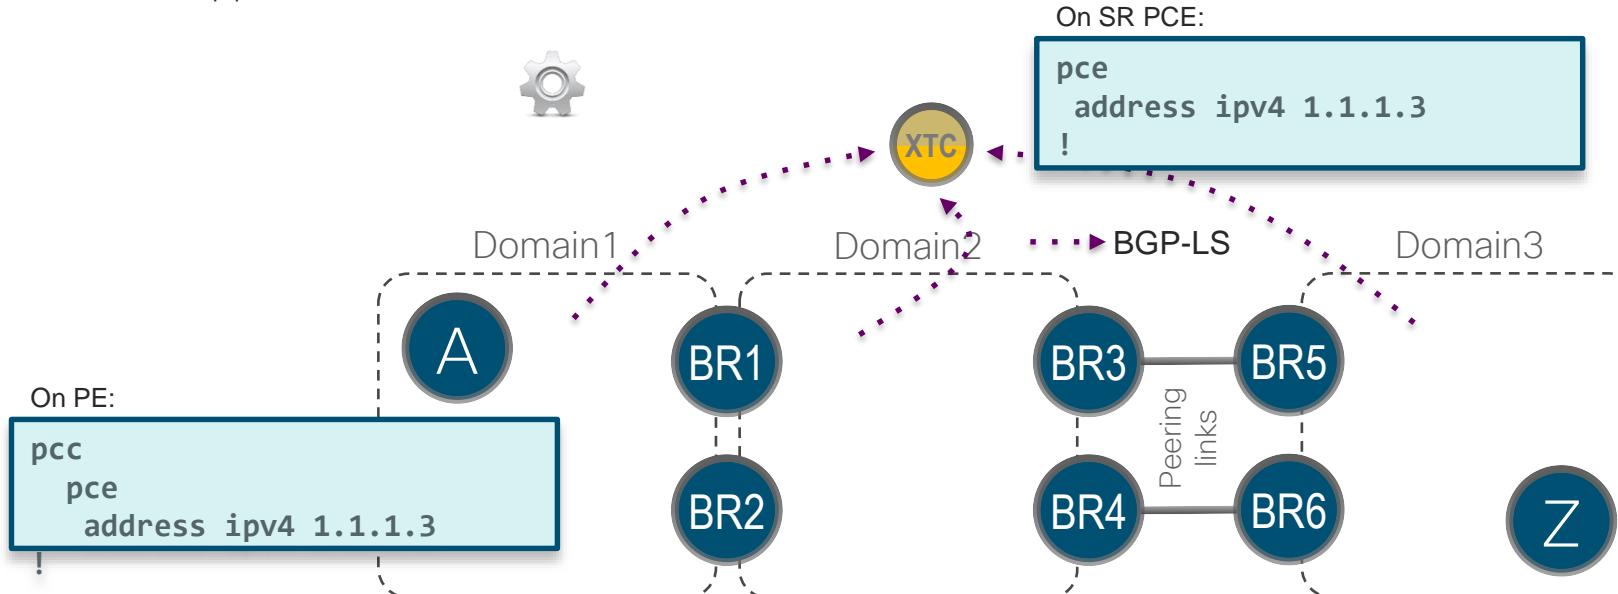


↑  
↑  
“Collection”  
BGP-LS  
ISIS / OSPF

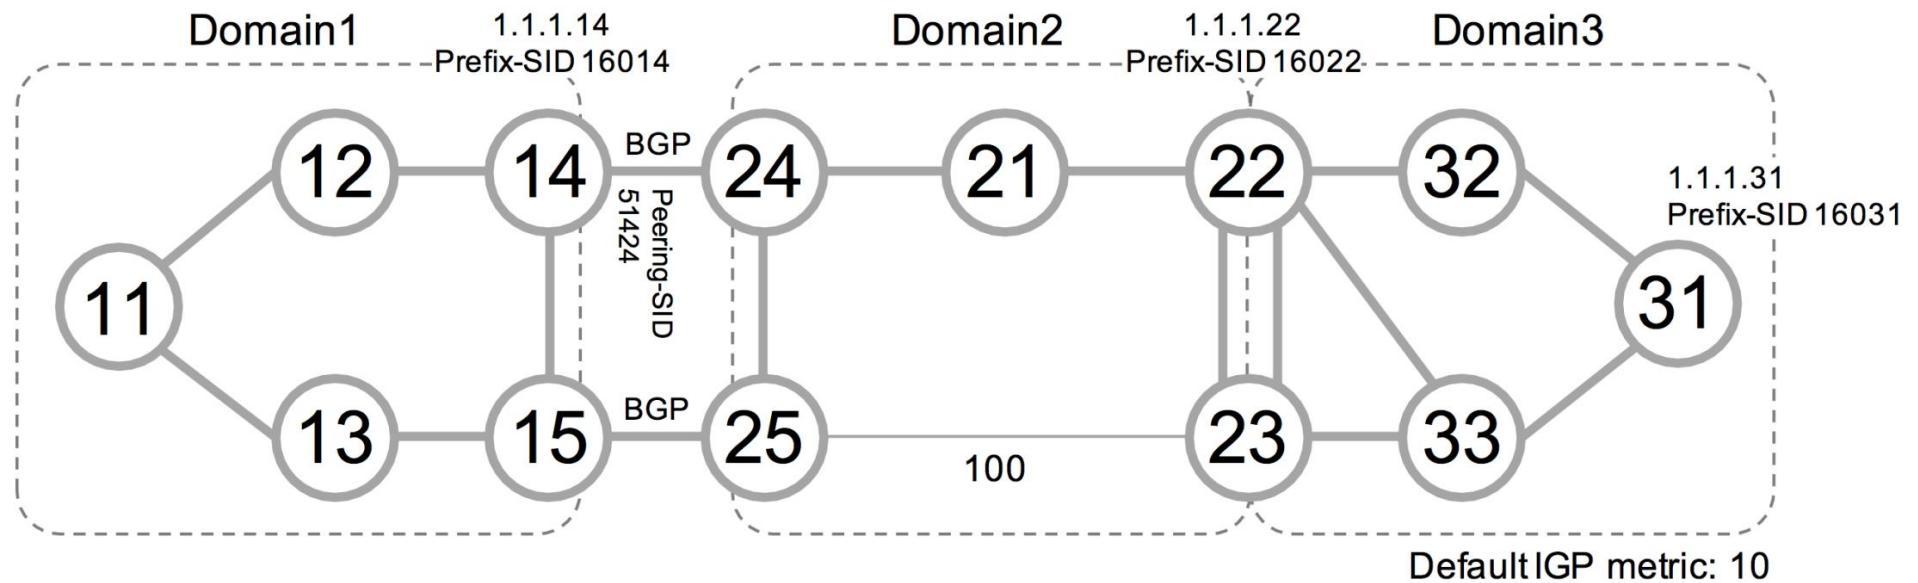
↑  
→  
“Deployment”  
PCEP/BGP-SRTE

# SR PCE Controller

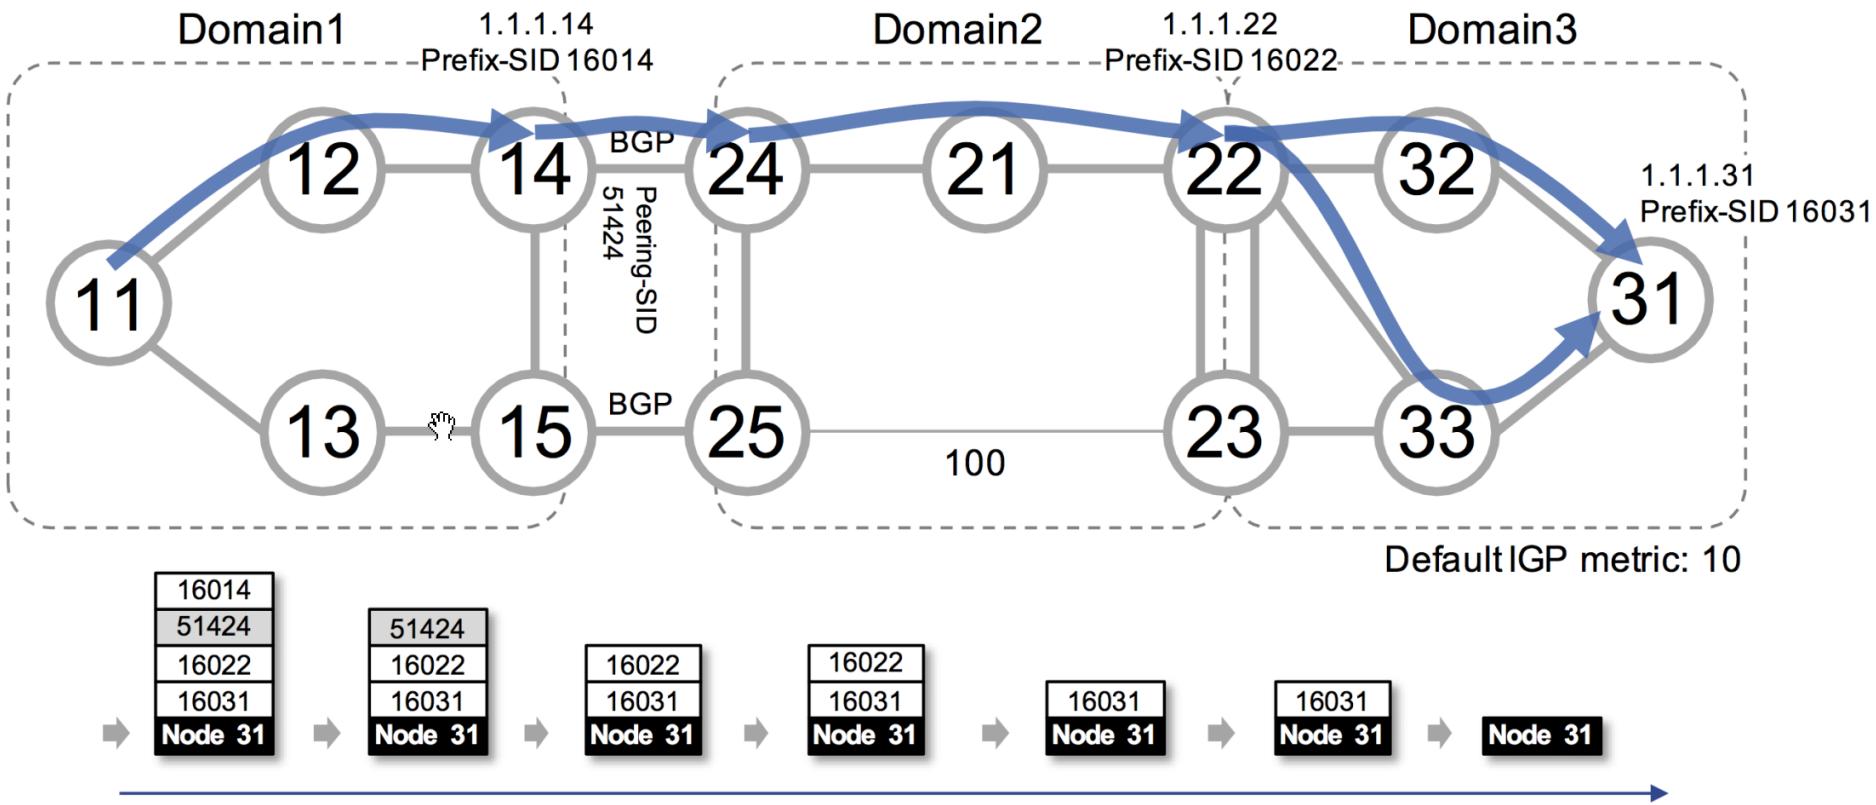
- SR PCE is an IOS XR multi-domain stateful SR Path Computation Element (PCE)
  - Fundamentally Distributed (RR-like Deployment)
  - Supports RSVP-TE



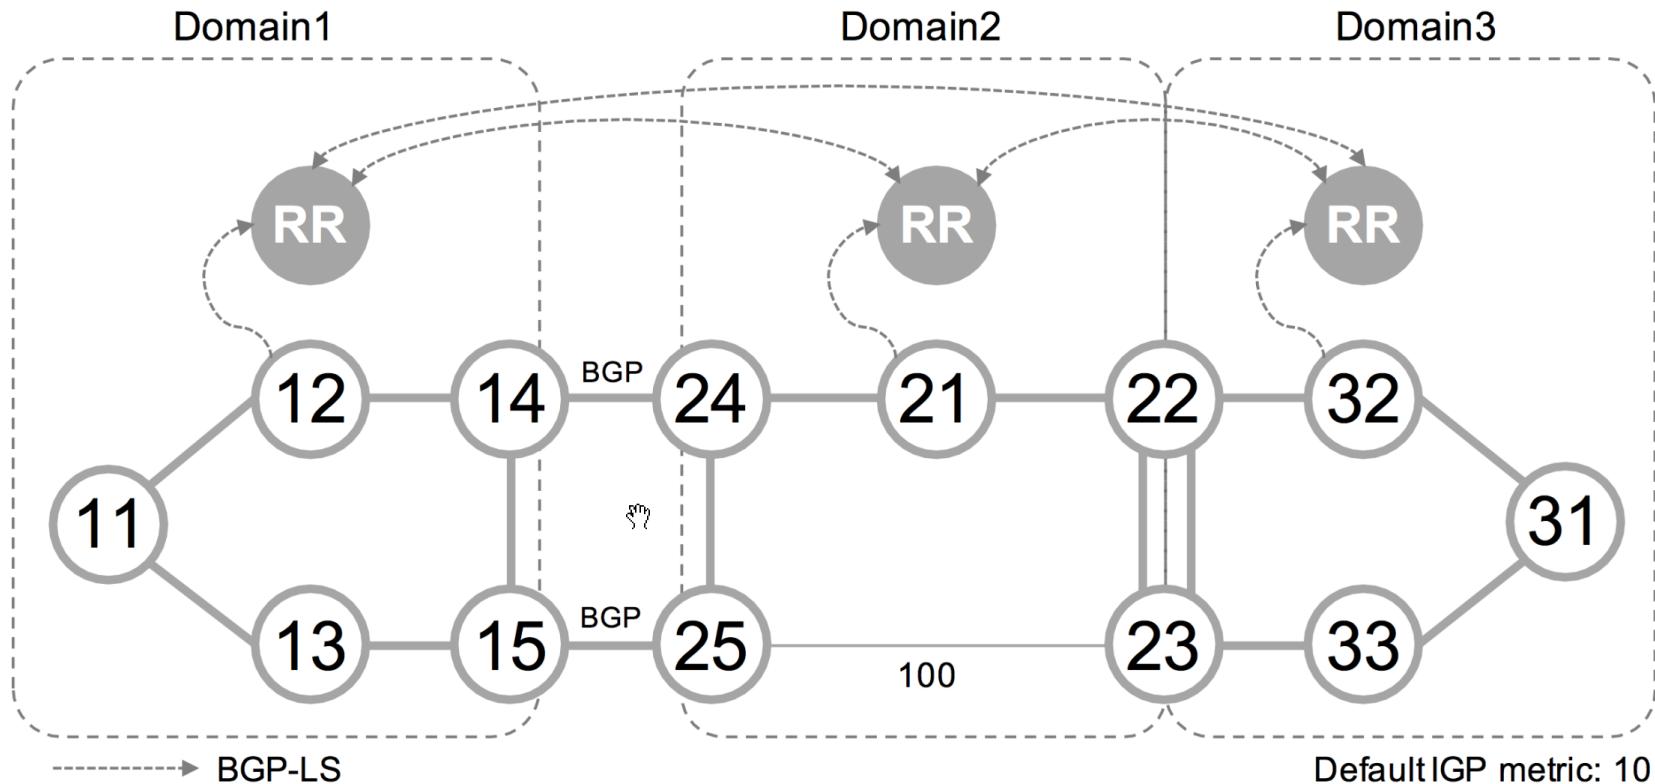
# Inter-domain : Seamless MPLS



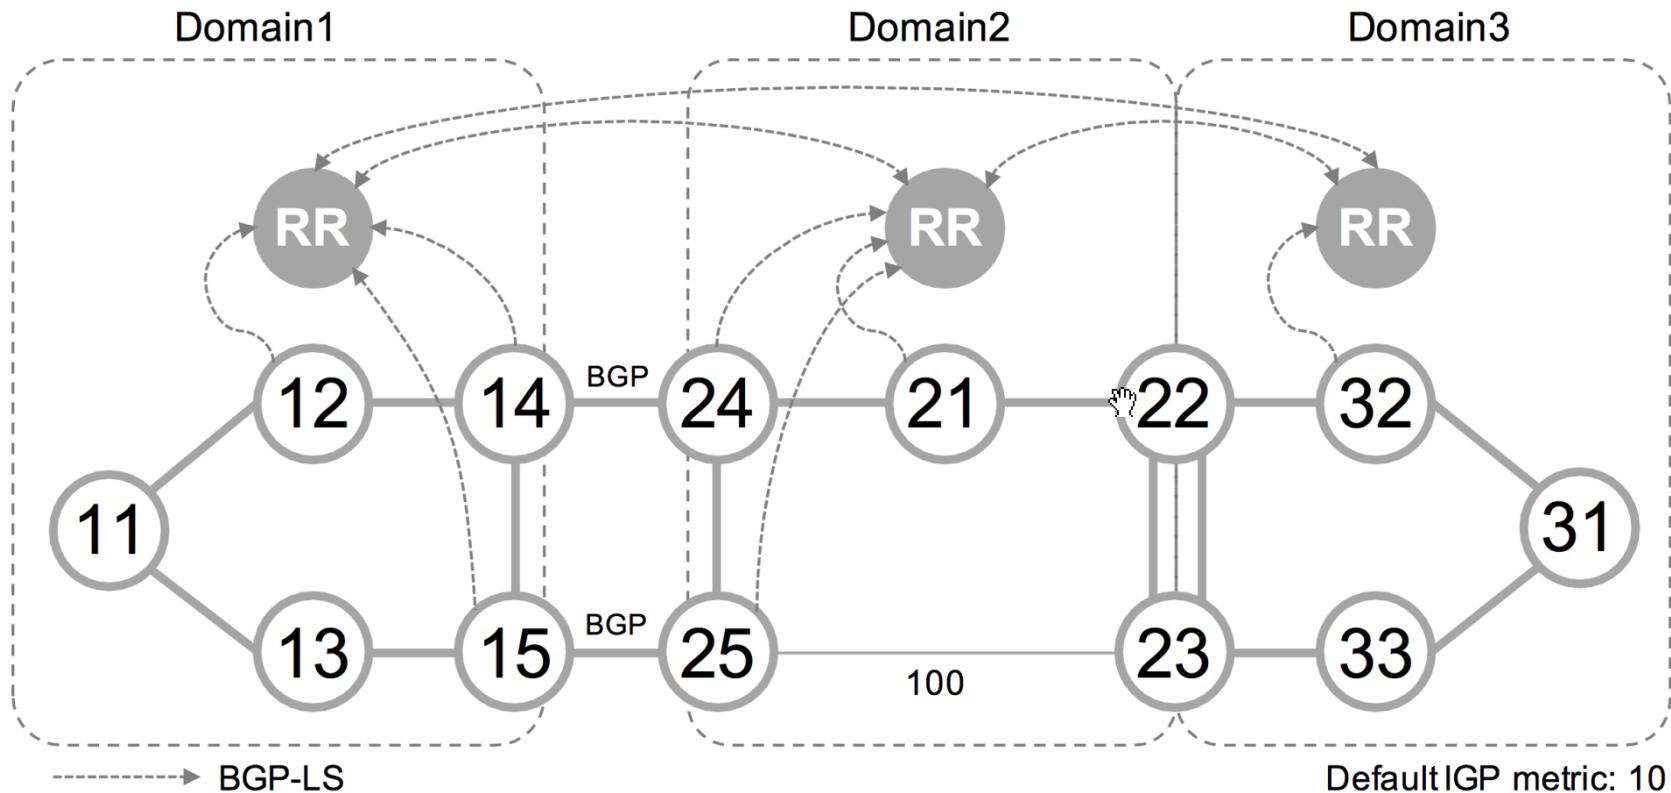
# End-to-End Inter-domain Explicit Path



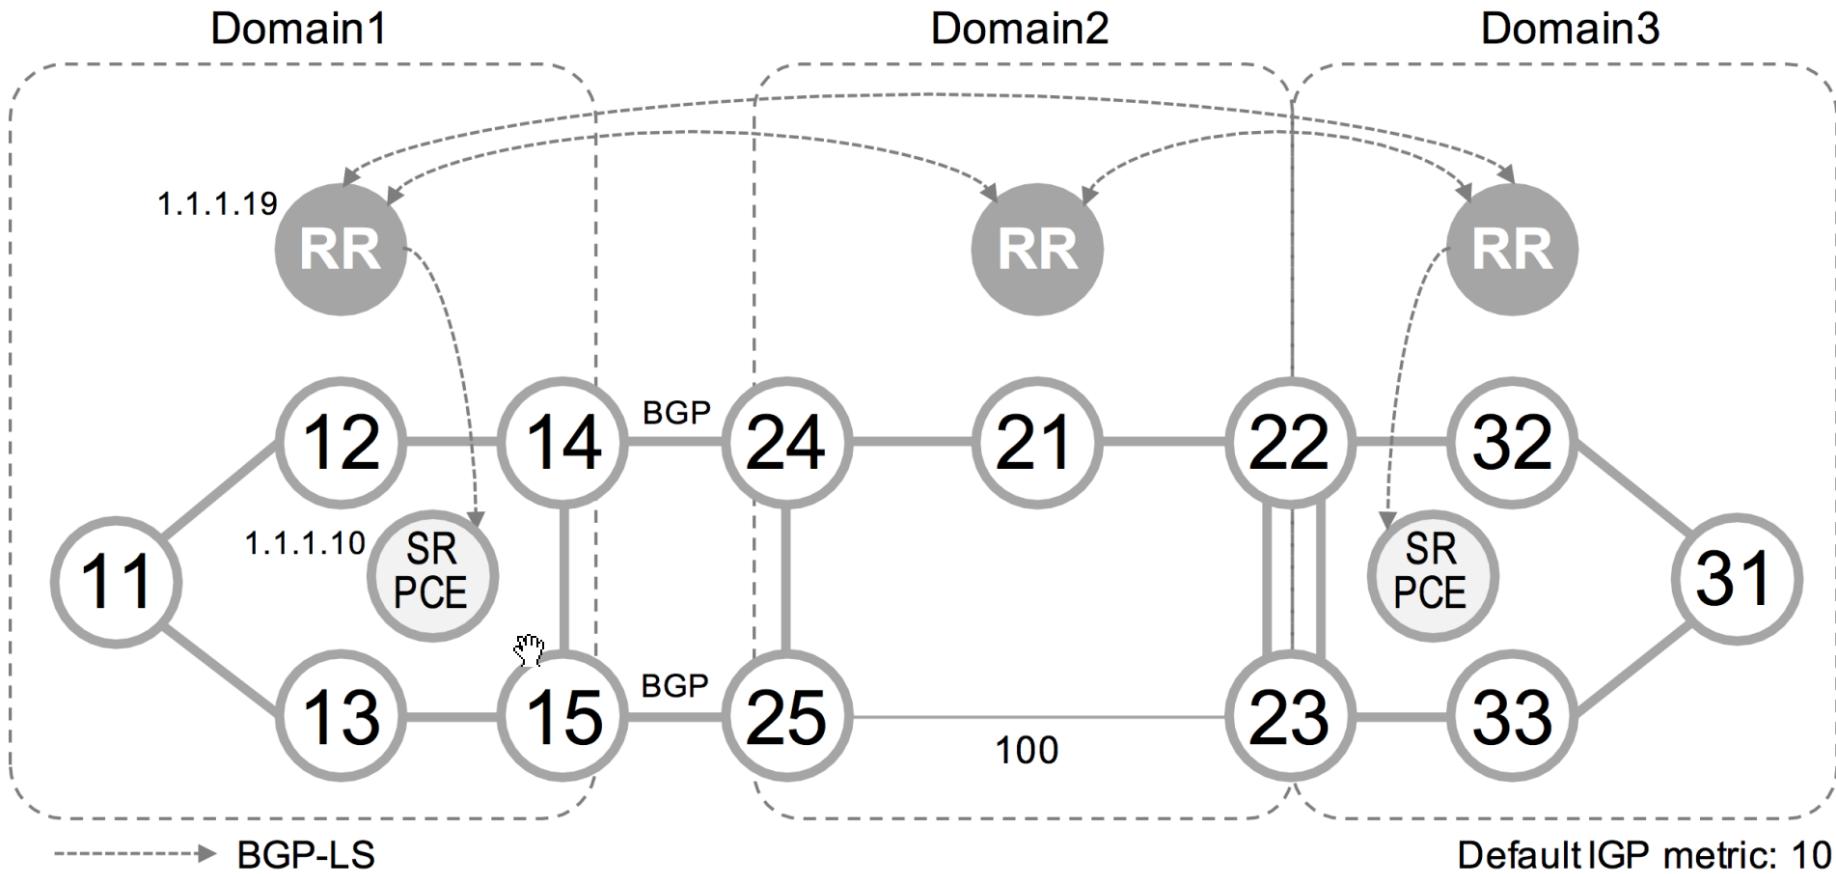
# IGP LS-DB in BGP-LS to RR



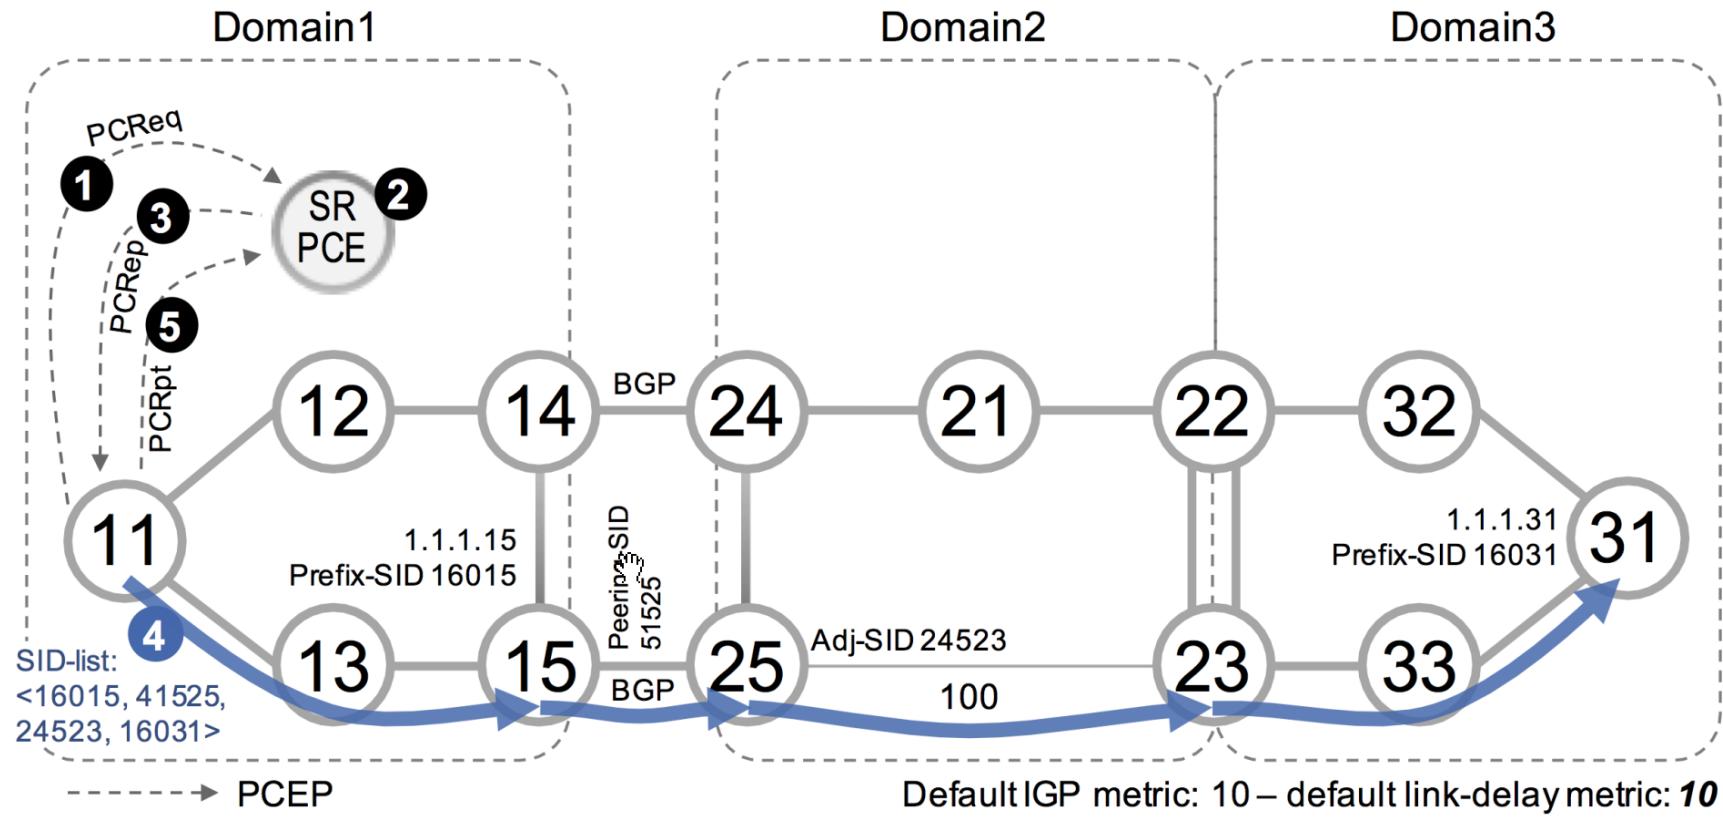
# EPE in BGP-LS to RR



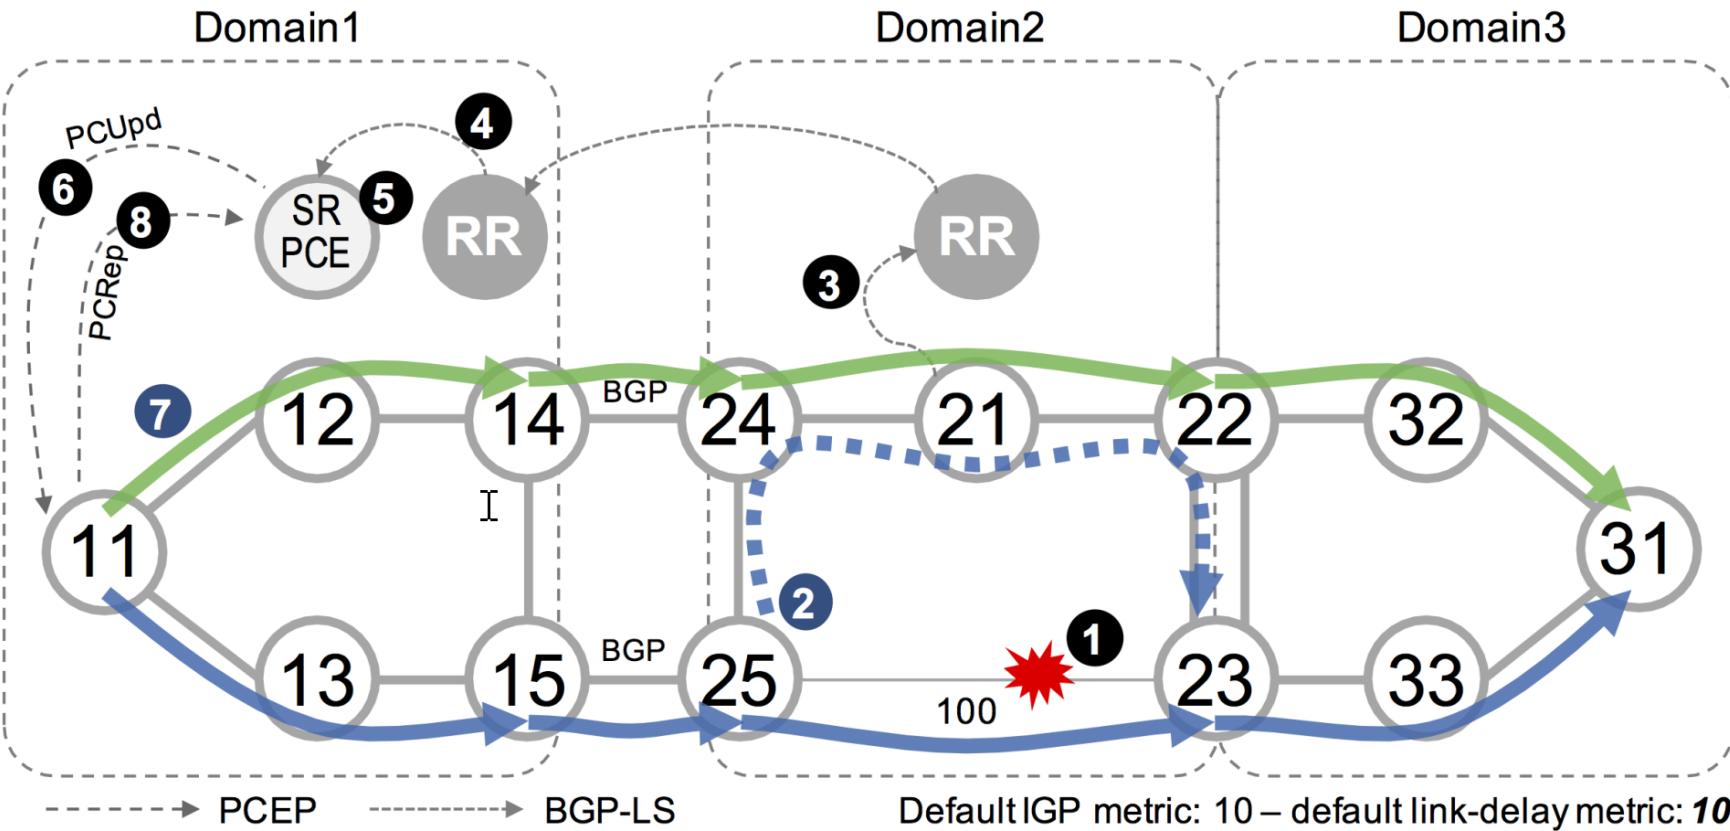
# SR PCE + BGP-LS to receive Multi-domain Topology



# Inter-domain Path PCEP

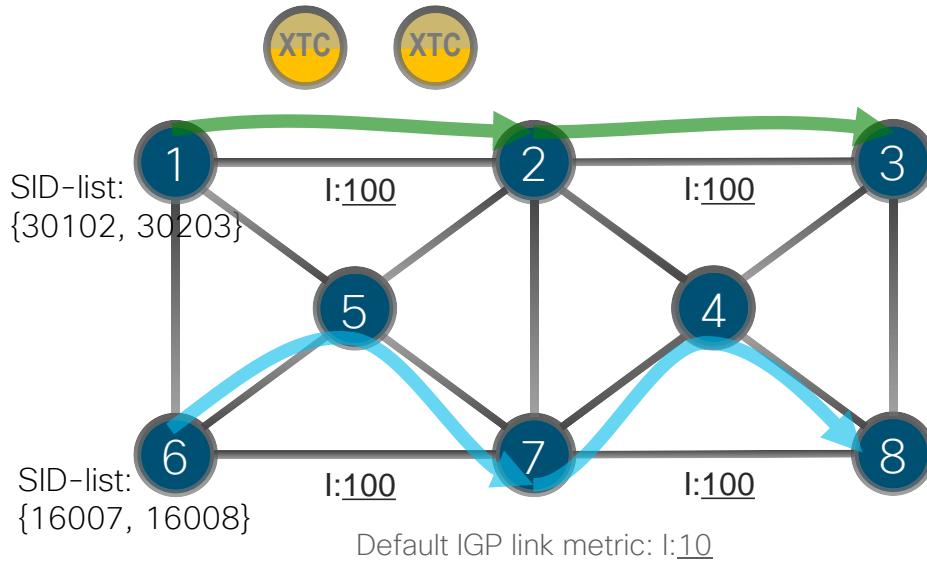


# Stateful PCE updates inter-domain path after failure



# Service Disjointness

Intra and inter domain



Node1

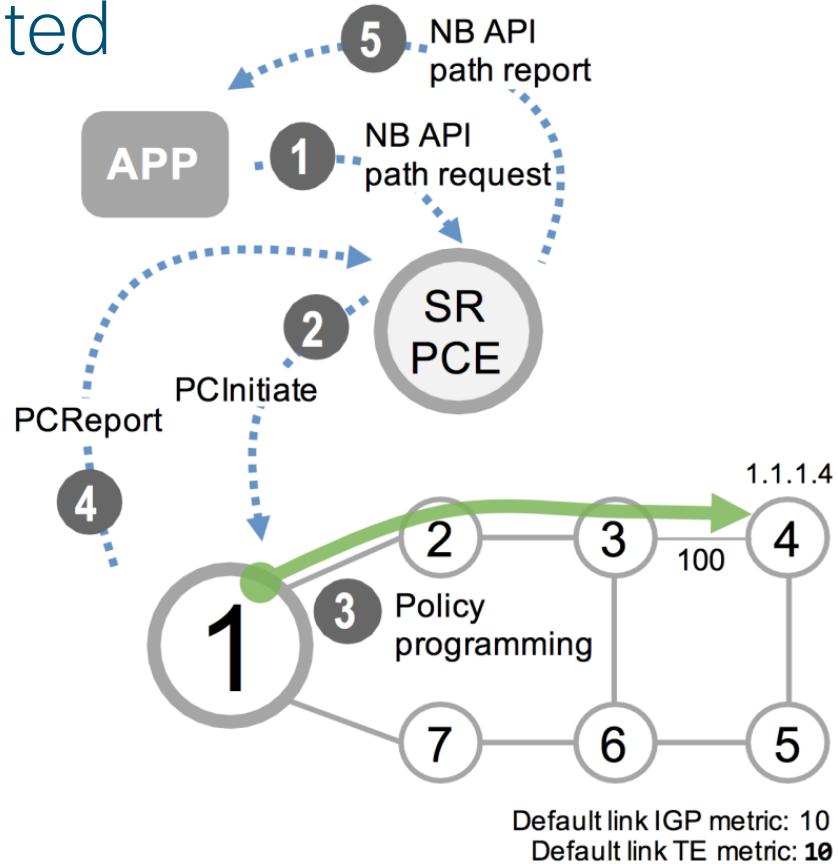
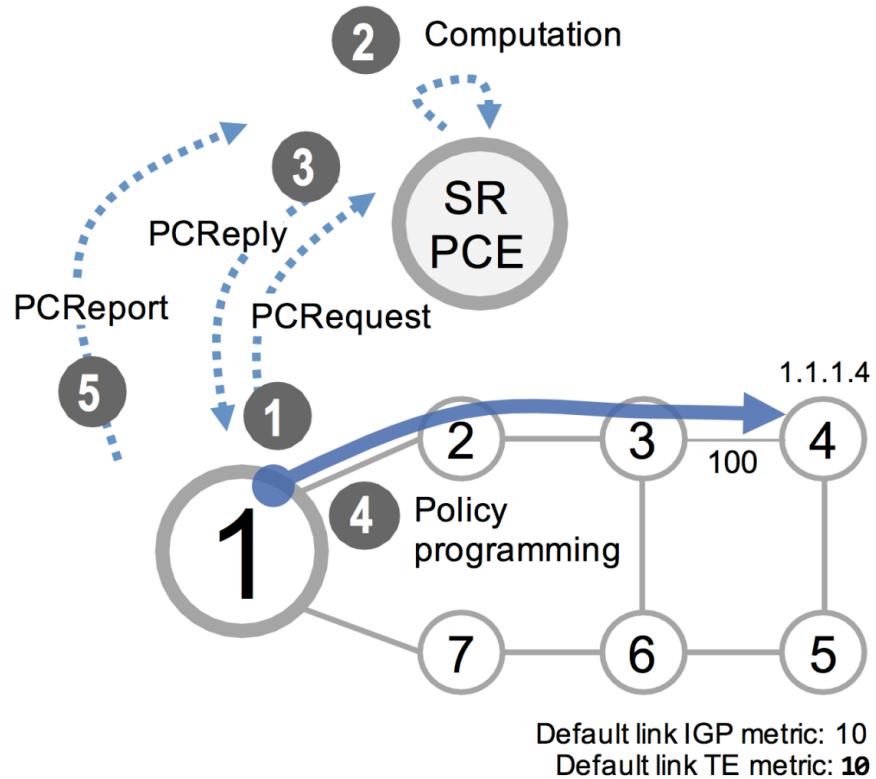
```
segment-routing
traffic-eng
policy POLICY1
color 20 end-point ipv4 1.1.1.3
candidate-paths
preference 100
dynamic mpls pce
metric
type igp
association group 1 type node
```

Node6

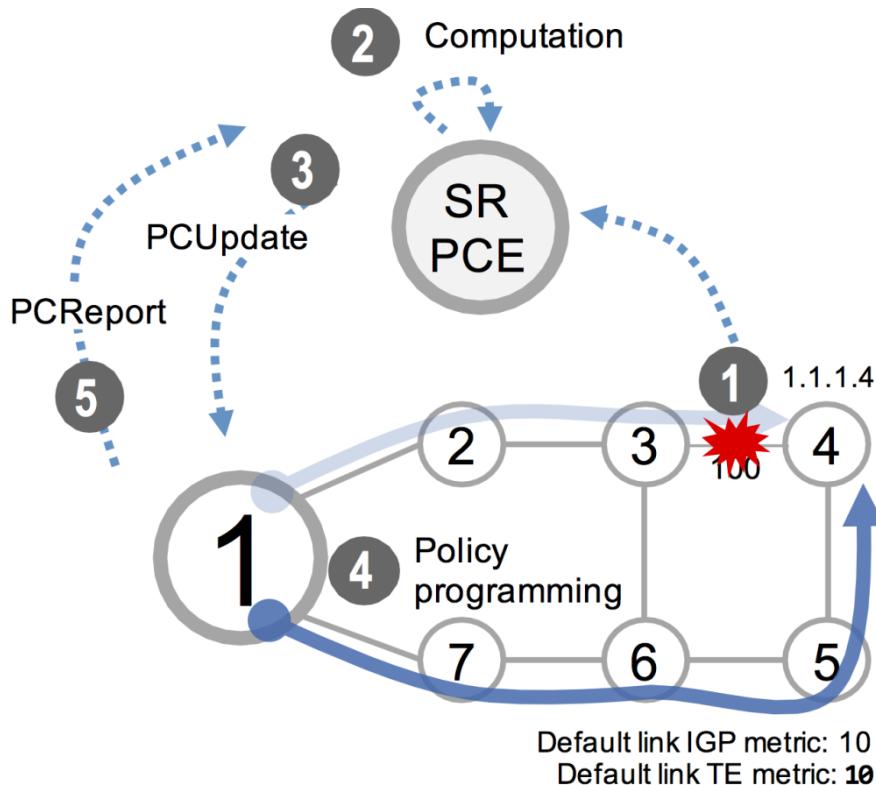
```
segment-routing
traffic-eng
policy POLICY2
color 20 end-point ipv4 1.1.1.8
candidate-paths
preference 100
dynamic mpls pce
metric
type igp
association group 1 type node
```

- Two dynamic paths between two different pairs of (head-end, end-point) must be disjoint from each other

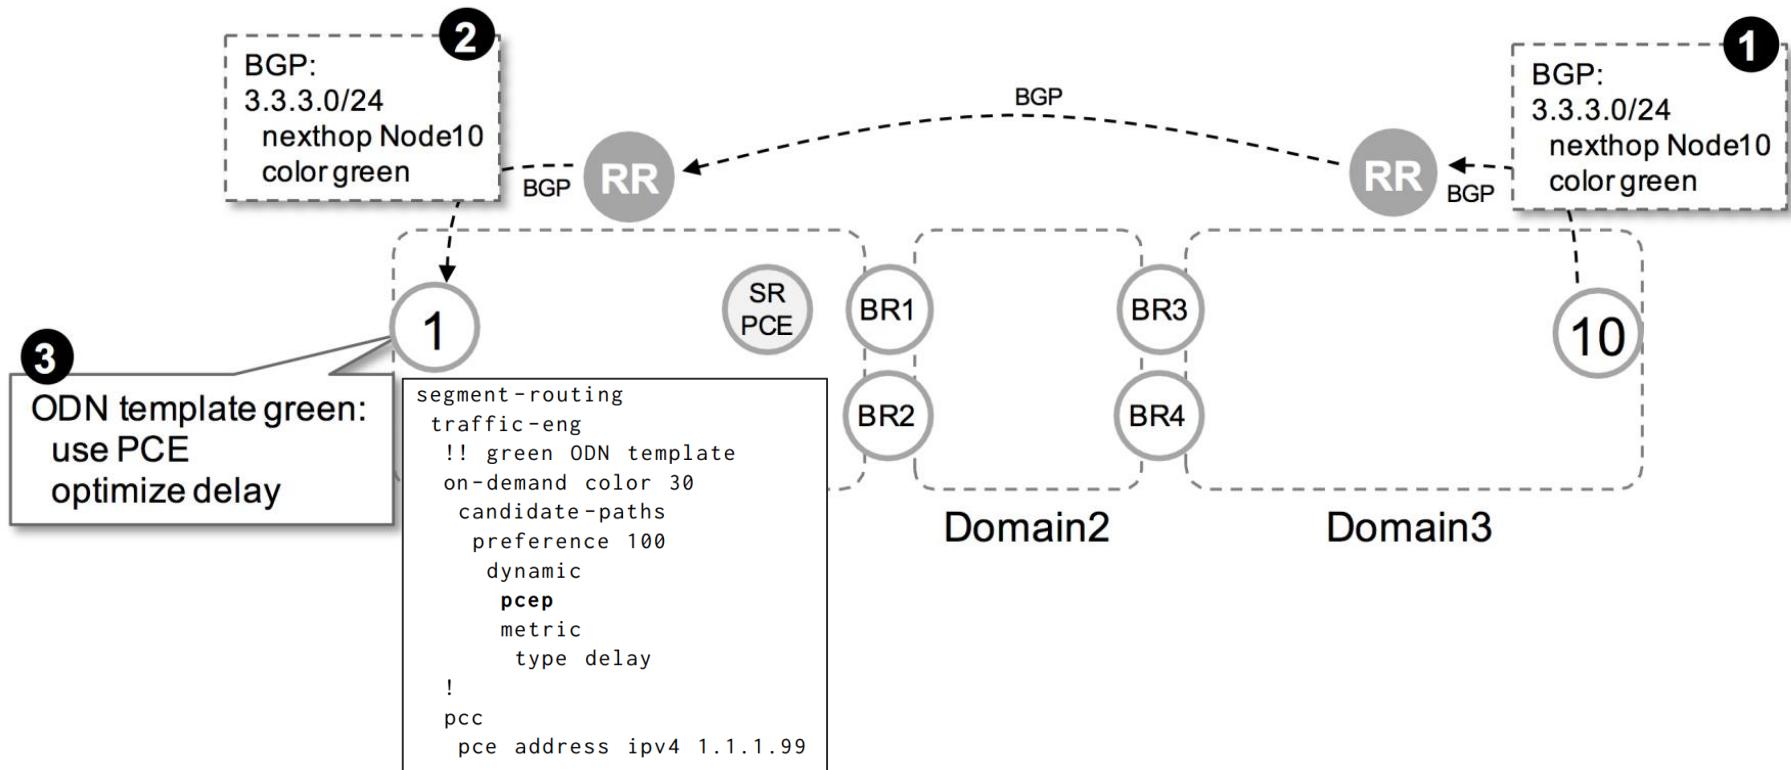
# PCEP SRTE : Headend initiated SR PCE initiated



# SR PCE Update



# SR ODN Multi-domain Topology





# Segment Routing PCE REST API

# SR PCE REST

- RP/0/0/CPU0:XTC-1#sh run pce
- Sun Feb 3 02:36:53.445 UTC
- pce
- address ipv4 1.1.1.201
- rest
- user cisco
- password encrypted 00071A150754
- !
- authentication basic
- !
- logging
- no-path
- fallback
- !
- state-sync ipv4 1.1.1.202
- password encrypted 121A0C041104

# SR PCE REST client – curl & browser

```
← → ⌂ ⓘ Not secure | 198.18.1.31:8080/lsp/delete/simple?peer=1.1.1.4&name=rest_4_5

delete-lsp "rest_4_5" (SR) on peer 1.1.1.4 (Success)

htjun@HTJUN-L90X6 MINGW64 ~
$ curl -u cisco:cisco --raw -vN "http://198.18.1.31:8080/lsp/create/simple?peer=1.1.1.5&name=rest_5_4&type=sr&source=1.1.1.5&destination=1.1.1.4&color=10&hop1-node-sid=16001&hop1-ip-local=1.1.1.1&hop2-node-sid=16002&hop2-ip-local=1.1.1.2&hop3-node-sid=16004&hop3-ip-local=1.1.1.4"
* Trying 198.18.1.31...
* TCP_NODELAY set
% Total    % Received % Xferd  Average Speed   Time     Time      Time  Current
          Dload  Upload   Total Spent  Left  Speed
 0          0       0       0      0  0 --::-- --::-- --::-- 035
create-lsp "rest_5_4" (sr) on peer 1.1.1.5 (Success)
0

create-lsp "rest_5_4" (sr) on peer 1.1.1.5 (Success) * Connected to 198.18.1.31 (198.18.1.31) port 8080 (#0)
* Server auth using Basic with user 'cisco'
> GET /lsp/create/simple?peer=1.1.1.5&name=rest_5_4&type=sr&source=1.1.1.5&destination=1.1.1.4&color=10&hop1-node-sid=16001&hop1-ip-local=1.1.1.1&hop2-node-sid=16002&hop2-ip-local=1.1.1.2&hop3-node-sid=16004&hop3-ip-local=1.1.1.4 HTTP/1.1
> Host: 198.18.1.31:8080
> Authorization: Basic Y2lzY286Y2lzY28=
> User-Agent: curl/7.57.0
> Accept: */*
>
< HTTP/1.1 200 OK
< Cache-Control: no-cache, no-store
< Content-Type: text/json; charset=utf-8
< Expires: -1
< Transfer-Encoding: chunked
< Connection: keep-alive
<
{: [64 bytes data]
100 53 0 53 0 0 53 0 --::-- --::-- --::-- 679
* Connection #0 to host 198.18.1.31 left intact
```

# SR PCE REST client - postman

The screenshot shows the Postman application interface. The title bar reads "Postman". The menu bar includes "File", "Edit", "View", and "Help". The top navigation bar has buttons for "New", "Import", "Runner", and "My Workspace". A status bar indicates "IN SYNC". A yellow banner at the top says "Hey! You're running a very old version of Postman. Our newest app has a lot more features. [Learn more](#)".

The left sidebar shows a "History" section with several collections: "NETWORK-Slice-D..." (9 requests), "NSO-GNS3" (9 requests), "SR EPE" (3 requests), "XTC" (4 requests), "SR PCE GET Topol...", "Node-4 CREATE S...", "U\_Node-4 CREATE..", and "Node-4 DELETE S...".

The main workspace displays a collection titled "U\_Node-4 CREATE SR Policy 4\_to\_5 Color 20". The request method is "GET" with the URL "http://192.168.1.3:18080/lsp/create/simple?peer=1.1.1.4&name=rest\_4\_5&type=pe=sr&source=1.1.4&destination=1.1.1.5&color=20". The "Auth" tab is selected, showing "about auth" and a note about secure variables. The "Params" tab shows the query parameters listed above. The "Send" button is highlighted.

The response panel shows a "Status: 200 OK" with "Time: 100 ms" and "Size: 212 B". The "Body" tab is selected, displaying the response content: "create-lsp "rest\_4\_5" (sr) on peer 1.1.1.4 (Success)". Other tabs include "Cookies", "Headers (5)", and "Test Results".

At the bottom, there are icons for "BUILD", "Lightbulb", "Code", "Network", and "Help".

# SR PCE REST create\_lsp parameter - 1

/lsp/create/simple – rest url

peer=1.1.1.2 – pcc (tunnel headend router) address

name=rest\_2\_4 – tunnel name, max name length is restricted to 128 characters. Name cannot contains spaces or special characters (allowed are '-', '\_', '(', ')')

type=sr – setup type. If not specified SR is assumed. Allowed values are sr or rsvp

source=1.1.1.2 – tunnel source address. If not specified "peer" address will be used instead

allow-xtc-reoptimization=0 – REST client prevents XTC to update lsp with new path e.g. in case of link down and REST client has to update path itself. Default value is 1, in that case XTC is allowed to modify lsp.

destination=1.1.1.4 – tunnel destination address

color=10 – SR policy color, mandatory when creating SR policy (not needed for rsvp lsps)

preference=10 – SR policy path preference, default value is 100, (not needed for rsvp lsps)

# SR PCE REST create\_lsp parameter -2

metric-igp=10 – using igp metric as path computation constraint

metric-te=10 – using te metric as path computation constraint

metric-latency=10 – using latency metric as path computation constraint

hop1-adj-sid=24001&hop1-ip-local=99.1.2.2&hop1-ip-remote=99.1.2.1 – path's first hop defined by adj-sid

hop2-adj-sid=28099&hop2-ip-local=99.1.3.1&hop2-ip-remote=99.1.3.3 – path's second hop defined by adj-sid

hop3-adj-sid=24003&hop3-ip-local=99.3.4.3&hop3-ip-remote=99.3.4.4 – path's third hop defined by adj-sid

hopX-node-sid=30001&hopX-ip-local=1.1.1.1

Example: (no\_space)

```
curl --raw -vN "http://<xtc-ip-addr>:8080/lsp/create/simple?peer=1.1.1.2&name=rest_2_4&type=sr
&source=1.1.1.2&destination=1.1.1.4&color=10
&hop1-adj-sid=24001&hop1-ip-local=99.1.2.2&hop1-ip-remote=99.1.2.1
&hop2-adj-sid=28099&hop2-ip-local=99.1.3.1&hop2-ip-remote=99.1.3.3
&hop3-adj-sid=24003&hop3-ip-local=99.3.4.3&hop3-ip-remote=99.3.4.4"
```

# SR PCE REST update\_lsp parameter

/lsp/update/simple – rest url

peer=1.1.1.2 – pcc (tunnel headend router) address

name=rest\_2\_4 – tunnel name, max name length is restricted to 128 characters. Name cannot contain spaces or special characters (allowed are '-', '\_', '(', ')')

type=sr – setup type. If not specified SR is assumed

hop1-adj-sid=24001&hop1-ip-local=99.1.2.2&hop1-ip-remote=99.1.2.1 – path's first hop defined by adj-sid

hop2-adj-sid=28099&hop2-ip-local=99.1.3.1&hop2-ip-remote=99.1.3.3 – path's second hop defined by adj-sid

hop3-adj-sid=24003&hop3-ip-local=99.3.4.3&hop3-ip-remote=99.3.4.4 – path's third hop defined by adj-sid

Example: (no\_space)

```
curl --raw -vN "http://<xtc-ip-addr>:8080/lsp/update/simple?peer=1.1.1.2&name=rest_2_4
&hop1-adj-sid=24001&hop1-ip-local=99.1.2.2&hop1-ip-remote=99.1.2.1
&hop2-adj-sid=28099&hop2-ip-local=99.1.3.1&hop2-ip-remote=99.1.3.3
&hop3-adj-sid=24003&hop3-ip-local=99.3.4.3&hop3-ip-remote=99.3.4.4"
```

# SR PCE REST delete\_lsp parameter

/lsp/delete/simple – rest url

peer=1.1.1.2 – pcc (tunnel headend router) address

name=rest\_2\_4 – tunnel name, max name length is restricted to 128 characters. Name cannot contains spaces or special characters (allowed are '-', '\_', '(', ')')

Example:

```
curl --raw -vN "http://<xtc-ip-addr>:8080/lsp/delete/simple?peer=1.1.1.2&name=rest_2_4"
```



# Steering Traffic over SR Policy

# Steering Traffic over SR Policy

- Automated Steering
- On Demand Nexthop (ODN)
- Static route via local SR Policy
- SR Policy Autoroute
- VPWS, VPLS and PWHE Preferred-path over SR Policy
- EVPN-VPWS Preferred-path over SR Policy

# Steering Traffic over SR Policy Static Route via Local SR Policy

```
router static
  address-family ipv4 unicast
    202.202.1.0/24 sr-policy POLICY-31
```

```
segment-routing
  traffic-eng
    policy POLICY-31
      color 3001 end-point ipv4 100.100.4.20 < ----- local SR Policy has (color:3001, EP:100.100.4.20)
      candidate-paths
        preference 100
        explicit segment-list Path-31
```

# Steering Traffic over SR Policy SR Policy Autoroute

```
segment-routing
traffic-eng
policy POLICY-31
color 31 end-point ipv4 100.100.4.20
    < ----- L3VPN or EVPN next-hop
autoroute
include all
!
candidate-paths
preference 100
explicit segment-list Path-31
!
```

```
segment-routing
traffic-eng
policy POLICY-31
color 31 end-point ipv4 100.100.4.20
autoroute
include ipv4 98.1.1.0/24
include ipv4 99.1.1.0/24
    < ----- IGP subnets behind 100.100.4.20 router !
candidate-paths
preference 100
explicit segment-list Path-31
!
```

# Steering Traffic over SR Policy

## VPWS, VPLS and PWHE Preferred-path over SR Policy

```
segment-routing
traffic-eng
policy POLICY-100
color 100 end-point ipv4 2.2.2.2
candidate-paths
preference 100
explicit segment-list PE1-P1-PE2
!
policy POLICY-200
color 200 end-point ipv4 2.2.2.2
candidate-paths
preference 100
explicit segment-list PE1-P2-PE2
!
policy POLICY-300
color 300 end-point ipv4 2.2.2.2
candidate-paths
preference 100
explicit segment-list PE1-P3-PE2
!
```

```
l2vpn
pw-class PW-100
encapsulation mphls
preferred-path sr-te policy POLICY-100
!
pw-class PW-200
encapsulation mpls
preferred-path sr-te policy POLICY-200
!
pw-class PW-300
encapsulation mpls
preferred-path sr-te policy POLICY-300
!
```

# Steering Traffic over SR Policy

## VPWS, VPLS and PWHE Preferred-path over SR Policy

```
xconnect group PWHE-100 < ----- PWHE preferred-path over SR Policy example
p2p PWHE-100
interface PW-Ether1
neighbor ipv4 2.2.2.2 pw-id 1003
pw-class PW-300
!
xconnect group VPWS-100 < ----- VPWS preferred-path over SR Policy exmaple
p2p VPWS-100
interface TenGigE0/9/0/1.1001
neighbor ipv4 2.2.2.2 pw-id 1001
pw-class PW-100
!
bridge group BD-100 < ----- VPLS preferred-path over SR Policy exmaple
bridge-domain BG-100
interface TenGigE0/9/0/1.1002
!
neighbor 2.2.2.2 pw-id 1002
pw-class PW-200
!
```

# Steering Traffic over SR Policy

## EVPN-VPWS Preferred-path over SR Policy

```
segment-routing
traffic-eng
policy POLICY-400
color 400 end-point ipv4 2.2.2.2
candidate-paths
preference 100
explicit segment-list PE1-P4-PE2
!
l2vpn
pw-class PW-400
encapsulation mpls
preferred-path sr-te policy POLICY-400
!
xconnect group EVPN-VPWS-100
p2p EVPN-VPWS-100
interface TenGigE0/9/0/1.1004
neighbor evpn evi 10 target 5004 source 1004 < ----- EVPN-VPWS preferred-path over SR Policy exmaple
pw-class PW-400
!
!
```

